

r1: Risk-First Software Development

Rob Moffat

Risk-First Series

Risk First: The Menagerie

Book one of the **Risk-First** series argues the case for viewing *all* of the activities on a software project through the lens of *managing risk*. It introduces the menagerie of different risks you're likely to meet on a software project, naming and classifying them so that we can try to understand them better.

Risk First: Tools and Practices

Book two of the **Risk First** series explores the relationship between software project risks and the tools and practices we use to mitigate them. Due for publication in 2020.

Online

Material for the books is freely available to read, drawn from risk-first.org.

Published By

Kite9 Ltd.
14 Manor Close
Colchester
CO6 4AR

Risk First: The Menagerie

By Rob Moffat

Copyright © 2018 Kite9 Ltd.

All rights reserved. No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law. For permission requests, write to the publisher, addressed “Attention: Permissions Coordinator,” at the address below.

ISBN: tbd.

Credits

tbd

Cover Images: Biodiversity Heritage Library. *Biologia Centrali-Americanana. Insecta. Rhynchota. Hemiptera-Homoptera. Volume 1 (1881-1905)*

Cover Design By P. Moffat (peter@petermoffat.com)

Thanks to:

Dedicated to blan

Contents

Preface	v
Executive Summary	ix
I Introduction	1
1 A Simple Scenario	3
2 Development Process	7
3 All Risk Management	13
4 Software Project Scenario	17
5 Risk Theory	21
6 Meeting Reality	29
7 Cadence	35
8 A Conversation	39
9 De Risking	43
II Risk	45
10 Risk Landscape	47
11 Feature Risk	51
12 Complexity Risk	59
13 Communication Risk	73
14 Dependency Risk	93
15 Schedule Risk	103

16 Software Dependency Risk	111
17 Process Risk	125
18 Boundary Risk	139
19 Agency Risk	155
20 Coordination Risk	161
21 Map And Territory Risk	181
22 Operational Risk	195
23 Staging And Classifying	205
24 Stories Of Failure	209
III Preview	211
Glossary	215

Preface

Welcome to Risk-First!

Let's cover some of the big questions up-front: The why, what, who, how and where of *The Menagerie*.

Why

"Scrum, Waterfall, Lean, Prince2: what do they all have in common?"

I've started this because, on my career journey, I've noticed that the way I do things doesn't seem to match up with the way the books say it should be done. And, I found this odd and wanted to explore it further. Hopefully, you, the reader, will find something of use in this.

I started with this observation: *Development Teams* put a lot of faith in methodology. Sometimes, this faith is often so strong it borders on religion. (Which in itself is a concern.) For some, this is **Prince2**. For others, it might be **Lean** or **Agile**.

Developers put a lot of faith in *particular tools* too. Some developers are pro-or-anti-Java, others are pro-or-anti-XML. All of them have their views coloured by their *experiences* (or lack of) with these tools. Was this because their past projects *succeeded* or *failed* because of them?

As time went by, I came to see that the choice of methodology, process or tool was contingent on the problem being solved, and the person solving the problem. We don't face a shortage of tools in IT, or a shortage of methodologies, or a shortage of practices. Essentially, that all the tools and methodologies that the industry had supplied were there to help *minimize the risk of my project failing*.

This book considers that perspective: that building software is all about *managing risk*, and that these methodologies are acknowledgements of this fact, and they differ because they have *different ideas* about which are the most important *risks to manage*.

What This Is

Hopefully, after reading this, you'll come away with:

- An appreciation of how risk underpins everything we do as developers, whether we want it to or not.

- A framework for evaluating methodologies, tools and practices and choosing the right one for the task-at-hand.
- A recontextualization of the software process as being an exercise in mitigating different kinds of risk.
- The tools to help you decide when a methodology or tool is *letting you down*, and the vocabulary to argue for when it's a good idea to deviate from it.

This is not intended to be a rigorously scientific work: I don't believe it's possible to objectively analyze a field like software development in any meaningful, statistically significant way. (For one, things just change **too fast**.)

"I have this Pattern"

Does that diminish it? If you have visited the **TVTropes** website, you'll know that it's a set of web-pages describing *common patterns* of narrative, production, character design etc. to do with fiction. For example:

tbd.

Is it scientific? No. Is it correct? Almost certainly. TVTropes is a set of *empirical patterns* for how stories on TV and other media work. It's really useful, and a lot of fun. (Warning: it's also incredibly addictive).

In the same way, tbd, the tbd published a book called "Design Patterns: tbd". Which shows you patterns of *structure* within Object-Oriented programming:

tbd.

Patterns For Practitioners

This book aimed to be a set of *useful* patterns which practitioners could use in their software to achieve certain goals. "I have this pattern" was a phrase used to describe how they had seen a certain set of constraints before, and how they had solved it in software.

This book was a set of experts handing down their battle-tested practices for other developers to use, and, whether you like patterns or not, knowing them is an important part of being a software developer, as you will see them used everywhere you go and probably use them yourself.

In the same way, this book aims to be a set of *Patterns for Software Risk*. Hopefully after reading this book, you will see where risk hides in software projects, and have a name for it when you see it.

Towards a Periodic Table

In the latter chapters of "The Menagerie" we try to assemble these risk patterns into a cohesive whole. Projects fail because of risks, and risks arise from predictable sources.

What This is Not

This is not intended to be a rigorously scientific work: I don't believe it's possible to objectively analyze a field like software development in any meaningful, statistically significant way. (For one, things just change **too fast**.)

Neither is this site isn't going to be an exhaustive guide of every possible software development practice and methodology. That would just be too long and tedious.

Neither is this really a practitioner's guide to using any particular methodology: If you've come here to learn the best way to do **Retrospectives**, then you're in the wrong place. There are plenty of places you can find that information already. Where possible, this site will link to or reference concepts on Wikipedia or the wider internet for further reading on each subject.

Who

This work is intended to be read by people who work on software projects, and especially those who are involved in managing software projects.

If you work collaboratively with other people in a software process, you should find Risk-First a useful lexicon of terms to help describe the risks you face.

But here's a warning: This is going to be a depressing book to read. It is book one of a two-book series, but in **Book One** you only get to meet the bad guy.

While **Book Two** is all about *how to succeed*, This book is all about how projects *fail*. In it, we're going to try and put together a framework for understanding the risk of failure, in order that we can reconstruct our understanding of our activities on a project based on avoiding it.

So, if you are interested in *avoiding your project failing*, this is probably going to be useful knowledge.

For Developers

Risk-First is a tool you can deploy to immediately improve your ability to plan your work.

Frequently, as developers we find software methodologies “done to us” from above. Risk-First is a toolkit to help *take apart* methodologies like **Scrum**, **Lean** and **Prince2**, and understand them. Methodologies are *bicycles*, rather than *religions*. Rather than simply *believing*, we can take them apart and see how they work.

For Project Managers and Team Leads

All too often, Project Managers don't have a full grasp of the technical details of their projects. And this is perfectly normal, as the specialization belongs below them. However, projects fail because risks materialize, and risks materialize because the devil is in those details.

This seems like a lost cause, but there is hope: the ways in which risks materialize on technical projects is the same every time. With Risk-First we are attempting to name each of these types

of risk, which allows for a dialog with developers about which risks they face, and the order they should be tackled.

Risk-First allows a project manager to pry open the black box of development and talk with developers about their work, and how it will affect the project. It is another tool in the (limited) arsenal of techniques a project manager can bring to bear on the task of delivering a successful project.

How

One of the original proponents of the **Agile Manifesto**, Kent Beck, begins his book **Extreme Programming** by stating:

“It’s all about risk” > Kent Beck

This is a promising start. From there, he introduces his methodology, **Extreme Programming**, and explains how you can adopt it in your team, the features to observe and the characteristics of success and failure. However, while *Risk* has clearly driven the conception of **Extreme Programming**, there is no clear model of software risk underpinning the work, and the relationship between the practices he espouses and the risks he is avoiding are hidden.

In this book, we are going to introduce a model of software project risk. This means that in **Book Two** (Risk-First: Tools and Practices), we can properly analyse **Extreme Programming** (and Scrum, Waterfall, Lean and all the others) and *understand* what drives them. Since they are designed to deliver successful software projects, they must be about mitigate risks, and we will uncover *exactly which risks are mitigated and how they do it*.

Where

All of the material for this book is available Open Source on github.com¹, and at the risk-first.org² website. Please visit, your feedback is appreciated.

There is no compulsion to buy a print or digital version of the book, but we’d really appreciate the support. So, if you’ve read this and enjoyed it, how about buying a copy for someone else to read?

A Note on References

Where possible, references are to the Wikipedia³ website. Wikipedia is not perfect. There is a case for linking to the original articles and papers, but by using Wikipedia references are free and easy for everyone to access, and hopefully will exist for a long time into the future.

On to **The Executive Summary**

¹<https://github.com>

²<https://risk-first.org>

³<https://wikipedia.org>

Executive Summary

1. There are Lots of Ways of Running Software Projects

There are lots of different ways to look at a project. For example, metrics such as “number of open tickets”, “story points”, “code coverage” or “release cadence” give us a numerical feel for how things are going and what needs to happen next. We also judge the health of projects by the practices used on them - **Continuous Integration**, **Unit Testing** or **Pair Programming**, for example.

Software methodologies, then, are collections of tools and practices: “Agile”, “Waterfall”, “Lean” or “Phased Delivery” (for example) all suggest different approaches to running a project, and are opinionated about the way they think projects should be done and the tools that should be used.

None of these is necessarily more “right” than another- they are suitable on different projects at different times.

A key question then is: **how do we select the right tools for the job?**

2. We can Look at Projects in Terms of Risks

One way to examine a project in-flight is by looking at the risks it faces.

Commonly, tools such as RAID logs and RAG status reporting are used. These techniques should be familiar to project managers and developers everywhere.

However, the Risk-First view is that we can go much further: that each item of work being done on the project is mitigating a particular risk. Risk isn’t something that just appears in a report, it actually drives *everything we do*.

For example:

- A story about improving the user login screen can be seen as reducing *the risk of users not signing up*.
- A task about improving the health indicators could be seen as mitigating *the risk of the application failing and no-one reacting to it*.
- Even a task as basic as implementing a new function in the application is mitigating *the risk that users are dissatisfied and go elsewhere*.

One assertion of Risk-First** therefore, is that every action you take on a project is to mitigate some risk.**

3. We Can Break Down Risks on a Project Methodically

Although risk is usually complicated and messy, other industries have found value in breaking down the types of risks that affect them and addressing them individually.

For example:

- In manufacturing, *tolerances* allow for calculating the likelihood of defects in production.
- In finance, reserves are commonly set aside for the risks of stock-market crashes, and teams are structured around monitoring these different risks.
- The insurance industry is founded on identifying particular risks and providing financial safety-nets for when they occur, such as death, injury, accident and so on.

Software risks are difficult to quantify, and mostly, the effort involved in doing so *exactly* would outweigh the benefit. Nevertheless, there is value in spending time building *classifications of risk for software*. That's what Risk-First does: describes the set of *risk patterns* we see every day on software projects.

With this in place, we can:

- Talk about the types of risks we face on our projects, using an appropriate language.
- Expose **Hidden Risks** that we hadn't considered before.
- Weigh the risks against each other, and decide which order to tackle them.

4. We Can Analyse Tools and Techniques in Terms of how they Mitigate Risk

If we accept the assertion above that *all* the actions we take on a project are about mitigating risks, then it stands to reason that the tools and techniques available to us on a project are there for mitigating different types of risks.

For example:

- If we do a **Code Review**, we are partly trying to mitigate the risks of bugs slipping through into production, and also mitigate the **Key-Man Risk** of knowledge not being widely-enough shared.
- If we write **Unit Tests**, we're also mitigating the risk of bugs going to production, but we're also mitigating against future changes breaking our existing functionality.
- If we enter into a **contract with a supplier**, we are mitigating the risk of the supplier vanishing and leaving us exposed. With the contract in place, we have legal recourse against this risk.

Different tools are appropriate for mitigating different types of risks.

5. Different Methodologies for Different Risk Profiles

In the same way that our tools and techniques are appropriate to dealing with different risks, the same is true of the methodologies we use on our projects. We can use a Risk-First approach to examine the different methodologies, and see which risks they address.

For example:

- **Agile** methodologies prioritise mitigating the risk that requirements capture is complicated, error-prone and that requirements change easily.
- **Waterfall** takes the view that coding effort is an expensive risk, and that we should build plans up-front to avoid it.
- **Lean** takes the view that risk lies in incomplete work and wasted work, and aims to minimize that.

Although many developers have a methodology-of-choice, the argument here is that there are tradeoffs with all of these choices. Methodologies are like *bicycles*, rather than *religions*. Rather than simply *believing*, we can take them apart and see how they work.

We can place methodologies within a framework, and show how choice of methodology is contingent on the risks faced.

6. Driving Development With a Risk-First Perspective

We have described a model of risk within software projects, looking something like this:

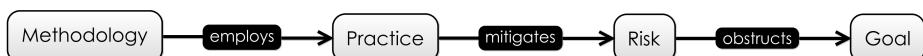


Figure 1: Methdologies, Risks, Practices

How do we take this further?

The first idea we explore is that of the **Risk Landscape**: Although the software team can't remove risk from their project, they can take actions that move them to a place in the **Risk Landscape** where the risks on the project are more favourable than where they started.

From there, we examine basic risk archetypes you will encounter on the software project, to build up a **Taxonomy of Software Risk**, and look at which specific tools you can use to mitigate each kind of risk.

Then, we look at different software practices, and how they mitigate various risks. Beyond this we examine the question: *how can a Risk-First approach inform the use of this technique?*

For example:

- If we are introducing a **Sign-Off** in our process, we have to balance the risks it *mitigates* (coordination of effort, quality control, information sharing) with the risks it *introduces* (delays and process bottlenecks).
- If we have **Redundant Systems**, this mitigates the risk of a *single point of failure*, but introduces risks around *synchronizing data and communication* between the systems.
- If we introduce **Process**, this may make it easier to *coordinate as a team* and *measure performance* but may lead to bureaucracy, focusing on the wrong goals or over-rigid interfaces to those processes.

Risk-First aims to provide a framework in which we can *analyse these choices* and weigh up *accepting* versus *mitigating* risks.

Still interested? Then dive into reading the introduction.

Part I

Introduction

Chapter 1

A Simple Scenario

First up, I'm going to introduce a simple model for thinking about risk.

A Simple Scenario

Lets for a moment forget about software completely, and think about *any endeavor at all* in life. It could be passing a test, mowing the lawn or going on holiday. Choose something now. I'll discuss from the point of view of "cooking a meal for some friends", but you can play along with your own example.

Goal In Mind

Now, in this endeavour, we want to be successful. That is to say, we have a **Goal In Mind**: we want our friends to go home satisfied after a decent meal, and not to feel hungry. As a bonus, we might also want to spend time talking with them before and during the meal. So, now to achieve our **Goal In Mind** we *probably* have to do some tasks.

If we do nothing, our friends will turn up and maybe there's nothing in the house for them to eat. Or maybe, the thing that you're going to cook is going to take hours and they'll have to sit around and wait for you to cook it and they'll leave before it's ready. Maybe you'll be some ingredients short, or maybe you're not confident of the steps to prepare the meal and you're worried about messing it all up.

Attendant Risk

These *nagging doubts* that are going through your head I'll call the **Attendant Risks**: they're the ones that will occur to you as you start to think about what will happen.

When we go about preparing this wonderful evening, we can with these risks and try to mitigate them: shop for the ingredients in advance, prepare parts of the meal, maybe practice the cooking in advance. Or, we can wing it, and sometimes we'll get lucky.

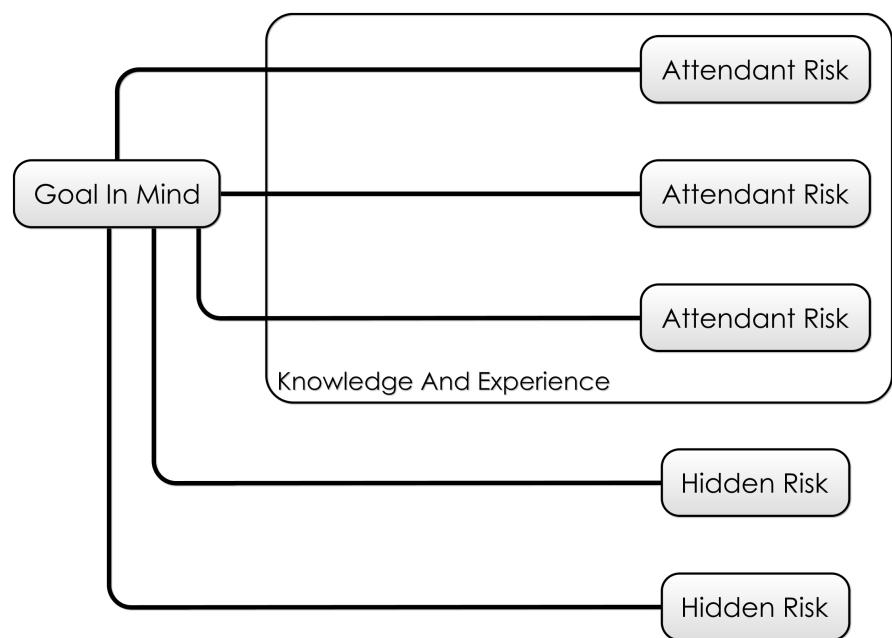


Figure 1.1: Goal In Mind, with the risks you know about

How much effort we expend on mitigating **Attendant Risks** depends on how great we think they are: for example, if you know it's a 24-hour shop, you'll probably not worry too much about getting the ingredients well in advance (although, the shop *could still be closed*).

Hidden Risks

There are also hidden **Attendant Risks** that you might not know about: if you're poaching eggs for dinner, you might know that fresh eggs poach best. These are the "Unknown Unknowns" of Rumsfeld's model¹.

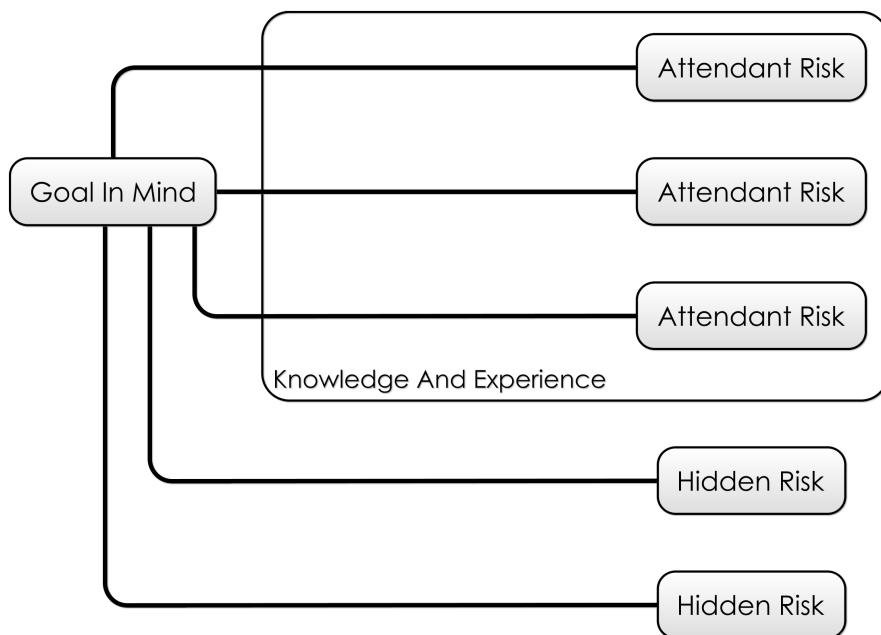


Figure 1.2: Goal In Mind, the risks you know about and the ones you don't

Different people will evaluate the risks differently. (That is, worry about them more or less.) They'll also *know* about different risks. They might have cooked the recipe before, or organised lots more dinner parties than you.

How we evaluate the risks, and which ones we know about depends on our **knowledge** and **experience**, then. And that varies from person to person (or team to team). Lets call this our **Internal Model**, and it's something we build on and improve with experience (of organising dinner parties, amongst everything else).

¹https://en.wikipedia.org/wiki/There_are_known_unknowns

Model Meets Reality

As the dinner party gets closer, we make our preparations, and the inadequacies of the **Internal Model** become apparent, and we learn what we didn't know. The **Hidden Risks** reveal themselves; things we were worried about may not materialise, things we thought would be minor risks turn out to be greater.

Our model is forced into contact with reality, and the model changes.

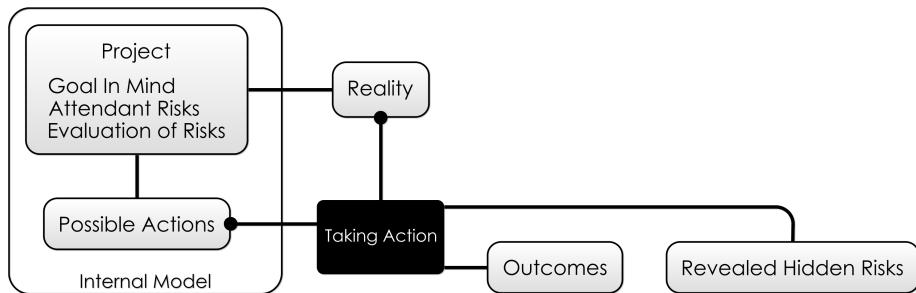


Figure 1.3: How taking action affects Reality, and also changes your internal model

If we had a good model, and took the right actions, we should see positive outcomes. If we failed to mitigate risks, or took inappropriate actions, we'll probably see negative outcomes.

On To Software

In this website, we're going to look at the risks in the software process and how these are mitigated by the various methodologies you can choose from.

Let's examine the scenario of a new software project, and expand on the simple model being outlined above: instead of a single person, we are likely to have a team, and our model will not just exist in our heads, but in the code we write.

On to Development Process

Chapter 2

Development Process

In the **previous section** we looked at a simple model for risks on any given activity.

Now, let's look at the everyday process of developing *a new feature* on a software project, and see how our risk model informs it.

An Example Process

Let's ignore for now the specifics of what methodology is being used - we'll come to that later. Let's say your team have settled for a process something like the following:

1. **Specification:** A new feature is requested somehow, and a business analyst works to specify it.
 2. **Code And Unit Test:** A developer writes some code, and some unit tests.
 3. **Integration:** They integrate their code into the code base.
 4. **UAT:** They put the code into a User Acceptance Test (UAT) environment, and user(s) test it.
- ... All being well, the code is released to production.

Now, it might be waterfall, it might be agile, we're not going to commit to specifics at this stage. It's probably not perfect, but let's just assume that *it works for this project* and everyone is reasonably happy with it.

I'm not saying this is the *right* process, or even a *good* process: you could add code review, a pilot, integration testing, whatever. We're just doing some analysis of *what process gives us*.

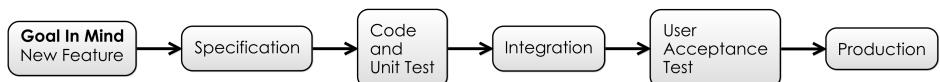


Figure 2.1: A Simple Development Process

What's happening here? Why these steps?

Minimizing Risks - Overview

I am going to argue that this entire process is *informed by software risk*:

1. We have a *business analyst* who talks to users and fleshes out the details of the feature properly. This is to minimize the risk of **building the wrong thing**.
2. We *write unit tests* to minimize the risk that our code **isn't doing what we expected, and that it matches the specifications**.
3. We *integrate our code* to minimize the risk that it's **inconsistent with the other, existing code on the project**.
4. We have *acceptance testing* and quality gates generally to **minimize the risk of breaking production**, somehow.

We could skip all those steps above and just do this:

1. Developer gets wind of new idea from user, logs onto production and changes some code directly.

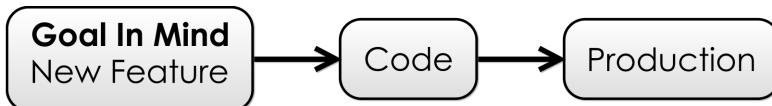


Figure 2.2: A Dangerous Development Process

We can all see this would be a disaster, but why?

Two reasons:

1. You're meeting reality all-in-one-go: all of these risks materialize at the same time, and you have to deal with them all at once.
2. Because of this, at the point you put code into the hands of your users, your **Internal Model** is at its least-developed. All the **Hidden Risks** now need to be dealt with at the same time, in production.

Applying the Model

Let's look at how our process should act to prevent these risks materializing by considering an unhappy path, one where at the outset, we have lots of **Hidden Risks** ready to materialize. Let's say a particularly vocal user rings up someone in the office and asks for new **Feature X** to be added to the software. It's logged as a new feature request, but:

- Unfortunately, this feature once programmed will break an existing **Feature Y**.
- Implementing the feature will use some api in a library, which contains bugs and have to be coded around.
- It's going to get misunderstood by the developer too, who is new on the project and doesn't understand how the software is used.
- Actually, this functionality is mainly served by **Feature Z...**
- which is already there but hard to find.



Figure 2.3: Development Process - Hidden Risks

-or-

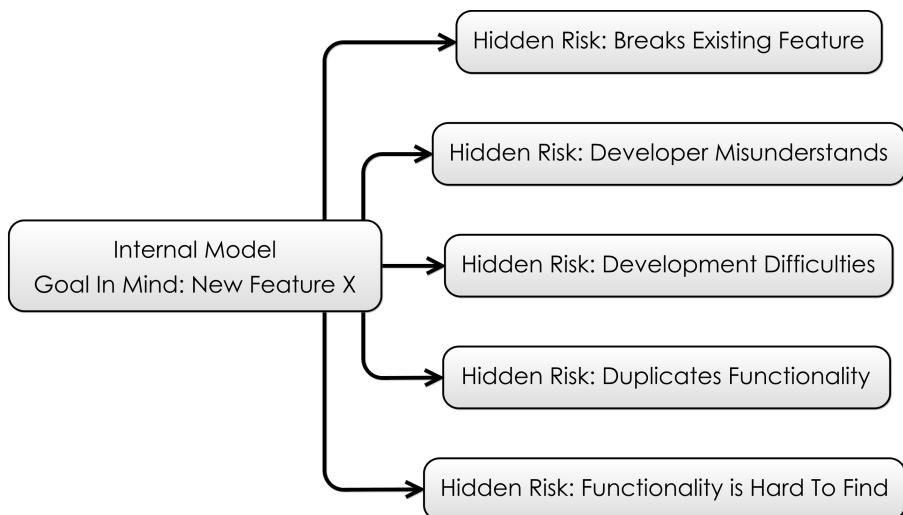


Figure 2.4: Development Process - Hidden Risks

This is a slightly contrived example, as you'll see. But let's follow our feature through the process and see how it meets reality slowly, and the hidden risks are discovered:

Specification

The first stage of the journey for the feature is that it meets the Business Analyst (BA). The purpose of the BA is to examine new goals for the project and try to integrate them with *reality*

as they understands it. A good BA might take a feature request and vet it against the internal logic of the project, saying something like:

- “This feature doesn’t belong on the User screen, it belongs on the New Account screen”
- “90% of this functionality is already present in the Document Merge Process”
- “We need a control on the form that allows the user to select between Internal and External projects”

In the process of doing this, the BA is turning the simple feature request *idea* into a more consistent, well-explained *specification* or *requirement* which the developer can pick up. But why is this a useful step in our simple methodology? From the perspective of our **Internal Model**, we can say that the BA is responsible for:

- Trying to surface **Hidden Risks**
- Trying to evaluate **Attendant Risk** and make it clear to everyone on the project.

Hopefully, after this stage, our **Internal Model** might look something like this:

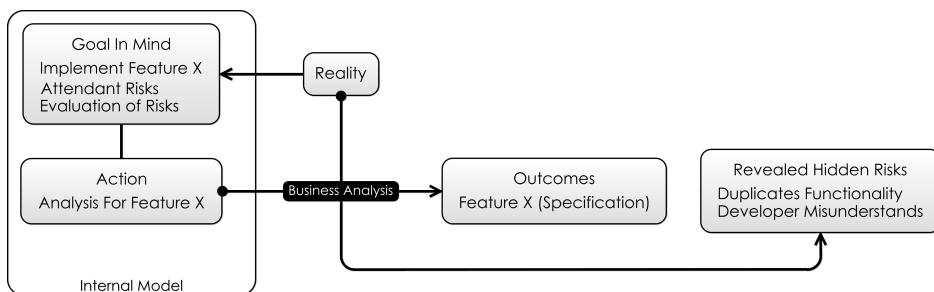


Figure 2.5: BA Specification: exposing hidden risks as soon as possible

In surfacing these risks, there is another outcome: while **Feature X** might be flawed as originally presented, the BA can “evolve” it into a specification, and tie it down sufficiently to reduce the risks. The BA does all this by simply *thinking about it, talking to people and writing stuff down*.

This process of evolving the feature request into a requirement is the BAs job. From our risk-first perspective, it is *taking an idea and making it meet reality*. Not the *full reality* of production (yet), but something more limited. After its brush with reality, the **goal in mind** has *evolved* from being **Feature X (Idea)** to **Feature X (Specification)**.

Code And Unit Test

The next stage for our feature, **Feature X (Specification)** is that it gets coded and some tests get written. Let’s look at how our **goal in mind** meets a new reality: this time it’s the reality of a pre-existing codebase, which has it’s own internal logic.

As the developer begins coding the feature in the software, she will start with an **Internal Model** of the software, and how the code fits into it. But, in the process of implementing it, she is likely to learn about the codebase, and her **Internal Model** will develop.

To a large extent, this is the whole point of *type safety*: to ensure that your **Internal Model** stays consistent with the reality of the codebase. If you add code that doesn't fit the reality of the codebase, you'll know about it with compile errors.

The same thing is true of writing unit tests: again you are testing your **Internal Model** against the reality of the system being built, running in your development environment. Hopefully, this will surface some new hidden risks, and again, because the **goal in mind** has met reality, it is changed, to **Feature X (Code)**.

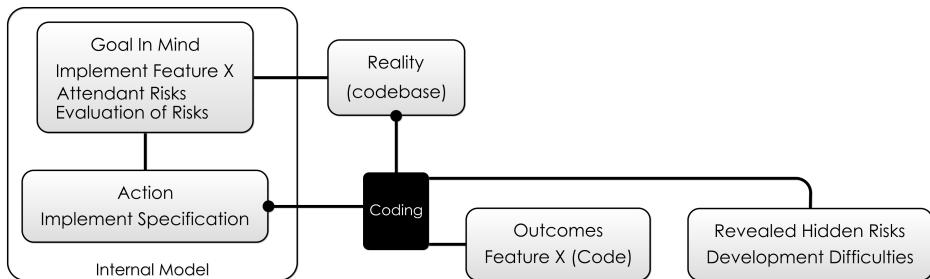


Figure 2.6: Coding Process: exposing more hidden risks as you code

Integration

Integration is where we run *all* the tests on the project, and compile *all* the code in a clean environment: the “reality” of the development environment can vary from one developer’s machine to another.

So, this stage is about the developer’s committed code meeting a new reality: the clean build.

At this stage, we might discover the **Hidden Risk** that we’d break **Feature Y**

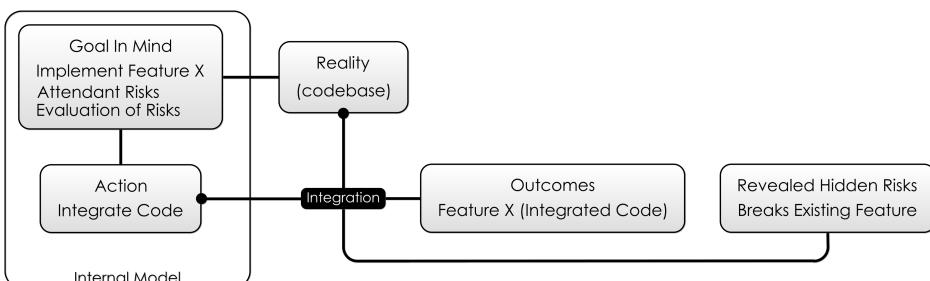


Figure 2.7: Integration testing exposes hidden risks before you get to production

UAT

Is where our feature meets another reality: *actual users*. I think you can see how the process works by now. We're just flushing out yet more **Hidden Risks**:

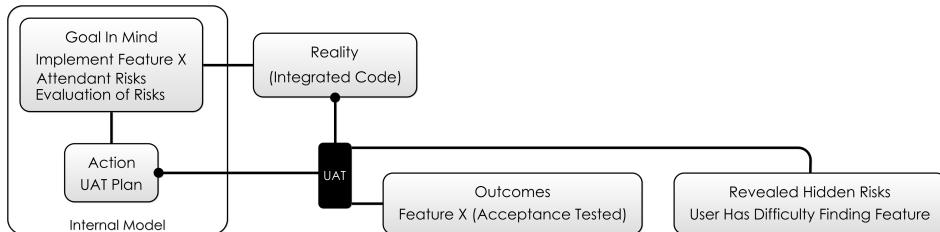


Figure 2.8: UAT - putting tame users in front of your software is better than real ones, where the risk is higher

Observations

A couple of things:

First, the people setting up the development process *didn't know* about these *exact* risks, but they knew the *shape that the risks take*. The process builds “nets” for the different kinds of hidden risks without knowing exactly what they are. Part of the purpose of this site is to help with this and try and provide a taxonomy for different types of risks.

Second, are these really risks, or are they *problems we just didn't know about*? I am using the terms interchangeably, to a certain extent. Even when you know you have a problem, it's still a risk to your deadline until it's solved. So, when does a risk become a problem? Is a problem still just a schedule-risk, or cost-risk? It's pretty hard to draw a line and say exactly.

Third, the real take-away from this is that all these risks exist because we don't know 100% how reality is. Risk exists because we don't (and can't) have a perfect view of the universe and how it'll develop. Reality is reality, *the risks just exist in our head*.

Fourth, hopefully you can see from the above that really *all this work is risk management*, and *all work is testing ideas against reality*.

Conclusion?

Could it be that *everything* you do on a software project is risk management? This is an idea explored in **the next section**.

Chapter 3

All Risk Management

In this section, I am going to introduce the idea that everything you do on a software project is Risk Management.

In the **last section**, we observed that all the activities in a simple methodology had a part to play in exposing different risks. They worked to manage risk prior to them creating bigger problems in production.

Here, we'll look at one of the tools in the Project Manager's toolbox, the RAID Log¹, and observe how risk-centric it is.

RAID Log

Many project managers will be familiar with the RAID Log². It's simply four columns on a spreadsheet:

- Risks
- Actions
- Issues
- Decisions

Let's try and put the following **Attendant Risk** into the RAID Log:

Debbie needs to visit the client to get them to choose the logo to use on the product, otherwise we can't size the screen areas exactly.

- So, is this an **action**? Certainly. There's definitely something for Debbie to do here.
- Is it an **issue**? Yes, because it's holding up the screen-areas sizing thing.
- Is it a **decision**? Well, clearly, it's a decision for someone.
- Is it a **risk**? Probably: Debbie might go to the client and they *still* don't make a decision. What then?

¹<http://pmtips.net/blog-new/raid-logs-introduction>

²<http://pmtips.net/blog-new/raid-logs-introduction>

Let's Go Again

This is a completely made-up example, deliberately chosen to be hard to categorize. Normally, items are more one thing than another. But often, you'll have to make a choice between two categories, if not all four.

This hints at the fact that at some level it's All Risk:

Every Action Mitigates Risk

The reason you are *taking* an action is to mitigate a risk. For example, if you're coing up new features in the software, this is mitigating **Feature Risk**. If you're getting a business sign-off for something, this is mitigating a **Too Many Cooks**-style *stakeholder risk*.

Every Action Carries Risk.

- How do you know if the action will get completed?
- Will it overrun on time?
- Will it lead to yet more actions?

Consider *coding a feature* (as we did in the earlier **Development Process** section). We saw here how the whole process of coding was an exercise in learning what we didn't know about the world, uncovering problems and improving our **Internal Model**. That is, flushing out the **Attendant Risk** of the **Goal In Mind**.

And, as we saw in the **Introduction**, even something *mundane* like the Dinner Party had risks.

An Issue is Just A Type of Risk

- Because issues need to be solved...
- And solving an issue is an action...
- Which, as we just saw also carry risk.

One retort to this might be to say: an issue is a problem I have now, whereas a risk is a problem that *might* occur. I am going to try and *break* that mindset in the coming pages, but I'll just start with this:

- Do you know *exactly* how much damage this issue will do?
- Can you be sure that the issue might not somehow go away?

Issues then, just seem more "definite" and "now" than *risks*, right? This classification is arbitrary: they're all just part of the same spectrum, so stop agonising over which column to put them in.

Every Decision is a Risk.

- By the very nature of having to make a decision, there's the risk you'll decide wrongly.
- And, there's the time it takes to make the decision.
- And what's the risk if the decision doesn't get made?

Failure

tbd.

What To Do?

It makes it much easier to tackle the RAID log if there's only one list: all you do is pick the worst risk on the list, and deal with it. (In **Risk Theory** we look at how to figure out which one that is).

OK, so maybe that *works* for a RAID log (or a Risk log, since we've thrown out the others), but does it scale to a whole project?

In the next section, **Software Project Scenario** I will make a slightly stronger case for the idea that it does.

Chapter 4

Software Project Scenario

Where do the risks of the project lie?

How do we decide what *needs to be done today* on a software project?

Let's look again at the simple risk framework from the **introduction** and try to apply it at the level of the *entire project*.

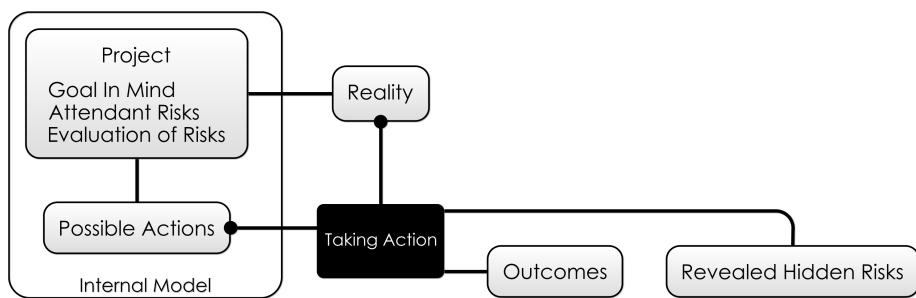


Figure 4.1: Taking action changes reality, but it changes your perception of the attendant risks too

Goal In Mind

How should we decide how to spend our time today?

What actions should we take? (In **Scrum** terminology, what is our *Sprint Goal*?).

If we want to take the right actions, we need to have a good **Internal Model**.

Sometimes, we will know that our model is deficient, and our time should be spent *improving* it, perhaps by talking to our clients, or the support staff, or other developers, or reading.

But let's say for example, today our **Goal In Mind** is to grow our user base.

Attendant Risks

What are the **Attendant Risks** that come with that goal? Here are some to get us started:

1. The users can't access the system
2. The data gets lost, stolen.
3. The data is wrong or corrupted
4. There are bugs that prevent the functionality working
5. The functionality isn't there that the user needs (**Feature Risk**).
6. Our **Internal Model** of the market is poor, and we could be building the wrong thing.

I'm sure you can think of some more.

Evaluating The Risks

Next, we can look at each of these risks and consider the threat they represent. Usually, when **evaluating a risk** we consider both its **impact** and **likelihood**.

The same **Attendant Risks** will be evaluated differently depending on the *nature of the project* and the mitigations you already have in place. For example:

- If they **can't access it**, does that mean that they're stuck unable to get on the train? Or they can't listen to music?
- If the **data is lost**, does this mean that no one can get on the plane? Or that the patients have to have their CAT scans done again? Or that people's private information is scattered around the Internet?
- If the **data is wrong**, does that mean that the wrong people get sent their parcels? Do they receive the wrong orders? Do they end up going to the wrong courses?
- If there are **bugs**, does it mean that their pictures don't end up on the internet? Does it mean that they have to restart the program? Does it mean that they'll waste time, or that they end up thinking they have insurance but haven't?
- If there is **missing functionality**, will they not buy the system? Will they use a competitor's product? Will they waste time doing things a harder or less optimal way?
- If our **Internal Model** is wrong, then is there a chance we are building something for a non-existent market? Or annoying our customers? Or leaving an opportunity for competitors?

Outcomes

As part of evaluating the risks, we can also *predict* the negative outcomes if these risks materialise and we don't take action.

- Losing Revenue
- Legal Culpability
- Losing Users
- Bad Reputation
- etc.

A Single Attendant Risk: Getting Hacked

Let's consider a single risk: that the website gets hacked, and sensitive data is stolen. How we evaluate this risk is going to depend on a number of factors:

- How many users we have
- The importance of the data
- How much revenue will be lost
- Risk of litigation
- etc.

Ashley Maddison

We've seen in the example of hacks on LinkedIn and Ashley Maddison¹ that passwords were not held as hashes in the database. (A practice which experienced developers mainly would see as negligent).

How does our model explain what happened here?

- It's possible that *at the time of implementing the password storage*, hashing was considered, but the evaluation of the risk was low: Perhaps, the risk of not shipping quickly was deemed greater. And so they ignored this concern.
- It's also possible that for the developers in question this was a **Hidden Risk**, and they hadn't even considered it.
- However, as the number of users of the sites increased, the risk increased too, but there was no re-evaluation of the risk otherwise they would have addressed it. This was a costly *failure to update the Internal Model*.

Possible Action

When exposing a service on the Internet, it's now a good idea to *look for trouble*: you should go out and try and improve your **Internal Model**.

Thankfully, this is what sites like OWASP² are for: they *tell you about the Attendant Risks* and further, try to provide some evaluation of them to guide your actions.

¹<https://www.acunetix.com/blog/articles/password-hashing-and-the-ashley-madison-hack/>

²https://www.owasp.org/index.php/Top_10-2017_Top_10

Actions

So, this gives us a guide for one potential action we could take *today*. But on its own, this isn't helpful: we would need to consider this action against the actions we could take to mitigate the other risks. Can we answer this question:

Which actions give us the biggest benefit in terms of mitigating the **Attendant Risks**?

That is, we consider for each possible action:

- The Impact and Likelihood of the **Attendant Risks** it mitigates
- The Cost of the Action

For example, it's worth considering that if we're just starting this project, risks 1-4 are *negligible*, and we're only going to spend time building functionality or improving our understanding of the market. (Which makes sense, right?)

Tacit and Explicit Modelling

As we saw in the example of the **Dinner Party**, creating an internal model is something *we just do*: we have this functionality in our brains already. When we scale this up to a whole project team, we can expect the individuals on the project to continue to do this, but we might also want to consider *explicitly* creating a **risk register for the whole project**.

Whether we do this explicitly or not, we are still individually following this model.

In the next section, we're going to take a quick aside into looking at some **Risk Theory**.

Chapter 5

Risk Theory

Here, I am going to recap on some pre-existing knowledge about risk, generally, in order to set the scene for the next section on **Meeting Reality**.

Risk Registers

In the previous section **Software Project Scenario** we saw how you try to look across the **Attendant Risks** of the project, in order to decide what to do next.

A Risk Register¹ can help with this. From Wikipedia:

A typical risk register contains:

- A risk category to group similar risks
- The risk breakdown structure identification number
- A brief description or name of the risk to make the risk easy to discuss
- The impact (or consequence) if event actually occurs rated on an integer scale
- The probability or likelihood of its occurrence rated on an integer scale
- The Risk Score (or Risk Rating) is the multiplication of Probability and Impact and is often used to rank the risks.
- Common mitigation steps (e.g. within IT projects) are Identify, Analyze, Plan Response, Monitor and Control.

This is Wikipedia's example:

Some points about this description:

This is a Bells-and-Whistles Description

Remember back to the Dinner Party example at the start: the Risk Register happened *entirely in your head*. There is a continuum all the way from "in your head" to Wikipedia's Risk Register

¹https://en.wikipedia.org/wiki/Risk_register

Category	Name	RBS ID	Probability	Impact	Mitigation	Contingency	Risk Score after Mitigation	Action By	Action When
Guests	The guests find the party boring	1.1.	low	medium	Invite crazy friends, provide sufficient liquor	Bring out the karaoke	2		within 2hrs
Guests	Drunken brawl	1.2.	medium	low	Don't invite crazy friends, don't provide too much liquor	Call 911	x		Now
Nature	Rain	2.1.	low	high	Have the party indoors	Move the party indoors	0		10mins
Nature	Fire	2.2.	highest	highest	Start the party with instructions on what to do in the event of fire	Implement the appropriate response plan	1	Everyone	As per plan

Figure 5.1: Wikipedia Risk Register

description. Most of the time, it's going to be in your head, or in discussion with the team, rather than written down.

Most of the value of the **Risk-First** approach is *in conversation*. Later, we'll have an example to show how this can work out.

Probability And Impact

Sometimes, it's better to skip these, and just figure out a Risk Score. This is because if you think about "impact", it implies a definite, discrete event occurring, or not occurring, and asks you then to consider the probability of that occurring.

Risk-First takes a view that risks are a continuous quantity, more like *money* or *water*: by taking an action before delivering a project you might add a degree of **Schedule Risk**, but decrease the **Production Risk** later on by a greater amount.

Graphical Analysis

The Wikipedia page² also includes this wonderful diagram showing you risks of a poorly run barbecue party:

This type of graphic is *helpful* in deciding what to do next, although personally I prefer to graph the overall **Risk Score** against the **Cost of Mitigation**: easily mitigated, but expensive risks can therefore be dealt with first (hopefully).

²https://en.wikipedia.org/wiki/Risk_register

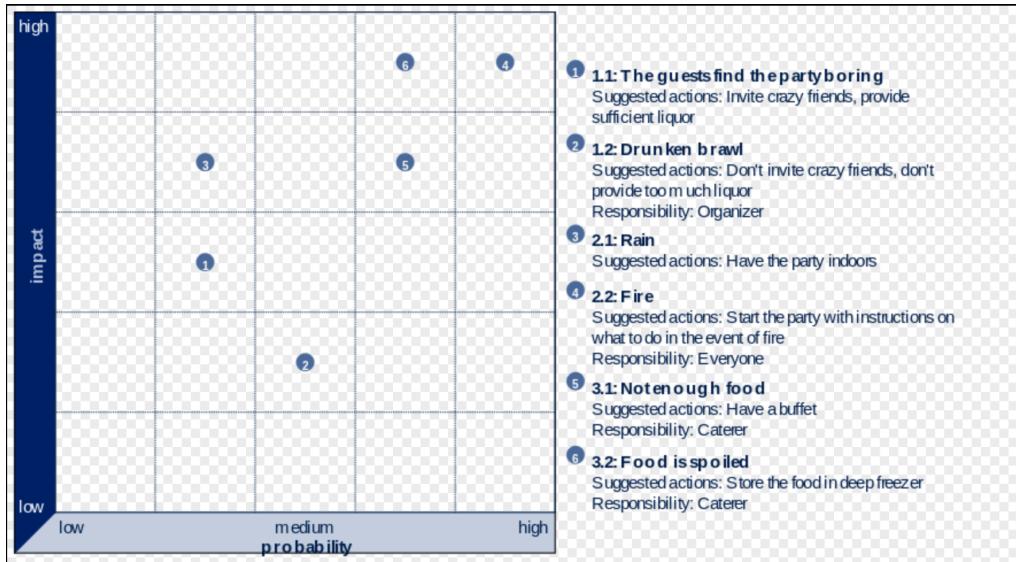


Figure 5.2: Wikipedia Risk Register

Unknown Unknowns

In Wikipedia's example, this fictitious BBQ has high fire risk, so one should begin mitigating there.

But, does this feel right? One of the criticisms of the **Risk Register** approach is that of **mistaking the map for the territory**. That is, mistakenly believing that what's on the Risk Register is *all there is*.

In the preceding discussions, I have been careful to point out the existence of **Hidden Risks** for that very reason. Or, to put another way:

What we don't know is what usually gets us killed - Petyr Baelish

Donald Rumsfeld's famous Known Knowns³ is also a helpful conceptualization.

Risk And Uncertainty

Arguably, this site uses the term 'Risk' wrongly: most literature suggests risk can be measured⁴ whereas uncertainty represents things that cannot.

I am using **risk** everywhere because later we will talk about specific risks (e.g. **Boundary Risk** or **Complexity Risk**), and it doesn't feel grammatically correct to talk about those as **uncertainties**, especially given the pre-existing usage in Banking of terms like **Operational**

³https://en.wikipedia.org/wiki/There_are_known_knowns

⁴<https://keydifferences.com/difference-between-risk-and-uncertainty.html>

Risk⁵ or Reputational risk⁶ which are also not really a-priori measurable.

The Opposite Of Risk Management

Let's look at the classic description of Risk Management:

Risk Management is the process of thinking out corrective actions before a problem occurs, while it's still an abstraction.

The opposite of risk management is crisis management, trying to figure out what to do about the problem after it happens. - Waltzing With Bears, Tom De Marco & Tim Lister

This is not how **Risk-First** sees it:

First, we have the notion that Risks are discrete events, again. Some risks *are* (like gambling on a horse race), but most *aren't*. In the **Dinner Party**, for example, bad preparation is going to mean a *worse* time for everyone, but how good a time you're having is a spectrum, it doesn't divide neatly into just "good" or "bad".

Second, the opposite of "Risk Management" (or trying to minimize the "Downside") is either "Upside Risk Management", (trying to maximise the good things happening), or it's trying to make as many bad things happen as possible. Humans tend to be optimists (especially when there are lots of **Hidden Risks**), hence our focus on Downside Risk. Sometimes though, it's good to stand back and look at a scenario and think: am I capturing all the Upside Risk here?

Finally, Crisis Management is *still just Risk Management*: the crisis (Earthquake, whatever) has *happened*. You can't manage it because it's in the past. All you can do is Risk Manage the future (minimize further casualties and human suffering, for example).

Yes, it's fine to say "we're in crisis", but to assume there is a different strategy for dealing with it is a mistake: this is the Fallacy of Sunk Costs⁷.

Invariances #1: Panic Invariance

You would expect then, that any methods for managing software delivery should be *invariant* to the level of crisis in the project. If, for example, a project proceeds using **Scrum** for eight months, and then the deadline looms and everyone agrees to throw Scrum out of the window and start hacking, then *this implies there is a problem with Scrum*. Or at least, the way it was being implemented.

I call this **Panic Invariance**: the methodology shouldn't need to change given the amount of pressure or importance on the table.

⁵https://en.wikipedia.org/wiki/Operational_risk

⁶<https://www.investopedia.com/terms/r/reputational-risk.asp>

⁷https://en.wikipedia.org/wiki/Sunk_costs

Invariances #2: Scale Invariance

Another test of a methodology is that it shouldn't fall down when applied at different *scales*. Because, if it does, this implies that there is something wrong with the methodology. The same is true of physical laws: if they don't apply under all circumstances, then that implies something is wrong. For example, Newton's Laws of Motion fail to calculate the orbital period of Mercury, and this was an early win for Einstein's Relativity.

Some methodologies are designed for certain scales: Extreme Programming is designed for small, co-located teams. And, that's useful. But the fact it doesn't scale tells us something about it: chiefly, that it considers certain *kinds* of risk, while ignoring others. At small scales, that works ok, but at larger scales, the bigger risks increase too fast for it to work.

tbd.

So ideally, a methodology should be applicable at *any* scale:

- A single class or function
- A collection of functions, or a library
- A project team
- A department
- An entire organisation

If the methodology *fails at a particular scale*, this tells you something about the risks that the methodology isn't addressing. It's fine to have methodologies that work at different scales, and on different problems. One of the things that I am exploring with Risk First is trying to place methodologies and practices within a framework to say *when* they are applicable.

Value

“Upside Risk” isn't a commonly used term: industry tends to prefer “value”, as in “Is this a value-add project?”. There is plenty of theory surrounding **Value**, such as Porter’s **Value Chain** and **Net Present Value**. This is all fine so long as we remember:

- **The pay-off is risky:** Since the **Value** is created in the future, we can't be certain about it happening - we should never consider it a done-deal. **Future Value** is always at risk. In finance, for example, we account for this in our future cash-flows by discounting them according to the risk of default.
- **The pay-off amount is risky:** Additionally, whereas in a financial transaction (like a loan, say), we might know the size of a future payment, in IT projects we can rarely be sure that they will deliver a certain return. On some fixed-contract projects this sometimes is not true: there may be a date when the payment-for-delivery gets made, but mostly we'll be expecting an uncertain pay-off.

Risk-First is a particular view on reality. It's not the only one. However, I am going to try and make the case that it's an underutilized one that has much to offer us.

Speed

For example, in **Rapid Development** by Steve McConnell we have the following diagram:

tbd. redraw this.

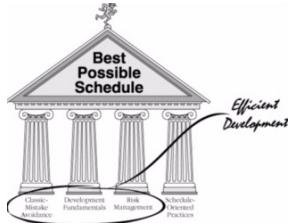


Figure 5.3: Rapid Development Pillars- From Steve McConnell

And, this is *fine*, McConnel is structuring the process from the perspective of *delivering as quickly as possible*. However, here, I want to turn this on its head. Exploring Software Development from a risk-first perspective is an under-explored technique, and I believe it offers some useful insights. So the aim here is to present the case for viewing software development like this:

tbd. risk-first diagram.

Net Present Risk

If we can view software delivery from the point of view of *value*, then why can't we apply the same tools to **Risk** too? In order to do this, let's review "Eisenhower's Box" model. This considers two variables:

- How valuable the work is (Importance)
- How soon it is needed (Urgency)

tbd. image from wikipedia. text from wikipedia.

Here, we're considering a synthesis of both *time* and *value*. But **Net Present Value** allows us to discount value in the future, which offers us a way to reconcile these two variables:

chart of discounting into the future tbd.

Let's do the same thing with risk? Let's introduce the concept of **Net Present Risk**, or NPR:

Net Present Risk is tbd.

Let's look at a quick example to see how this could work out. Let's say you had the following 3 risks:

- Risk A, which will cost you £50 in 5 year's time.
- Risk B, which will cost you £70 in 8 year's time.
- Risk C, which will cost you £120 in 18 year's time.

Which has the biggest NPR? Well, it depends on the discount rate that you apply. Let's assume we are discounting at 6% per year. A graph of the discounted risks looks like this:

tbd, see numbers

On this basis, the biggest risk is B, at about #45. If we *reduce* the discount factor to 3%, we get a different result:

tbd, see numbers.

Now, risk **C** is bigger.

Because this is *Net Present Risk*, we can also use it to make decisions about whether or not to mitigate risks. Let's consider the cost of mitigating each risk *now*:

- Risk **A** costs £20 to mitigate
- Risk **B** costs £50 to mitigate
- Risk **C** costs £100 to mitigate

Which is the best deal?

Well, under the 6% regime, only Risk **A** is worth mitigating, because you spend £20 today to get rid of #40 of risk (today). The NPR is positive at around £20, whereas for **B** and **C** mitigations it's under water.

tbd.

Under a 3% regime, risk **A** and **B** are *both* worth mitigating, as you can see in this graph:

Discounting the Future To Zero

I have worked in teams sometimes where the blinkers go down, and the only thing that matters is *now*. They may apply a rate of 60% per-day, which means that anything with a horizon over a week is irrelevant. Regimes of such **hyperinflation** are a sure sign that something has *really broken down* within a project. Consider in this case a Discount Factor of 60% per day, and the following risks:

- Risk A: £80 cost, happening *tomorrow*
- Risk B: £500 cost, happening in *5 days*.

Risk B is almost irrelevant under this regime, as this graph shows:

tbd.

Why do things like this happen? Often, the people involved are under incredible job-stress: usually they are threatened with the sack on a daily basis, and therefore feel they have to react. For publically-listed companies you can also

- more pressure, heavier discounting pooh bear procrastination

Is This Scientific?

Risk-First is an attempt to provide a practical framework, rather than a scientifically rigorous analysis. In fact, my view is that you should *give up* on trying to compute risk numerically. You *can't* work out how long a software project will take based purely on an analysis of (say) *function points*. (Whatever you define them to be).

- First, there isn't enough evidence for an approach like this. We *can* look at collected data about IT projects, but **techniques and tools change**.

- Second, IT projects have too many confounding factors, such as experience of the teams, technologies used etc. That is, the risks faced by IT projects are *too diverse* and *hard to quantify* to allow for meaningful comparison from one to the next.
- Third, as soon as you *publish a date* it changes the expectations of the project (see **Student Syndrome**).
- Fourth, metrics get first of all **misused** and then **gamed**.

Reality is messy. Dressing it up with numbers doesn't change that and you risk **fooling yourself**. If this is the case, is there any hope at all in what we're doing? I would argue yes: *forget precision*. You should, with experience be able to hold up two separate risks and answer the question, "is this one bigger than this one?"

Reality is Reality, **so let's meet it**.

Chapter 6

Meeting Reality

In this section, we will look at how exposing your **Internal Model** to reality is in itself a good risk management technique.

Revisiting the Model

In **A Simple Scenario**, we looked at a basic model for how **Reality** and our **Internal Model** interacted with each other: we take action based on our **Internal Model**, hoping to **change Reality** with some positive outcome.

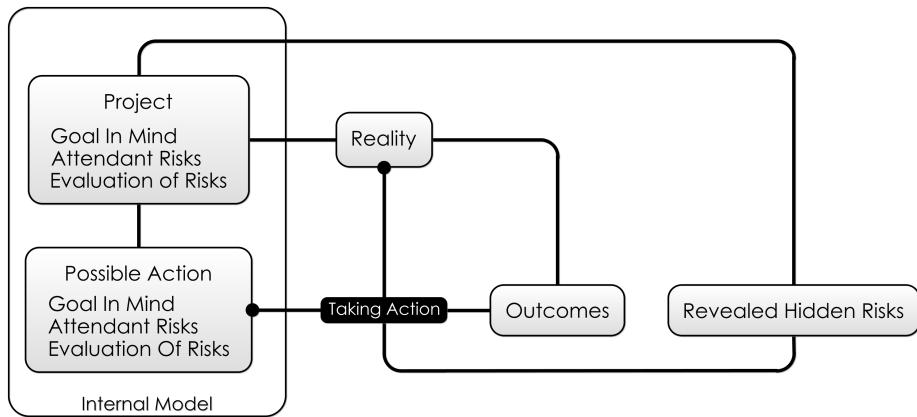
And, in **Development Process** we looked at how we can meet with reality in *different forms*: Analysis, Testing, Integration and so on, and saw how the model could work in each stage of a project.

Finally, in **Software Project Scenario** we looked at how we could use this model on a day-to-day basis to inform what we should do next.

So, it should be no surprise to see that there is a *recursive* nature about this:

1. The **actions we take** each day have consequences: they **expose new Hidden Risks******, which inform our **Internal Model**, and at the same time, they change reality in some way (otherwise, what would be the point of doing them?)
2. The actions we take towards achieving a **Goal In Mind** each have their *own Goal In Mind*. And because of this, when we take action, we have to consider and evaluate the **Hidden Risks** exposed by that action. That is, there are many ways to achieving a goal, and these different ways expose different **Hidden Risks**.

So, let's see how this kind of recursion looks on our model. Note that here, I am showing *just one possible action*, in reality, you'll have choices.



Hopefully, if you've read along so far, this model shouldn't be too hard to understand. But, how is it helpful?

“Navigating the Risk Landscape”

So, we often have multiple ways of achieving a **Goal In Mind**.

What's the best way?

I would argue that the best way is the one which accrues the *least risk* to get it done: each action you take in trying to achieve the overall **Goal In Mind** will have its **Attendant Risks**, and it's the experience you bring to bear on these that will help you navigate through them smoothly.

Ideally, when you take an action, you are trading off a big risk for a smaller one. Take Unit Testing for example. Clearly, writing Unit Tests adds to the amount of development work, so on its own, it adds **Schedule Risk**. However, if you write *just enough* of the right Unit Tests, you should be short-cutting the time spent finding issues in the User Acceptance Testing (UAT) stage, so you're hopefully trading off a larger **Schedule Risk** from UAT and adding a smaller risk to **Development**.

Sometimes, in solving one problem, you can end up somewhere worse: the actions you take to solve a higher-level **Attendant Risk** will leave you with a worse **Attendant Risks**. Almost certainly, this will have been a **Hidden Risk** when you embarked on the action, otherwise you'd not have chosen it.

An Example: Automation

diagram of how automation reduces process risk, but increases complexity?

Another Quick Example: MongoDB

On a recent project in a bank, we had a requirement to store a modest amount of data and we needed to be able to retrieve it fast. The developer chose to use MongoDB¹ for this. At the time, others pointed out that other teams in the bank had had lots of difficulty deploying MongoDB internally, due to licensing issues and other factors internal to the bank.

Other options were available, but the developer chose MongoDB because of their *existing familiarity* with it: therefore, they felt that the **Hidden Risks** of MongoDB were *lower* than the other options, and disregarded the others' opinions.

The data storage **Attendant Risk** was mitigated easily with MongoDB. However, the new **Attendant Risk** of licensing bureaucracy eventually proved too great, and MongoDB had to be abandoned after much investment of time.

This is not a criticism of MongoDB: it's simply a demonstration that sometimes, the cure is worse than the disease. Successful projects are *always* trying to *reduce Attendant Risks*.

The Cost Of Meeting Reality

Meeting reality is *costly*, for example. Going to production can look like this:

- Releasing software
- Training users
- Getting users to use your system
- Gathering feedback

All of these steps take a lot of effort and time. But you don't have to meet the whole of reality in one go - sometimes that is expensive. But we can meet it in "limited ways".

In all, to de-risk, you should try and meet reality:

- **Sooner**, so you have time to mitigate the hidden risks it uncovers
- **More Frequently**: so the hidden risks don't hit you all at once
- **In Smaller Chunks**: so you're not overburdened by hidden risks all in one go.
- **With Feedback**: if you don't collect feedback from the experience of meeting reality, hidden risks *stay hidden*.

tbd this is what testing is all about.

YAGNI

As a flavour of what's to come, let's look at YAGNI², an acronym for You Aren't Gonna Need It. Martin Fowler says:

Yagni originally is an acronym that stands for "You Aren't Gonna Need It". It is a mantra from ExtremeProgramming that's often used generally in agile software

¹<https://www.mongodb.com>

²<https://www.martinfowler.com/bliki/Yagni.html>

teams. It's a statement that some capability we presume our software needs in the future should not be built now because "you aren't gonna need it".

This principle was first discussed and fleshed out on Ward's Wiki³

The idea makes sense: if you take on extra work that you don't need, *of course* you'll be accreting **Attendant Risks**.

But, there is always the opposite opinion: You Are Gonna Need It⁴. As a simple example, we often add log statements in our code as we write it, though following YAGNI strictly says we should leave it out.

Which is right?

Now, we can say: do the work *if it mitigates your Attendant Risks*.

- Logging statements are *good*, because otherwise, you're increasing the risk that in production, no one will be able to understand *how the software went wrong*.
- However, adding them takes time, which might introduce **Schedule Risk**.

So, it's a trade-off: continue adding logging statements so long as you feel that overall, you're reducing risk.

Do The Simplest Thing That Could Possibly Work

Another mantra from Kent Beck (originator of the **Extreme Programming** methodology, is "Do The Simplest Thing That Could Possibly Work", which is closely related to YAGNI and is about looking for solutions which are simple. Our risk-centric view of this strategy would be:

- Every action you take on a project has its own **Attendant Risks**.
- The bigger or more complex the action, the more **Attendant Risk** it'll have.
- The reason you're taking action *at all* is because you're trying to reduce risk elsewhere on the project
- Therefore, the biggest payoff is whatever action *works* to remove that risk, whilst simultaneously picking up the least amount of new **Attendant Risk**.

So, "Do The Simplest Thing That Could Possibly Work" is really a helpful guideline for Navigating the **Risk Landscape**.

Summary

So, here we've looked at Meeting Reality, which basically boils down to taking actions to manage risk and seeing how it turns out:

- Each Action you take is a step on the Risk Landscape
- Each Action is a cycle around our model.
- Each cycle, you'll expose new **Hidden Risks**, changing your **Internal Model**.

³<http://wiki.c2.com/?YouArentGonnaNeedIt>

⁴<http://wiki.c2.com/?YouAreGonnaNeedIt>

- Preferably, each cycle should reduce the overall **Attendant Risk** of the Goal
- Surely, the faster you can do this, the better? **Let's investigate...**

Chapter 7

Cadence

Let's go back to the model again, introduced in **Meeting Reality**:

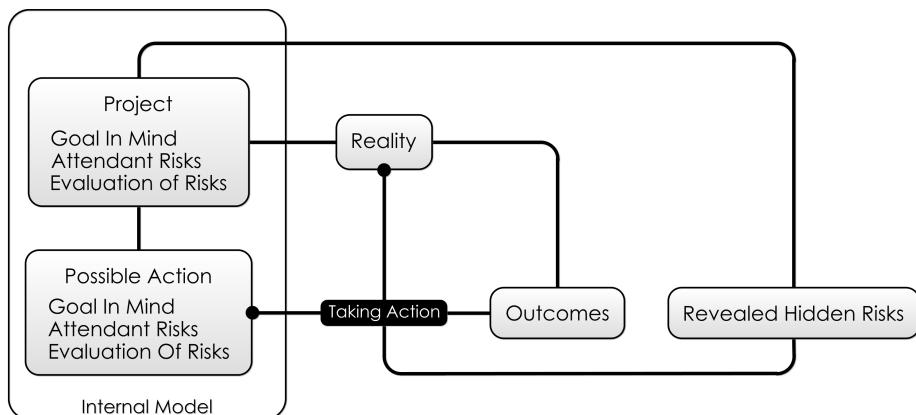


Figure 7.1: Meeting Reality: reality is changed and so is your internal model.

As you can see, it's an idealized **Feedback Loop**.

How *fast* should we go round this loop? Is there a right answer? The longer you leave your **goal in mind**, the longer it'll be before you find out how it really stacks up against reality.

Testing your **goals in mind** against reality early and safely is how you'll manage risk effectively, and to do this, you need to set up **Feedback Loops**. e.g.

- **Bug Reports and Feature Requests** tell you how the users are getting on with the software.
- **Monitoring Tools and Logs** allow you to find out how your software is doing in reality.
- **Dog-fooding** i.e using the software you write yourself might be faster than talking to users.

- **Continuous Delivery** (CD) is about putting software into production as soon as it's written.
- **Integration Testing** is a faster way of meeting *some* reality than continually deploying code and re-testing it manually.
- **Unit Testing** is a faster feedback loop than Integration Testing.
- **Compilation** warns you about logical inconsistencies in your code.

.. and so on.

Time / Reality Trade-Off

This list is arranged so that at the top, we have the most visceral, most *real* feedback loop, but at the same time, the slowest.

At the bottom, a good IDE can inform you about errors in your **Internal Model** in real time, by way of highlighting compilation errors . So, this is the fastest loop, but it's the most *limited* reality.

Imagine for a second that you had a special time-travelling machine. With it, you could make a change to your software, and get back a report from the future listing out all the issues people had faced using it over its lifetime, instantly.

That'd be neat, eh? If you did have this, would there be any point at all in a compiler? Probably not, right?

The whole *reason* we have tools like compilers is because they give us a short-cut way to get some limited experience of reality *faster* than would otherwise be possible. Because, cadence is really important: the faster we test our ideas, the more quickly we'll find out if they're correct or not.

Development Cycle Time

One thing that often astounds me is how developers can ignore the fast feedback loops at the bottom of the list, because the ones nearer the top *will do*. In the worst cases, changing two lines of code, running the build script, deploying and then manually testing out a feature. And then repeating.

If you're doing it over and over, this is a terrible waste of time. And, you get none of the benefit of a permanent suite of tests to run again in the future.

The Testing Pyramid¹ hints at this truth:

- **Unit Tests** have a *fast feedback loop*, so have *lots of them*.
- **Integration Tests** have a slightly *slower feedback loop*, so have *few of them*. Use them when you can't write unit tests (at the application boundaries).
- **Manual Tests** have a *very slow feedback loop*, so have *even fewer of them*. Use them as a last resort.

¹<http://www.agilenutshell.com/episodes/41-testing-pyramid>

Production

You could take this section to mean that **Continuous Delivery** (CD) is always and everywhere a good idea. I guess that's not a bad take-away, but it's clearly more nuanced than that.

Yes, CD will give you faster feedback loops, but getting things into production is not the whole story: the feedback loop isn't complete until people have used the code, and reported back to the development team.

The right answer is to use the fastest feedback loop possible, *which actually does give you feed back.*

Recap

Let's look at the journey so far:

- In **A Simple Scenario** we looked at how risk pervades every goal we have in life, big or small. We saw that risk stems from the fact that our **Internal Model** of the world couldn't capture everything about reality, and so some things were down to chance.
- In the **Development Process** we looked at how common software engineering conventions like Unit Testing, User Acceptance Testing and Integration could help us manage the risk of taking an idea to production, by *gradually* introducing it to reality in stages.
- In **It's All Risk Management** we took a leap of faith: Could *everything* we do just be risk management? And we looked at the RAID log and thought that maybe it could be.
- Next, in **A Software Project Scenario** we looked at how you could treat the project-as-a-whole as a risk management exercise, and treat the goals from one day to the next as activities to mitigate risk.
- **Some Risk Theory** was an aside, looking at some terminology and the useful concept of a Risk Register.
- Then, generalizing the lessons of the Development Process article, we examined the idea that **Meeting Reality** frequently helps flush out **Hidden Risks** and improve your **Internal Model**.
- Finally, above, we looked at **Cadence**, and how feedback loops allow you Navigate the Risk Landscape more effectively, by showing you more quickly when you're going wrong.

What this has been building towards is supplying us with a vocabulary with which to communicate to our team-mates about which Risks are important to us, which actions we believe are the right ones, and which tools we should use.

Let's have a **look at an example** of how this might work:

Chapter 8

A Conversation

After so much theory, it seems like it's time to look at how we can apply these principles in the real world.

The following is based the summary of an issue from just a few weeks ago. It's heavily edited and anonymized, and I've tried to add the **Risk-First** vocabulary along the way, but otherwise, it's real.

Some background: **Synergy** is an online service with an app-store, and **Eve** and **Bob** are developers working for **Large Corporation LTD**, which wants to have an application accepted into Synergy's app-store.

Synergy's release means that the app-store refresh will happen in a few weeks, so this is something of a hard deadline: if we miss it, the next release will be four months away.

A Risk Conversation

Eve: We've got a problem with the Synergy security review.

Bob: Tell me.

Eve: Well, you know Synergy did their review and asked us to upgrade our Web Server to only allow TLS version 1.1 and greater?

Bob: Yes, I remember: We discussed it as a team and thought the simplest thing would be to change the security settings on the Web Server, but we all felt it was pretty risky. We decided that in order to flush out **Hidden Risk**, we'd upgrade our entire production site to use it *now*, rather than wait for the app launch.

Eve: Right, and it *did* flush out **Hidden Risk**: some people using Windows 7, downloading Excel spreadsheets on the site, couldn't download them: for some reason, that combination didn't support anything greater than TLS version 1.0. So, we had to back it out.

Bob: Ok, well I guess it's good we found out *now*. It would have been a disaster to discover this after the go-live.

Eve: Yes. So, what's our next-best action to mitigate this?

Bob: Well, we could go back to Synergy and ask them for a reprieve, but I think it'd be better to mitigate this risk now if we can... they'll definitely want it changed at some point.

Eve: How about we run two web-servers? One for the existing content, and one for our new Synergy app? We'd have to get a new external IP address, handle DNS setup, change the firewalls, and then deploy a new version of the Web Server software on the production boxes.

Bob: This feels like there'd be a lot of **Attendant Risk**: and all of this needs to be handled by the Networking Team, so we're picking up a lot of **Bureaucracy Risk**. I'm also worried that there are too many steps here, and we're going to discover loads of **Hidden Risks** as we go.

Eve: Well, you're correct on the first one. But, I've done this before not that long ago for a Chinese project, so I know the process - we shouldn't run into any new **Hidden Risk**.

Bob: Ok, fair enough. But isn't there something simpler we can do? Maybe some settings in the Web Server?

Eve: Well, if we were using Apache, yes, it would be easy to do this. But, we're using Baroque Web Server, and it *might* support it, but the documentation isn't very clear.

Bob: Ok, and upgrading it is a *big* risk, right? We'd have to migrate all of our **configuration**...

Eve: Yes, let's not go there. But if we changing the settings on Baroque, we have the **Attendant Risk** that it's not supported by the software and we're back where we started. Also, if we isolate the Synergy app stuff now, we can mess around with it at any point in future, which is a big win in case there are other **Hidden Risks** with the security changes that we don't know about yet.

Bob: Ok, I can see that buys us something, but time is really short and we have holidays coming up.

Eve: Yes. How about for now, we go with the isolated server, and review next week? If it's working out, then great, we continue with it. Otherwise, if we're not making progress next week, then it'll be because our isolation solution is meeting more risk than we originally thought. We can try the settings change in that case.

Bob: Fair enough, it sounds like we're managing the risk properly, and because we can hand off a lot of this to the Networking Team, we can get on with mitigating our biggest risk on the project, the authentication problem, in the meantime.

Eve: Right. I'll check in with the Networking Team each day and make sure it doesn't get forgotten.

Aftermath

Hopefully, this type of conversation will feel familiar. It should. There's nothing ground-breaking at all in what we've covered so far; it's more-or-less just Risk Management theory.

If you can now apply it in conversation, like we did above, then that's one extra tool you have for delivering software.

So with the groundwork out of the way, let's get on to Part 2 and investigate **The Risk Landscape**.

Chapter 9

De Risking

Describe the basic process here.

Avoiding, evading, luck etc.

Part II

Risk

Chapter 10

Risk Landscape

Risk is messy. It's not always easy to tease apart the different components of risk and look at them individually. Let's look at a high-profile recent example to see why.

The Financial Crisis

In the Financial Services¹ industry, whole *departments* exist to calculate things like:

- Market Risk²: the risk that the amount some asset you hold/borrow/have loaned is going to change in value.
- Credit Risk³: the risk that someone who owes you a payment at a specific point in time might not pay it back.
- Liquidity Risk⁴: the risk that you can't find a market to sell/buy something, usually leading to a shortage of ready cash.

... and so on. But, we don't need to know the details exactly to understand this story.

They get expressed in ways like this:

“we have a 95% chance that today we'll lose less than £100”

In the financial crisis, though, these models of risk didn't turn out to be much use. Although there are lots of conflicting explanations of what happened, one way to look at it is this:

- Liquidity difficulties (i.e. amount of cash you have for day-to-day running of the bank) caused some banks to not be able to cover their interest payments.
- This caused credit defaults (the thing that Credit Risk⁵ measures were meant to guard against) even though the banks *technically* were solvent.

¹https://en.wikipedia.org/wiki/Financial_services

²https://en.wikipedia.org/wiki/Market_risk

³https://en.wikipedia.org/wiki/Credit_risk

⁴https://en.wikipedia.org/wiki/Liquidity_risk

⁵https://en.wikipedia.org/wiki/Credit_risk

- That meant that, in time, banks got bailed out, share prices crashed and there was lots of Quantitative Easing⁶.
- All of which had massive impacts on the markets in ways that none of the Market Risk⁷ models foresaw.

All the **Risks** were correlated⁸. That is, they were affected by the *same underlying events, or each other*.

The Risk Landscape Again

It's like this with software risks, too, sadly.

In **Meeting Reality**, we looked at the concept of the **Risk Landscape**, and how a software project tries to *navigate* across this landscape, testing the way as it goes, and trying to get to a position of *more favourable risk*.

It's tempting to think of our **Risk Landscape** as being like a Fitness Landscape⁹. That is, you have a "cost function" which is your height above the landscape, and you try and optimise by moving downhill in a Gradient Descent¹⁰ fashion.

However, there's a problem with this: As we said in **Risk Theory**, we don't have a cost function. We can only guess at what risks there are. And, we have to go on our *experience*. For this reason, I prefer to think of the **Risk Landscape** as a terrain which contains *fauna* and *obstacles* (or, specifically **Boundaries**).

I am going to try and show you some of the fauna of the **Risk Landscape**. We know every project is different, so every **Risk Landscape** is also different. But, just as I can tell you that the landscape outside your window will probably will have some roads, trees, fields, forests, buildings, and that the buildings are likely to be joined together by roads, I can tell you some general things about risks too.

In fact, we're going to try and categorize the kinds of things we see on this **Risk Landscape**. But, this isn't going to be perfect:

- One risk can "blend" into another just like sometimes a "field" is also a "car-park" or a building might contain some trees (but isn't a forest).
- There is *correlation* between different risks: one risk may cause another, or two risks may be due to the same underlying cause.
- As we saw in **Part 1**, mitigating one risk can give rise to another, so risks are often *inversely correlated*.

Why Should We Categorize The Risks?

This is a "spotters' guide" to risks, not an in-depth encyclopedia.

⁶https://en.wikipedia.org/wiki/Quantitative_easing

⁷https://en.wikipedia.org/wiki/Market_risk

⁸<https://www.investopedia.com/terms/c/correlation.asp>

⁹https://en.wikipedia.org/wiki/Fitness_landscape

¹⁰https://en.wikipedia.org/wiki/Gradient_descent

If we were studying insects, this might be a guide giving you a description and a picture of each insect, telling you where to find it and what it does. That doesn't mean that this is *all* there is to know. Just as a scientist could spend her entire life studying a particular species of bee, each of the risks we'll look at really has a whole sub-discipline of Computer Science attached to it, which we can't possibly hope to cover all of.

As software developers, we can't hope to know the detailed specifics of the whole discipline of Complexity Theory¹¹, or Concurrency Theory¹². But, we're still required to operate in a world where these things exist. So, we may as well get used to them, and ensure that we respect their primacy. We are operating in *their* world, so we need to know the rules.

We're all Naturalists Now

This is a new adventure. There's a long way to go. Just as naturalists are able to head out and find new species of insects and plants, we should expect to do the same. This is by no means a complete picture - it's barely a sketch.

It's a big, crazy, evolving world of software. Help to fill in the details. Report back what you find.

Our Tour Itinerary

Below is a table outlining the different risks we'll see. There *is* an order to this: the later risks are written assuming a familiarity with the earlier ones. Hopefully, you'll stay to the end and see everything, but you're free to choose your own tour if you want to.

Risk	Description
Feature Risk	When you haven't built features the market needs, or they contain bugs, or the market changes underneath you.
Complexity Risk	Your software is so complex it makes it hard to change, understand or run.
Communication Risk	Risks associated with getting messages heard and understood.
Dependency Risk	Risks of depending on other people, products, software, functions, etc. This is a general look at dependencies, before diving into specifics like...
Schedule Risk	Risks associated with having a dependency on time (or money).
Software Dependency Risk	When you choose to depend on a software library, service or function.
Process Risk	When you depend on a business process, or human process to give you something you need.

¹¹https://en.wikipedia.org/wiki/Complexity_theory

¹²[https://en.wikipedia.org/wiki/Concurrency_\(computer_science\)](https://en.wikipedia.org/wiki/Concurrency_(computer_science))

Risk	Description
Boundary Risk	Risks due to making decisions that limit your choices later on. Sometimes, you go the wrong way on the Risk Landscape and it's hard to get back to where you want to be.
Agency Risk	Risks that staff have their own Goals , which might not align with those of the project or team.
Coordination Risk	Risks due to the fact that systems contain multiple agents, which need to work together.
Map And Territory Risk	Risks due to the fact that people don't see the world as it really is. (After all, they're working off different, imperfect Internal Models .)
Operational Risk	Software is embedded in a system containing people, buildings, machines and other services. Operational risk considers this wider picture of risk associated with running a software service or business in the real world.

On each page we'll start by looking at the category of the risk *in general*, and then break this down into some specific subtypes. At the end, in **Staging and Classifying** we'll have a recap about what we've seen and make some guesses about how things fit together.

So, let's get started with **Feature Risk**.

Chapter 11

Feature Risk

Feature Risk is the category of software risk to do with features that have to be in your software. It is the risk that you face by *not having features that your clients need*.

In a way, **Feature Risk** is very fundamental: if there were *no* feature risk, the job would be done already, either by you, or by another product, and the product would be perfect!

As a simple example, if your needs are served perfectly by Microsoft Excel, then you don't have any **Feature Risk**. However, the day you find Microsoft Excel wanting, and decide to build an Add-On is the day when you first appreciate some **Feature Risk**.

Not considering **Feature Risk** means that you might be building the wrong functionality, for the wrong audience or at the wrong time. And eventually, this will come down to lost money, business, acclaim, or whatever else reason you are doing your project for. So let's unpack this concept into some of its variations.

Variations

Feature Fit Risk

This is the one we've just discussed above: the feature that you (or your clients) want to use in the software *isn't there*. Now, as usual, you could call this an issue, but we're calling it a **Risk** because it's not clear exactly *how many* people are affected, or how badly.

- This might manifest itself as complete *absence* of something you need, e.g “Where is the word count?”
- It could be that the implementation isn't complete enough, e.g “why can't I add really long numbers in this calculator?”

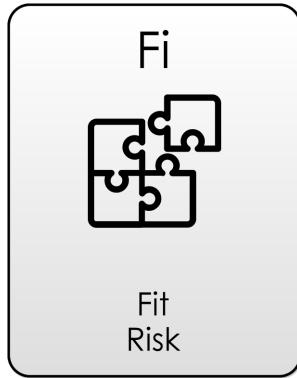
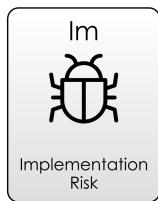


Figure 11.1: Feature Risk

Features Don't Work Properly

Feature Risk also includes things that don't work as expected: That is to say, bugs¹. Although the distinction between “a missing feature” and “a broken feature” might be worth making in the development team, we can consider these both the same kind of risk: *the software doesn't do what the user expects*.



- Risk that the functionality you are providing doesn't match the features the client is expecting, due to poor or confusing implementation.

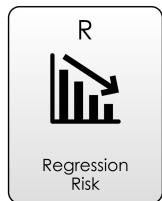
Figure 11.2: Implementation Risk

(At this point, it's worth pointing out that sometimes, *the user expects the wrong thing*. This is a different but related risk, which could be down to **Training** or **Documentation** or simply **Poor User Interface** and we'll look at that more in **Communication Risk**.)

Regression Risk

Regression Risk is basically risk of breaking existing features in your software when you add new ones. As with the previous risks, the eventual result is the same; customers don't have the features they expect. This can become a problem as your code-base **gains Complexity**, as it becomes impossible to keep a complete **Internal Model** of the whole thing.

¹https://en.wikipedia.org/wiki/Software_bug



- Risk that the functionality you provide changes for the worse, over time.

Figure 11.3: Regression Risk

Also, while delivering new features can delight your customers, breaking existing ones will annoy them. This is something we'll come back to in **Reputation Risk**.

Conceptual Integrity Risk



- Risk that the software you provide is too complex, or doesn't match the expectations of your clients' internal models.

Figure 11.4: Conceptual Integrity Risk

Sometimes, users *swear blind* that they need some feature or other, but it runs at odds with the design of the system, and plain *doesn't make sense*. Often, the development team can spot this kind of conceptual failure as soon as it enters the **Backlog**. Usually, it's in coding that this becomes apparent.

tbd: feature phones.

Sometimes, it can go for a lot longer. I once worked on some software that was built as a score-board within a chat application. However, after we'd added much-asked-for commenting and reply features to our score-board, we realised we'd implemented a chat application *within a chat application*, and had wasted our time enormously.

Which leads to Greenspun's 10th Rule:

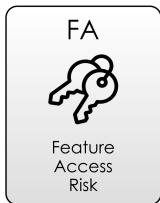
"Any sufficiently complicated C or Fortran program contains an ad-hoc, informally-specified, bug-ridden, slow implementation of half of Common Lisp." - Greenspun's 10th Rule, *Wikipedia*²

²https://en.wikipedia.org/wiki/Greenspun%27s_tenth_rule

This is a particularly pernicious kind of **Feature Risk** which can only be mitigated by good **Design**. Human needs are fractal in nature: the more you examine them, the more differences you can find. The aim of a product is to capture some needs at a *general* level: you can't hope to "please all of the people all of the time".

Conceptual Integrity Risk is the risk that chasing after features leaves the product making no sense, and therefore pleasing no-one.

Feature Access Risk



- Risks due to some clients not having access to some or all of the features in your product.

Figure 11.5: Feature Access Risk

Sometimes, features can work for some people and not others: this could be down to Accessibility³ issues, language barriers or localization.

You could argue that the choice of *platform* is also going to limit access: writing code for XBox-only leaves PlayStation owners out in the cold. This is *largely Feature Access Risk*, though **Dependency Risk** is related here.

In Marketing, minimizing **Feature Access Risk** is all about **Segmentation**: trying to work out *who* your product appeals to, and tailoring it to that particular market, but for technologists, increasing **Feature Access** means increasing complexity: you have to deliver the software on more platforms, localized in more languages, with different configurations of features at different price-points. Mitigating **Feature Access Risk** therefore means increased effort and complexity (which we'll come to later).

Market Risk

Feature Access Risk is related, of course, to **Market Risk**, which I introduced on the **Risk Landscape** page as being the value that the market places on a particular asset. Since the product you are building is your asset, it makes sense that you'll face **Market Risk** on it:

"Market risk is the risk of losses in positions arising from movements in market prices." - Market Risk, Wikipedia⁴

³<https://en.wikipedia.org/wiki/Accessibility>

⁴https://en.wikipedia.org/wiki/Market_risk



- Risk that the value your clients place on the features you supply will change, over time.

Figure 11.6: Market Risk

I face market risk when I own (i.e. have a *position* in) some Apple⁵ stock. Apple's⁶ stock price will decline if a competitor brings out an amazing product, or if fashions change and people don't want their products any more.

In the same way, you have **Market Risk** on the product or service you are building: the *market* decides what it is prepared to pay for this, and it tends to be outside your control.

Feature Drift Risk



- Risk that the features required by clients will change and evolve over time.

Figure 11.7: Feature Drift Risk

Feature Drift is the tendency that the features people need *change over time*. For example, at one point in time, supporting IE6 was right up there for website developers, but it's not really relevant anymore. Although that change took *many years* to materialize, other changes are more rapid.

The point is: **Requirements captured today** might not make it to *tomorrow*, especially in the fast-paced world of IT. This is partly because the market *evolves* and becomes more discerning. This happens in several ways: - Features present in competitor's versions of the software become *the baseline*, and they're expected to be available in your version. - Certain ways of interacting become the norm (e.g. **qwerty** keyboards, or the control layout in cars: these don't

⁵<http://apple.com>

⁶<http://apple.com>

change with time). - Features decline in usefulness: *Printing* is less important now than it was, for example.

Feature Drift Risk is *not the same thing* as **Requirements Drift**, which is the tendency projects have to expand in scope as they go along. There are lots of reasons they do that, a key one being the **Hidden Risks** uncovered on the project as it progresses.

Fashion

Fashion plays a big part in IT, as this infographic on website design shows⁷. True, websites have got easier to use as time has gone by, and users now expect this. Also, bandwidth is greater now, which means we can afford more media and code on the client side. However, *fashion* has a part to play in this.

By being *fashionable*, websites are communicating: *this is a new thing, this is relevant, this is not terrible*: all of which is mitigating a **Communication Risk**. Users are all-too-aware that the Internet is awash with terrible, abandon-ware sites that are going to waste their time. How can you communicate that you're not one of them to your users?

Delight

If this breakdown of **Feature Risk** seems reductive, then try not to think of it that way: the aim of course should be to delight users, and turn them into fans. That's a laudable **Goal**, but should be treated in the usual Risk-First way: *pick the biggest risk you can mitigate next*.

Consider **Feature Risk** from both the down-side and the up-side:

- What are we missing?
- How can we be *even better*?

Hopefully, this has given you some ideas about what **Feature Risk** involves. Hopefully, you might be able to identify a few more specific varieties. But, it's time to move on and look in more detail at **Complexity Risk** and how it affects what we build.

Analysis

At this point, it would be easy to stop and say, look, here are a bunch of **Feature Risk** issues that you could face. But, it turns out that we're going to be relying heavily on **Feature Risk** as we go on in order to build our understanding of other risks, so it's probably worth spending a bit of time up front to classify what we've found.

The **Feature Risks** identified here basically exist in a 3-dimensional space: - **Fit**: How well the features fit for a particular client. - **Audience**: The range of clients (the *market*) that may be able to use this feature. - **Evolution**: The way the fit and the audience changes and evolves as time goes by.

⁷<https://designers.hubspot.com/blog/the-history-of-web-design-infographic>

Fit

“Survival Of The Fittest” - Darwin, tbd.

Darwin’s conception of fitness was not one of athletic prowess, but how well an organism worked within the landscape.

tbd: definition of biological fitness

Fit Risk, Conceptual Integrity Risk and Implementation Risk all hint at different aspects of this “fitness”. We can conceive of the relationships between them in the following way:

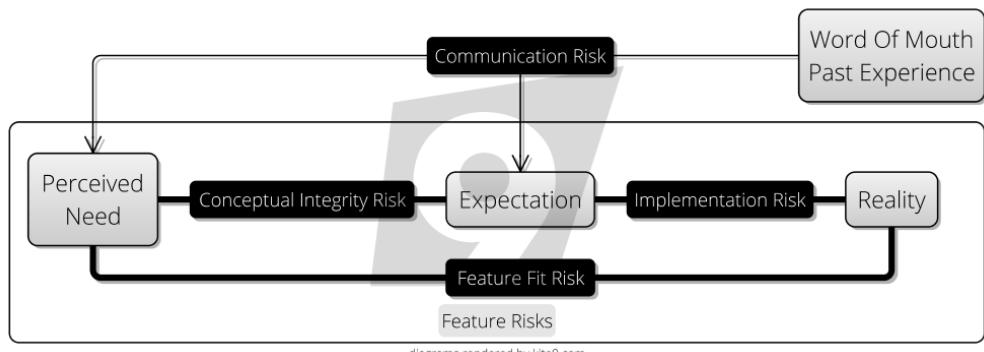


Figure 11.8: Feature Risks Assembled

For further reading, you can check out The Service Quality Model⁸, which this model is derived from. This model analyses the types of *quality gaps* in services, and how consumer expectations and perceptions of a service arise. In **Staging And Classifying**, we’ll come back and build on this model further.

Audience

- We’ve looked at the market of all possible users, and considered which of those will find our product suitable. This is the **Market** dimension.

tbd: show 5 users, each with slightly different requirements, how they intersect with the product. (market risk, feature access risk)

- We’ve looked at time, and how this can change both the market and the features. This is the **Change** dimension.

tbd. regression risk and feature drift risk.

Applying Feature Risk

Consider **Feature Risk** carefully next time you are grooming the backlog: - Can you judge which tasks mitigate the most **Feature Risk**? - Are you delivering features that are valuable

⁸<http://en.wikipedia.org/SERVQUAL>

to a large audience? How well do you understand your audience? How does the size of the audience for a task impact its importance in the backlog? - Does the audience *know* that the features exist? How do you communicate feature availability to them?

- How does writing a specification mitigate **Fit Risk**? For what other reasons are you writing specifications?

In the next section, we are going to unpack this third point further. Somewhere between “what the customer wants” and “what you give them” is a *dialog*. In using a software product, users are engaging in a *dialog* with its features. If the features don’t exist, hopefully they will engage in a dialog with the development team to get them added.

These dialogs are prone to risk, and this is the subject of the next section, **Communication-Risk**.

Gaps

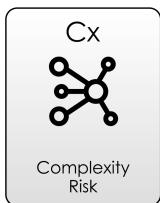
<https://community.verint.com/b/customer-engagement/posts/service-quality-gap-model>

GAP model' developed by a group of authors- Parasuraman, Zeithaml and Berry at Texas and North Carolina in 1985 ,

Chapter 12

Complexity Risk

Complexity Risk are the risks to your project due to its underlying “complexity”. Over the next few sections, we’ll break down exactly what we mean by complexity, looking at **Dependency Risk** and **Boundary Risk** as two particular sub-types of **Complexity Risk**. However, in this section, we’re going to be specifically focusing on *code you write*: the size of your code-base, the number of modules, the interconnectedness of the modules and how well-factored the code is.



- Risks caused by the weight of complexity in the systems we create, and their resistance to change and comprehension.

Figure 12.1: Complexity Risks

You could think of this section, then, as **Codebase Risk**: We’ll look at three separate measures of codebase complexity and talk about **Technical Debt**, and look at places in which **Codebase Risk** is at its greatest.

Kolmogorov Complexity

The standard Computer-Science definition of complexity, is Kolmogorov Complexity¹. This is:

“...is the length of the shortest computer program (in a predetermined programming language) that produces the object as output.” - Kolmogorov Complexity,

¹https://en.wikipedia.org/wiki/Kolmogorov_complexity

Wikipedia²

This is a fairly handy definition for us, as it means that to in writing software to solve a problem, there is a lower bound on the size of the software we write. In practice, this is pretty much impossible to quantify. But that doesn't really matter: the techniques for *moving in that direction* are all that we are interested in, and this basically amounts to compression.

Let's say we wanted to write a javascript program to output this string:

```
abcdabcdabcdabcdabcdabcdabcdabcdabcdabcd
```

We might choose this representation:

```
function out() {  
    return "abcdabcdabcdabcdabcdabcdabcdabcd"  
}  
(7 symbols)  
(45 symbols)  
(1 symbol)
```

... which contains 53 symbols, if you count `function`, `out` and `return` as one symbol each.

But, if we write it like this:

```
const ABCD="ABCD";  
  
function out() {  
    return ABCD+ABCD+ABCD+ABCD+ABCD+ABCD+ABCD+ABCD+ABCD+ABCD  
}  
(11 symbols)  
(7 symbols)  
(21 symbols)  
(1 symbol)
```

With this version, we now have 40 symbols. And with this version:

```
const ABCD="ABCD";  
  
function out() {  
    return ABCD.repeat(10)  
}  
(11 symbols)  
(7 symbols)  
(7 symbols)  
(1 symbol)
```

... we have 26 symbols.

Abstraction

What's happening here is that we're *exploiting a pattern*: we noticed that `ABCD` occurs several times, so we defined it a single time and then used it over and over, like a stamp. Separating the *definition* of something from the *use* of something as we've done here is called "abstraction". We're going to come across it over and over again in this part of the book, and not just in terms of computer programs.

By applying techniques such as Abstraction, we can improve in the direction of the Kolmogorov limit. And, by allowing ourselves to say that *symbols* (like `out` and `ABCD`) are worth one complexity point, we've allowed that we can be descriptive in our function name and `const`. Naming things is an important part of abstraction, because to use something, you have to be able to refer to it.

²https://en.wikipedia.org/wiki/Kolmogorov_complexity

Trade-Off

But we could go further down into Code Golf³ territory. This javascript program plays FizzBuzz⁴ up to 100, but is less readable than you might hope:

```
for(i=0;i<100;)document.write(((++i%3?'Fizz':(i%5?'Buzz':))||i)+"<br>")(66 symbols)
```

So there is at some point a trade-off to be made between **Complexity Risk** and **Communication Risk**. This is a topic we'll address more in that section. But for now, it should be said that **Communication Risk** is about *misunderstanding*: The more complex a piece of software is, the more difficulty users will have understanding it, and the more difficulty developers will have changing it.

Connectivity

A second, useful measure of complexity comes from graph theory, and that is the connectivity of a graph:

“...the minimum number of elements (nodes or edges) that need to be removed to disconnect the remaining nodes from each other” - Connectivity, *Wikipedia*⁵

To see this in action, have a look at the below graph:

It has 10 vertices, labelled a to j, and it has 15 edges (or links) connecting the vertices together. If any single edge were removed from this diagram, the 10 vertices would still be linked together. Because of this, we can say that the graph is *2-connected*. That is, to disconnect any single vertex, you'd have to remove *at least* two edges.

As a slight aside, let's consider the **Kolmogorov Complexity** of this graph, by inventing a mini-language to describe graphs. It could look something like this:

```
<item> : [<item>,]* <item>      # Indicates that the item before the colon  
                                # has a connection to all the items after the colon.  
  
a: b,c,d  
b: c,f,e  
c: f,d  
d: j  
e: h,j  
f: h  
g: j  
h: i  
i:j(39 symbols)
```

Let's remove some of those extra links:

³https://en.wikipedia.org/wiki/Code_golf

⁴https://en.wikipedia.org/wiki/Fizz_buzz

⁵[https://en.wikipedia.org/wiki/Connectivity_\(graph_theory\)](https://en.wikipedia.org/wiki/Connectivity_(graph_theory))

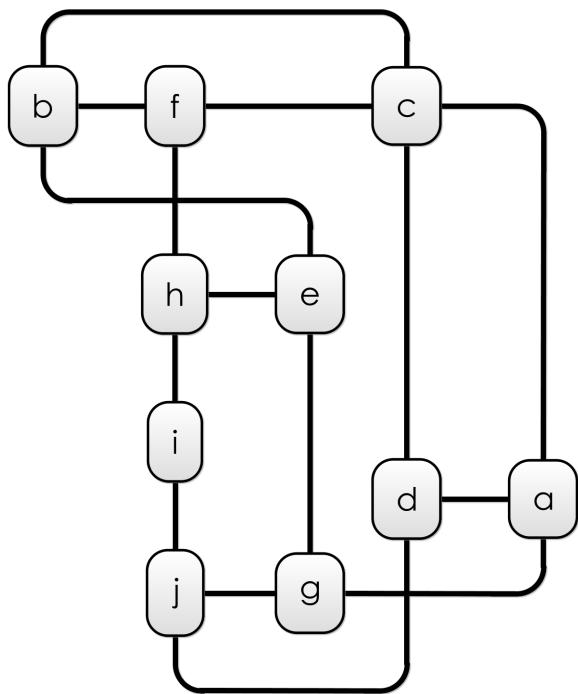


Figure 12.2: Graph 1

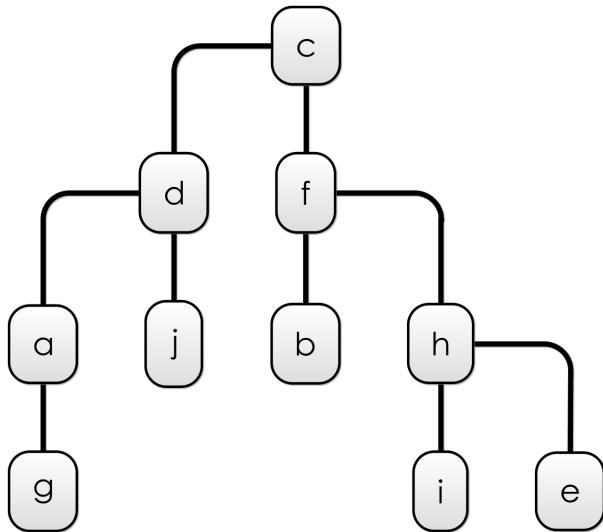


Figure 12.3: Graph 2

In this graph, I've removed 6 of the edges. Now, we're in a situation where if any single edge is removed, the graph becomes *unconnected*. That is, it's broken into distinct chunks. So, it's *1-connected*.

The second graph is clearly simpler than the first. And, we can show this by looking at the **Kolgomorov Complexity** in our little language:

a: d,g	
b: f	
c: d,f	
d: j	
f: h	
e: h	
h: i	(25 symbols)

Connectivity is also **Complexity**. Heavily connected programs/graphs are much harder to work with than less-connected ones. Even *laying out* the first graph sensibly is a harder task than the second (the second is a doddle). But the reason programs with greater connectivity are harder to work with is that changing one module potentially impacts many others.

Hierarchies and Modularization

In the second, simplified graph, I've arranged it as a hierarchy, which I can do now that it's only 1-connected. For 10 vertices, we need 9 edges to connect everything up. It's always:

```
edges = vertices - 1
```

Note that I could pick any hierarchy here: I don't have to start at **c** (although it has the nice property that it has two roughly even sub-trees attached to it).

How does this help us? Imagine if **a - j** were modules of a software system, and the edges of the graph showed communications between the different sub-systems. In the first graph, we're in a worse position: who's in charge? What deals with what? Can I isolate a component and change it safely? What happens if one component disappears? But, in the second graph, it's easier to reason about, because of the reduced number of connections and the new heirarchy of organisation.

On the downside, perhaps our messages have farther to go now: in the original **i** could send a message straight to **j**, but now we have to go all the way via **c**. But this is the basis of Modularization⁶ and Hierarchy⁷.

As a tool to battle complexity, we don't just see this in software, but everywhere in our lives. Society, business, nature and even our bodies:

- **Organelles** - such as Mitochondria⁸.
- **Cells** - such as blood cells, nerve cells, skin cells in the Human Body⁹.
- **Organs** - like hearts livers, brains etc.
- **Organisms** - like you and me.

The great complexity-reducing mechanism of modularization is that *you only have to consider your local environment*. Elements of the program that are “far away” in the hierarchy can be relied on not to affect you. This is somewhat akin to the **Principal Of Locality**:

“Spatial locality refers to the use of data elements within relatively close storage locations.” - Locality Of Reference, *Wikipedia*¹⁰

Cyclomatic Complexity

A variation on this graph connectivity metric is our third measure of complexity, Cyclomatic Complexity¹¹. This is:

$\text{Cyclomatic Complexity} = \text{edges} - \text{vertices} + 2P,$

Where **P** is the number of **Connected Components** (i.e. distinct parts of the graph that aren't connected to one another by any edges).

⁶https://en.wikipedia.org/wiki/Modular_programming

⁷<https://en.wikipedia.org/wiki/Hierarchy>

⁸<https://en.wikipedia.org/wiki/Mitochondrion>

⁹https://en.wikipedia.org/wiki/List_of_distinct_cell_types_in_the_adult_human_body

¹⁰https://en.wikipedia.org/wiki/Locality_of_reference

¹¹https://en.wikipedia.org/wiki/Cyclomatic_complexity

So, our first graph had a **Cyclomatic Complexity** of 7. ($15 - 10 + 2$), while our second was 1. ($9 - 10 + 2$).

Cyclomatic complexity is all about the number of different routes through the program. The more branches a program has, the greater its cyclomatic complexity. Hence, this is a useful metric in **Testing** and **Code Coverage**: the more branches you have, the more tests you'll need to exercise them all.

More Abstraction

Although we ended up with our second graph having a **Cyclomatic Complexity** of 1 (the minimum), we can go further through abstraction, because this representation isn't minimal from a **Kolmogorov Complexity** point-of-view. For example, we might observe that there are further similarities in the graph that we can "draw out":

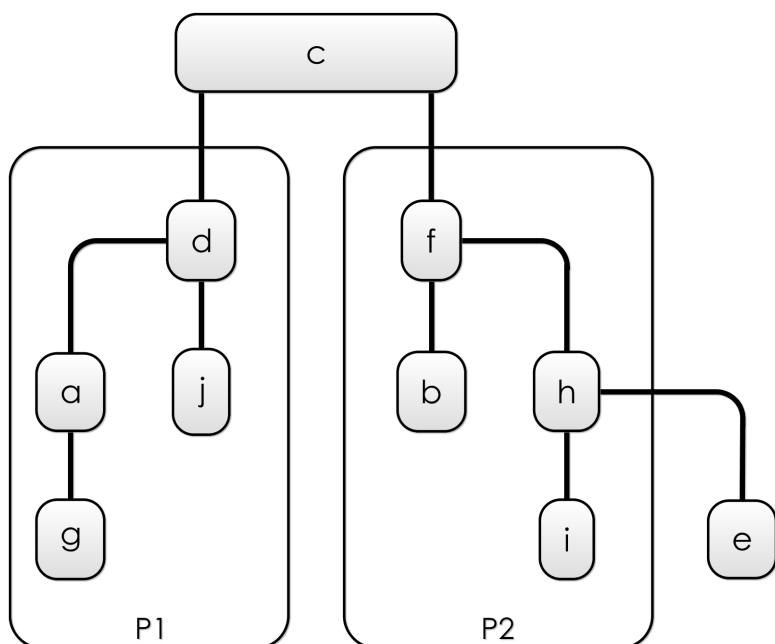


Figure 12.4: Complexity 3

Here, we've spotted that the structure of subgraphs **P1** and **P2** are the same: we can have the same functions there to assemble those. Noticing and exploiting patterns of repetition is one of the fundamental tools we have in the fight against **Complexity Risk**.

So, we've looked at some measures of software structure complexity, in order that we can say

“this is more complex than this”. However, we’ve not really said why complexity entails **Risk**. So let’s address that now by looking at two analogies, **Mass** and **Technical Debt**.

Complexity As Mass

The first way to look at complexity is as **Mass** or **Inertia**: a software project with more complexity has greater **Inertia** or **Mass** than one with less complexity.

Newton’s Second Law states:

“ $F = ma$, (Force = Mass x Acceleration)” - Netwon’s Laws Of Motion, *Wikipedia*¹²

That is, in order to move your project *somewhere new*, and make it do new things, you need to give it a push, and the more **Mass** it has, the more **Force** you’ll need to move (accelerate) it.

Inertia and **Mass** are equivalent concepts in physics:

“mass is the quantitative or numerical measure of a body’s inertia, that is of its resistance to being accelerated”. - *Inertia, Wikipedia*¹³

You could stop here and say that the more lines of code a project contains, the higher it’s mass. And, that makes sense, because in order to get it to do something new, you’re likely to need to change more lines.

But there is actually some underlying sense in which *this is real*, as discussed in this Veritasium¹⁴ video. To paraphrase:

“Most of your mass you owe due to $E=mc^2$, you owe to the fact that your mass is packed with energy, because of the **interactions** between these quarks and gluon fluctuations in the gluon field... what we think of as ordinarily empty space... that turns out to be the thing that gives us most of our mass.” - Your Mass is NOT From the Higgs Boson, *Veritasium*¹⁵

I’m not an expert in physics, *at all*, and so there is every chance that I am pushing this analogy too hard. But, substituting quarks and gluons for pieces of software we can (in a very handwaving-y way) say that more complex software has more **interactions** going on, and therefore has more mass than simple software.

The reason I am labouring this analogy is to try and make the point that **Complexity Risk** is really fundamental:

- **Feature Risk**: like **money**.
- **Schedule Risk**: like **time**.
- **Complexity Risk**: like **mass**.

At a basic level, **Complexity Risk** heavily impacts on **Schedule Risk**: more complexity means you need more force to get things done, which takes longer.

¹²https://en.wikipedia.org/wiki/Newton%27s_laws_of_motion

¹³https://en.wikipedia.org/wiki/Inertia#Mass_and_inertia

¹⁴<https://www.youtube.com/user/1veritasium>

¹⁵https://www.youtube.com/watch?annotation_id=annotation_3771848421&feature=iv&src_vid=Xo232kyTsOo&v=Ztc6QPNUqls

Technical Debt

The most common way we talk about unnecessary complexity in software is as **Technical Debt**:

“Shipping first time code is like going into debt. A little debt speeds development so long as it is paid back promptly with a rewrite... The danger occurs when the debt is not repaid. Every minute spent on not-quite-right code counts as interest on that debt. Entire engineering organizations can be brought to a stand-still under the debt load of an unconsolidated implementation, object-oriented or otherwise.” – Ward Cunningham, 1992¹⁶

Building a perfect first-time solution is a waste, because perfection takes a long time. You’re taking on more attendant **Schedule Risk** than necessary and **Meeting Reality** more slowly than you could.

A quick-and-dirty, over-complex implementation mitigates the same **Feature Risk** and allows you to **Meet Reality** faster (see **Prototyping**).

But, having mitigated the **Feature Risk**, you are now carrying more **Complexity Risk** than you necessarily need, and it’s time to think about how to **Refactor** the software to reduce this risk again.

Kitchen Analogy

It’s often hard to make the case for minimizing **Technical Debt**: it often feels that there are more important priorities, especially when technical debt can be “swept under the carpet” and forgotten about until later. (See **Discounting The Future**.)

One helpful analogy I have found is to imagine your code-base is a kitchen. After preparing a meal (i.e. delivering the first implementation), *you need to tidy up the kitchen*. This is just something everyone does as a matter of *basic sanitation*.

Now of course, you could carry on with the messy kitchen. When tomorrow comes and you need to make another meal, you find yourself needing to wash up saucepans as you go, or working around the mess by using different surfaces to chop on.

It’s not long before someone comes down with food poisoning.

We wouldn’t tolerate this behaviour in a restaurant kitchen, so why put up with it in a software project?

Feature Creep

In Brooks’ essay “No Silver Bullet – Essence and Accident in Software Engineering”, a distinction is made between:

- **Essence:** *the difficulties inherent in the nature of the software.*

¹⁶https://en.wikipedia.org/wiki/Technical_debt

- **Accident:** *those difficulties that attend its production but are not inherent.*
 - Fred Brooks, *No Silver Bullet*¹⁷

The problem with this definition is that we are accepting features of our software as *essential*.

The **Risk-First** approach is that if you want to mitigate some **Feature Risk** then you have to pick up **Complexity Risk** as a result. But, that's a *choice you get to make*.

Therefore, Feature Creep¹⁸ (or Gold Plating¹⁹) is a failure to observe this basic equation: instead of considering this trade off, you're building every feature possible. This has an impact on **Complexity Risk**, which in turn impacts **Communication Risk** and also **Schedule Risk**.

Sometimes, feature-creep happens because either managers feel they need to keep their staff busy, or the staff decide on their own that they need to **keep themselves busy**. But now, we can see that basically this boils down to bad risk management.

“Perfection is Achieved Not When There Is Nothing More to Add, But When There Is Nothing Left to Take Away” - Antoine de Saint-Exupery

Dead-End Risk



- The risk that a particular approach to a change will fail. Caused by the fact that at some level, our internal models are not a complete reflection of reality.

Figure 12.5: Dead-End Risk

Dead-End Risk is where you build functionality that you *think* is useful, only to find out later that actually, it was a dead-end, and is superceded by something else.

For example, let's say that the Accounting sub-system needed password protection (so you built this). Then the team realised that you needed a way to *change the password* (so you built that). Then, that you needed to have more than one user of the Accounting system so they would all need passwords (ok, fine).

Finally, the team realises that actually logging-in would be something that all the sub-systems would need, and that it had already been implemented more thoroughly by the Approvals sub-system.

At this point, you realise you're in a **Dead End**:

¹⁷https://en.wikipedia.org/wiki/No_Silver_Bullet

¹⁸https://en.wikipedia.org/wiki/Feature_creep

¹⁹[https://en.wikipedia.org/wiki/Gold_plating_\(software_engineering\)](https://en.wikipedia.org/wiki/Gold_plating_(software_engineering))

- **Option 1:** You carry on making minor incremental improvements to the accounting password system (carrying the extra **Complexity Risk** of the duplicated functionality).
- **Option 2:** You rip out the accounting password system, and merge in the Approvals system, surfacing new, hidden **Complexity Risk** in the process, due to the difficulty in migrating users from the old to new way of working.
- **Option 3:** You start again, trying to take into account both sets of requirements at the same time, again, possibly surfacing new hidden **Complexity Risk** due to the combined approach.

Sometimes, the path from your starting point to your goal on the **Risk Landscape** will take you to dead ends: places where the only way towards your destination is to lose something, and do it again another way.

This is because you surface new **Hidden Risk** along the way. And the source of a lot of this hidden risk will be unexpected **Complexity Risk** in the solutions you choose. This happens a lot.

Source Control

Version Control Systems²⁰ like Git²¹ are a useful mitigation of **Dead-End Risk**, because it means you can *go back* to the point where you made the bad decision and go a different way. Additionally, they provide you with backups against the often inadvertent **Dead-End Risk** of someone wiping the hard-disk.

The Re-Write

Option 3, Rewriting code or a whole project can seem like a way to mitigate **Complexity Risk**, but it usually doesn't work out too well. As Joel Spolsky says:

There's a subtle reason that programmers always want to throw away the code and start over. The reason is that they think the old code is a mess. And here is the interesting observation: they are probably wrong. The reason that they think the old code is a mess is because of a cardinal, fundamental law of programming:
It's harder to read code than to write it. - Things You Should Never Do, Part 1, Joel Spolsky²²

The problem that Joel is outlining here is that the developer mistakes hard-to-understand code for unnecessary **Complexity Risk**. Also, perhaps there is **Agency Risk** because the developer is doing something that is more useful to him than the project. We're going to return to this problem in again **Communication Risk**.

²⁰https://en.wikipedia.org/wiki/Version_control

²¹<https://en.wikipedia.org/wiki/Git>

²²<https://www.joelonsoftware.com/2000/04/06/things-you-should-never-do-part-i/>

Where Complexity Hides

Complexity isn't spread evenly within a software project. Some problems, some areas, have more than their fair share of issues. We're going to cover a few of these now, but be warned, this is not a complete list by any means:

- Memory Management
- Protocols / Types
- Algorithmic (Space and Time) Complexity
- Concurrency / Mutability
- Networks / Security

Memory Management

Memory Management is another place where **Complexity Risk** hides:

"Memory leaks are a common error in programming, especially when using languages that have no built in automatic garbage collection, such as C and C++." -
Memory Leak, *Wikipedia*²³

Garbage Collectors²⁴ (as found in Javascript or Java) offer you the deal that they will mitigate the **Complexity Risk** of you having to manage your own memory, but in return perhaps give you fewer guarantees about the *performance* of your software. Again, there are times when you can't accommodate this **Operational Risk**, but these are rare and usually only affect a small portion of an entire software-system.

Protocols And Types

Whenever two components of a software system need to interact, they have to establish a protocol for doing so. There are lots of different ways this can work, but the simplest example I can think of is where some component **a** calls some function **b**. e.g:

```
function b(a, b, c) {  
    return "whatever" // do something here.  
}  
  
function a() {  
    var bOut = b("one", "two", "three");  
    return "something "+bOut;  
}
```

If component **b** then changes in some backwards-incompatible way, say:

```
function b(a, b, c, d /* new parameter */) {  
    return "whatever" // do something here.  
}
```

²³https://en.wikipedia.org/wiki/Memory_leak

²⁴[https://en.wikipedia.org/wiki/Garbage_collection_\(computer_science\)](https://en.wikipedia.org/wiki/Garbage_collection_(computer_science))

Then, we can say that the protocol has changed. This problem is so common, so endemic to computing that we've had compilers that check function arguments since the 1960's²⁵. The point being is that it's totally possible for the compiler to warn you about when a protocol within the program has changed.

The same is basically true of Data Types²⁶: whenever we change the **Data Type**, we need to correct the usages of that type. Note above, I've given the javascript example, but I'm going to switch to typescript now:

```
interface BInput {  
    a: string,  
    b: string,  
    c: string,  
    d: string  
}  
  
function b(in: BInput): string {  
    return "whatever" // do something here.  
}
```

Now, of course, there is a tradeoff: we *mitigate Complexity Risk*, because we define the protocols / types *once only* in the program, and ensure that usages all match the specification. But the tradeoff is (as we can see in the typescript code) more *finger-typing*, which some people argue counts as **Schedule Risk**.

Nevertheless, compilers and type-checking are so prevalent in software that clearly, you have to accept that in most cases, the trade-off has been worth it: Even languages like Clojure²⁷ have been retro-fitted with type checkers²⁸.

We're going to head into much more detail on this in the section on **Protocol Risk**.

Space and Time Complexity

So far, we've looked at a couple of definitions of complexity in terms of the codebase itself. However, in Computer Science there is a whole branch of complexity theory devoted to how the software *runs*, namely Big O Complexity²⁹.

Once running, an algorithm or data structure will consume space or runtime dependent on its characteristics. As with Garbage Collectors³⁰, these characteristics can introduce **Performance Risk** which can easily catch out the unwary. By and large, using off-the-shelf data structures and algorithms helps, but you still need to know their performance characteristics.

The Big O Cheatsheet³¹ is a wonderful resource to investigate this further.

²⁵<https://en.wikipedia.org/wiki/Compiler>

²⁶https://en.wikipedia.org/wiki/Data_type

²⁷<https://clojure.org>

²⁸<https://github.com/clojure/core.typed/wiki/User-Guide>

²⁹https://en.wikipedia.org/wiki/Big_O_notation

³⁰[https://en.wikipedia.org/wiki/Garbage_collection_\(computer_science\)](https://en.wikipedia.org/wiki/Garbage_collection_(computer_science))

³¹<http://bigocheatsheet.com>

Concurrency / Mutability

Although modern languages include plenty of concurrency primitives, (such as the `java.util.concurrent`³² libraries), concurrency is *still* hard to get right.

Race conditions³³ and Deadlocks³⁴ *thrive* in over-complicated concurrency designs: complexity issues are magnified by concurrency concerns, and are also hard to test and debug.

Recently, languages such as Clojure³⁵ have introduced persistent collections³⁶ to alleviate concurrency issues. The basic premise is that any time you want to *change* the contents of a collection, you get given back a *new collection*. So, any collection instance is immutable once created. The tradeoff is again attendant **Performance Risk** to mitigate **Complexity Risk**.

An important lesson here is that choice of language can reduce complexity: and we'll come back to this in **Software Dependency Risk**.

Networking / Security

The last area I want to touch on here is networking. There are plenty of **Complexity Risk** perils in *anything* to do with networked code, chief amongst them being error handling and (again) **protocol evolution**.

In the case of security considerations, exploits *thrive* on the complexity of your code, and the weaknesses that occur because of it. In particular, Schneier's Law says, never implement your own crypto scheme:

“Anyone, from the most clueless amateur to the best cryptographer, can create an algorithm that he himself can't break. It's not even hard. What is hard is creating an algorithm that no one else can break, even after years of analysis.” - Bruce Schneier, 1998³⁷

Luckily, most good languages include crypto libraries that you can include to mitigate these **Complexity Risks** from your own code-base.

This is a strong argument for the use of libraries. But, when should you use a library and when should you implement yourself? This is again covered in the section on **Software Dependency Risk**.

tbd - next section.

costs associated with complexity risk

CHANGE is also more risky why?

³²<https://docs.oracle.com/javase/9/docs/api/java/util/concurrent/package-summary.html>

³³https://en.wikipedia.org/wiki/Race_condition

³⁴<https://en.wikipedia.org/wiki/Deadlock>

³⁵<https://clojure.org>

³⁶https://en.wikipedia.org/wiki/Persistent_data_structure

³⁷https://en.wikipedia.org/wiki/Bruce_Schneier#Cryptography

Chapter 13

Communication Risk

Communication Risk is the risk of communication between entities *going wrong*, due to loss or misunderstanding. Consider this: if we all had identical knowledge, there would be no need to do any communicating at all, and therefore and also no **Communication Risk**.



- Risks due to the difficulty of communicating with other entities, be they people, software, processes etc.

Figure 13.1: Communication Risk

But, people are not all-knowing oracles. We rely on our *senses* to improve our **Internal Models** of the world. There is **Communication Risk** here - we might overlook something vital (like an oncoming truck) or mistake something someone says (like “Don’t cut the green wire”).

Communication Risk isn’t just for people; it affects computer systems too.

A Model Of Communication

In 1948, Claude Shannon proposed this definition of communication:

“The fundamental problem of communication is that of reproducing at one point, either exactly or approximately, a message selected at another point.” - A Mathematical Theory Of Communication, *Claude Shannon*¹

¹https://en.wikipedia.org/wiki/A_Mathematical_Theory_of_Communication

And from this same paper, we get the following (slightly adapted) model.

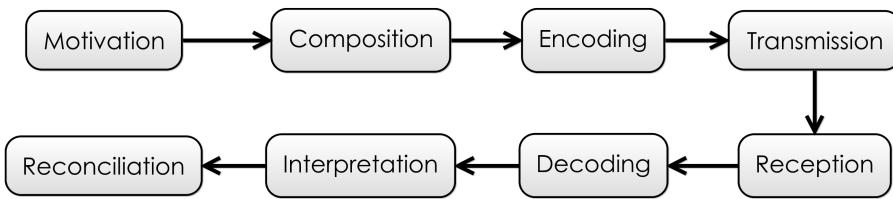


Figure 13.2: Communication Model

We move from top-left (“I want to send a message to someone”) to bottom left, clockwise, where we hope the message has been understood and believed. (I’ve added this last box to Shannon’s original diagram.)

One of the chief concerns in Shannon’s paper is the step between **Transmission** and **Reception**. He creates a theory of information (measured in **bits**), the upper-bounds of information that can be communicated over a channel and ways in which **Communication Risk** between these processes can be mitigated by clever **Encoding** and **Decoding** steps.

But it’s not just transmission. **Communication Risk** exists at each of these steps. Let’s imagine a short exchange where someone, **Alice** is trying to send a message to **Bob**:

- **Alice** might be **motivated** to send a message to tell **Bob** something, only to find out that he already knew it, or it wasn’t useful information for them.
- In the **composition** stage, **Alice** might mess up the *intent* of the message: instead of “Please buy chips” she might say, “Please buy chops”.
- In the **encoding** stage, **Alice** might not speak clearly enough to be understood, and...
- In the **transmission** stage, **Alice** might not say it loudly enough for **Bob** to...
- **receive** the message clearly (maybe there is background noise).
- Having heard **Alice** say something, can **Bob** **decode** what was said into a meaningful sentence?
- Then, assuming that, will they **interpret** correctly which type of chips (or chops) **Alice** was talking about? Does “Please buy chips” convey all the information they need?
- Finally, assuming *everything else*, will **Bob** believe the message? Will they **reconcile** the information into their **Internal Model** and act on it? Perhaps not, if **Bob** thinks that there are chips at home already.

Approach To Communication Risk

There is a symmetry about the steps going on in Shannon’s diagram, and we’re going to exploit this in order to break down **Communication Risk** into its main types.

To get inside **Communication Risk**, we need to understand **Communication** itself, whether between *machines, people or products*: we’ll look at each in turn. In order to do that, we’re going to examine four basic concepts in each of these settings:

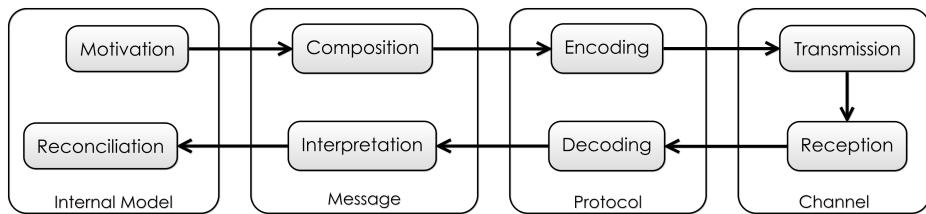


Figure 13.3: Communication Risk 2

- Channels², the medium via which the communication is happening.
- Protocols³ - the systems of rules that allow two or more entities of a communications system to transmit information.
- Messages⁴: The information we want to convey.
- **Internal Models**: the sources and destinations for the messages. Updating internal models (whether in our heads or machines) is the reason why we're communicating.

And, as we look at these four areas, we'll consider the **Attendant Risks** of each.

Channels

There are lots of different types of media for communicating (e.g. TV, Radio, DVD, Talking, Posters, Books, Phones, The Internet, etc.) and they all have different characteristics. When we communicate via a given medium, it's called a *channel*.

The channel *characteristics* depend on the medium, then. Some obvious ones are cost, utilisation, number of people reached, simplex or duplex (parties can transmit and receive at the same time), persistence (a play vs a book, say), latency (how long messages take to arrive) and bandwidth (the amount of information that can be transmitted in a period of time).

Channel characteristics are important: in a high-bandwidth, low-latency situation, **Alice** and **Bob** can *check* with each other that the meaning was transferred correctly. They can discuss what to buy, they can agree that **Alice** wasn't lying or playing a joke.

The channel characteristics also imply suitability for certain *kinds* of messages. A documentary might be a great way of explaining some economic concept, whereas an opera might not be.

Channel Risk

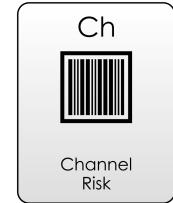
Shannon discusses that no channel is perfect: there is always the **risk of noise** corrupting the signal. A key outcome from Shannon's paper is that there is a tradeoff: within the capacity of

²https://en.wikipedia.org/wiki/Communication_channel

³https://en.wikipedia.org/wiki/Communication_protocol

⁴<https://en.wikipedia.org/wiki/Message>

the channel (the **Bandwidth**), you can either send lots of information with *higher* risk that it is wrong, or less information with *lower* risk of errors. And, rather like the **Kolgomorov complexity** result, the more *randomness* in the signal, the less compressible it is, and therefore the more *bits* it will take to transmit.



- Risks due to the inadequacy of the physical channel used to communicate our messages. e.g. noise, loss, interception, corruption.

Figure 13.4: Communication Channel Risk

But channel risk goes wider than just this mathematical example: messages might be delayed or delivered in the wrong order, or not be acknowledged when they do arrive. Sometimes, a channel is just an inappropriate way of communicating. When you work in a different time-zone to someone else on your team, there is *automatic* **Channel Risk**, because instantaneous communication is only available for a few hours' a day.

When channels are **poor-quality**, less communication occurs. People will try to communicate just the most important information. But, it's often impossible to know a-priori what constitutes "important". This is why **Extreme Programming** recommends the practice of **Pair Programming** and siting all the developers together: although you don't know whether useful communication will happen, you are mitigating **Channel Risk** by ensuring high-quality communication channels are in place.

At other times, channels can contain so much information that we can't hope to receive all the messages. In these cases, we don't even observe the whole channel, just parts of it. For example, you might have a few YouTube channels that you subscribe to, but hundreds of hours of video are being posted on YouTube every second, so there is no way you can keep up with all of it.

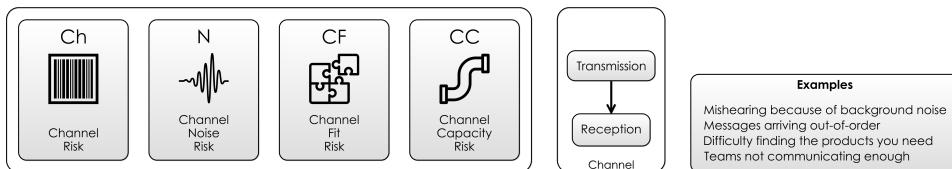


Figure 13.5: Communication Channels

Marketing Communications

When we are talking about a product or a brand, mitigating **Channel Risk** is the domain of Marketing Communications⁵. How do you ensure that the information about your (useful) project makes it to the right people? How do you address the right channels?

This works both ways. Let's look at some of the **Channel Risks** from the point of view of a hypothetical software tool, **D**, which would be really useful in my software:

- The concept that there is such a thing as **D** which solves my problem isn't something I'd even considered.
- I'd like to use something like **D**, but how do I find it?
- There are multiple implementations of **D**, which is the best one for the task?
- I know **D**, but I can't figure out how to solve my problem in it.
- I've chosen **D**, I now need to persuade my team that **D** is the correct solution...
- ... and then they also need to understand **D** to do their job too.

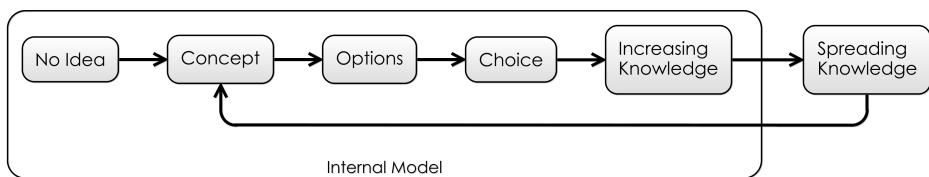


Figure 13.6: Communication Marketing

Internal Models don't magically get populated with the information they need: they fill up gradually, as shown in this diagram. Popular products and ideas *spread*, by word-of-mouth or other means. Part of the job of being a good technologist is to keep track of new **Ideas**, **Concepts** and **Options**, so as to use them as **Dependencies** when needed.

Protocols

In this section, I want to examine the concept of Communication Protocols⁶ and how they relate to **Abstraction**.

So, to do this, let's look in a bit of detail at how web pages are loaded. When considering this, we need to broaden our terminology. Although so far we've talked about **Senders** and **Receivers**, we now need to talk from the point of view of who-depends-on-who. If you're *depended on*, then you're a "Server", whereas if you require communication with something else, you're a "Client". Thus, clients depend on servers in order to load pages.

This is going to involve (at least) six separate protocols, the top-most one being the HTTP Protocol⁷. As far as the HTTP Protocol is concerned, a *client* makes an HTTP Request at a

⁵https://en.wikipedia.org/wiki/Marketing_communications

⁶https://en.wikipedia.org/wiki/Communication_protocol

⁷https://en.wikipedia.org/wiki/Hypertext_Transfer_Protocol

specific URL and the HTTP Response is returned in a predictable format that the browser can understand.

Let's have a quick look at how that works with a `curl` command, which allows me to load a web page from the command line. We're going to try and load Google's preferences page, and see what happens. If I type:

```
> curl -v http://google.com/preferences      # -v indicates verbose
```

1. DNS - Domain Name System

Then, the first thing that happens is this:

```
* Rebuilt URL to: http://google.com/
*   Trying 216.58.204.78...
```

At this point, `curl` has used DNS⁸ to *resolve* the address “google.com” to an IP address. This is some **Abstraction**: instead of using the machine's IP Address⁹ on the network, 216.58.204.78, I can use a human-readable address, `google.com`. The address `google.com` doesn't necessarily resolve to that same address each time: *They have multiple IP addresses for google.com*. But, for the rest of the `curl` request, I'm now set to just use this one.

2. IP - Internet Protocol

But this hints at what is beneath the abstraction: although I'm loading a web-page, the communication to the Google server happens by IP Protocol¹⁰ - it's a bunch of discrete “packets” (streams of binary digits). You can think of a packet as being like a real-world parcel or letter.

Each packet consists of two things:

- An address, which tells the network components (such as routers and gateways) where to send the packet, much like you'd write the address on the outside of a parcel.
- The *payload*, the stream of bytes for processing at the destination. Like the contents of the parcel.

But, even this concept of “packets” is an **Abstraction**. Although all the components of the network interoperate with this protocol, we might be using Wired Ethernet, or WiFi, 4G or *something else*.

3. 802.11 - WiFi Protocol

I ran this at home, using WiFi, which uses IEEE 802.11 Protocol¹¹, which allows my laptop to communicate with the router wirelessly, again using an agreed, standard protocol. But

⁸https://en.wikipedia.org/wiki/Domain_Name_System

⁹https://en.wikipedia.org/wiki/IP_address

¹⁰https://en.wikipedia.org/wiki/Internet_Protocol

¹¹https://en.wikipedia.org/wiki/IEEE_802.11

even *this* isn't the bottom, because this is actually probably specifying something like MIMO-OFDM¹², giving specifications about frequencies of microwave radiation, antennas, multiplexing, error-correction codes and so on. And WiFi is just the first hop: after the WiFi receiver, there will be protocols for delivering the packets via the telephony system.

4. TCP - Transmission Control Protocol

Anyway, the next thing that happens is this:

```
* TCP_NODELAY set  
* Connected to google.com (216.58.204.78) port 80 (#0)
```

The second obvious **Abstraction** going on here is that curl now believes it has a TCP¹³ connection. The TCP connection abstraction gives us the surety that the packets get delivered in the right order, and retried if they go missing. Effectively it *guarantees* these things, or that it will have a connection failure if it can't keep its guarantee.

But, this is a fiction - TCP is built on the IP protocol, packets of data on the network. So there are lots of packets floating around which say "this connection is still alive" and "I'm message 5 in the sequence" and so on in order to maintain this fiction. But that means that the HTTP protocol can forget about this complexity and work with the fiction of a connection.

5. HTTP - Hypertext Transfer Protocol

Next, we see this:

```
> GET /preferences HTTP/1.1      (1)  
> Host: google.com              (2)  
> User-Agent: curl/7.54.0       (3)  
> Accept: */*                   (4)  
>                                (5)
```

This is now the HTTP protocol proper, and these 5 lines are sending information *over the connection* to the Google server.

- (1) says what version of HTTP we are using, and the path we're loading (/preferences in this case).
- (2) to (4) are *headers*. They are name-value pairs, separated with a colon. The HTTP protocol specifies a bunch of these names, and later versions of the protocol might introduce newer ones.
- (5) is an empty line, which indicates that we're done with the headers, please give us the response. And it does:

```
< HTTP/1.1 301 Moved Permanently  
< Location: http://www.google.com/preferences  
< Content-Type: text/html; charset=UTF-8
```

¹²<https://en.wikipedia.org/wiki/MIMO-OFDM>

¹³https://en.wikipedia.org/wiki/Transmission_Control_Protocol

```

< Date: Sun, 08 Apr 2018 10:24:34 GMT
< Expires: Tue, 08 May 2018 10:24:34 GMT
< Cache-Control: public, max-age=2592000
< Server: gws
< Content-Length: 230
< X-XSS-Protection: 1; mode=block
< X-Frame-Options: SAMEORIGIN
<
<HTML><HEAD><meta http-equiv="content-type" content="text/html; charset=utf-8">
<TITLE>301 Moved</TITLE></HEAD><BODY>
<H1>301 Moved</H1>
The document has moved
</BODY></HTML>
* Connection #0 to host google.com left intact

```

There's a lot going on here, but we can break it down really easily into 3 chunks:

- The first line is the HTTP Status Code¹⁴. 301 is a code meaning that the page has moved.
- The next 9 lines are HTTP headers again (name-value pairs). The `Location:` directive tells us where the page has moved to. Instead of trying `http://google.com/preferences`, we should have used `http://www.google.com/preferences`.
- The lines starting `<HTML>` are now some HTML to display on the screen to tell the user that the page has moved.

6. HTML - Hypertext Markup Language

Although HTML¹⁵ is a language, a language is also a protocol. (After all, language is what we use to encode our ideas for transmission as speech.) In the example we gave, this was a very simple page telling the client that it's looking in the wrong place. In most browsers, you don't get to see this: the browser will understand the meaning of the 301 error and redirect you to the location.

Let's look at all the protocols we saw here:

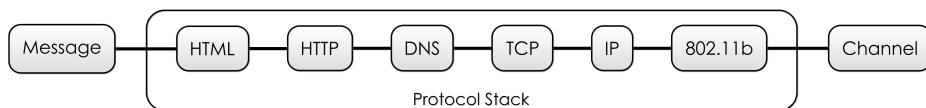


Figure 13.7: Protocol Stack

Each protocol “passes on” to the next one in the chain. On the left, we have the representation most suitable for the *messages*: HTTP is designed for browsers to use to ask for and receive web pages. As we move right, we are converting the message more and more into a form suitable for the **Channel**: in this case, microwave transmission.

¹⁴https://en.wikipedia.org/wiki/List_of_HTTP_status_codes

¹⁵<https://en.wikipedia.org/wiki/HTML>

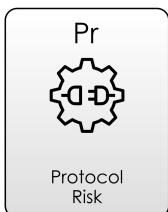
By having a stack of protocols, we are able to apply Separation Of Concerns¹⁶, each protocol handling just a few concerns:

- HTML Abstraction: A language for describing the contents of a web-page.
- HTTP Abstraction: Name-Value pairs, agreed on by both curl and Google, URLs and error codes.
- DNS Abstraction: Names of servers to IP Addresses.
- TCP Abstraction: The concept of a “connection” with guarantees about ordering and delivery.
- IP Abstraction: “Packets” with addresses and payloads.
- WiFi Abstraction: “Networks”, 802.11 flavours.
- Transmitters, Antennas, error correction codes, etc.

HTTP “stands on the shoulders of giants”. Not only does it get to use pre-existing protocols like TCP and DNS to make it’s life easier, it got 802.11 “for free” when this came along and plugged into the existing IP protocol. This is the key value of abstraction: you get to piggy-back on *existing* patterns, and use them yourself.

The protocol mediates between the message and the channel. Where this goes wrong, we have **Protocol Risk**. This is a really common issue for IT systems, but also sometimes for human communication too.

Protocol Risk



- Risks due to the failure of encoding or decoding messages between two parties in communication.

Figure 13.8: Protocol Risk

Generally, any time where you have different parts of a system communicating with each other, and one part can change incompatibly with another you have **Protocol Risk**.

Locally, (within our own project), where we have control, we can mitigate this risk using compile-time checking (as discussed already in **Complexity Risk**), which essentially forces all clients and servers to agree on protocol. But, the wider the group that you are communicating with, the less control you have and the more chance there is of **Protocol Risk**.

Let’s look at some types of **Protocol Risk**:

¹⁶https://en.wikipedia.org/wiki/Separation_of_concerns

Protocol Incompatibility Risk

The people you find it *easiest* to communicate with are your friends and family, those closest to you. That's because you're all familiar with the same protocols. Someone from a foreign country, speaking a different language and having a different culture, will essentially have a completely incompatible protocol for spoken communication to you.

Within software, there are also competing, incompatible protocols for the same things, which is maddening when your protocol isn't supported. Although the world seems to be standardizing, there used to be *hundreds* of different image formats. Photographs often use TIFF¹⁷, RAW¹⁸ or JPEG¹⁹, whilst we also have SVG²⁰ for vector graphics, GIF²¹ for images and animations and PNG²² for other bitmap graphics.

Protocol Versioning Risk

Even when systems are talking the same protocol, there can be problems. When we have multiple, different systems owned by different parties, on their own upgrade cycles, we have **Protocol Versioning Risk**: the risk that either client or server could start talking in a version of the protocol that the other side hasn't learnt yet. There are various mitigating strategies for this. We'll look at two now: **Backwards Compatibility** and **Forwards Compatibility**.

Protocol Complexity

tbd. abstraction - virtue between two vices. Postel's law.

Backward Compatibility

Backwards Compatibility mitigates **Protocol Versioning Risk**. Quite simply, this means, supporting the old format until it falls out of use. If a server is pushing for a change in protocol it either must ensure that it is Backwards Compatible with the clients it is communicating with, or make sure they are upgraded concurrently. When building web services²³, for example, it's common practice to version all APIs so that you can manage the migration. Something like this:

- Server publishes /api/v1/something.
- Clients use /api/v1/something.
- Server publishes /api/v2/something.
- Clients start using /api/v2/something.
- Clients (eventually) stop using /api/v2/something.
- Server retires /api/v2/something API.

¹⁷<https://en.wikipedia.org/wiki/TIFF>

¹⁸https://en.wikipedia.org/wiki/Raw_image_format

¹⁹<https://en.wikipedia.org/wiki/JPEG>

²⁰https://en.wikipedia.org/wiki/Scalable_Vector_Graphics

²¹<https://en.wikipedia.org/wiki/GIF>

²²https://en.wikipedia.org/wiki/Portable_Network_Graphics

²³https://en.wikipedia.org/wiki/Web_service

Forward Compatibility

HTML and HTTP provide “graceful failure” to mitigate **Protocol Risk**: while it’s expected that all clients can parse the syntax of HTML and HTTP, it’s not necessary for them to be able to handle all of the tags, attributes and rules they see. The specification for both these standards is that if you don’t understand something, ignore it. Designing with this in mind means that old clients can always at least cope with new features, but it’s not always possible.

JavaScript *can’t* support this: because the meaning of the next instruction will often depend on the result of the previous one.

Does human language support this? To some extent! New words are added to our languages all the time. When we come across a new word, we can either ignore it, guess the meaning, ask or look it up. In this way, human language has **Forward Compatibility** features built in.

Protocol Implementation Risk

A second aspect of **Protocol Risk** exists in heterogenous computing environments, where protocols have been independently implemented based on standards. For example, there are now so many different browsers, all supporting different levels of HTTP, HTML and JavaScript that it becomes impossible to test comprehensively over all the different versions. To mitigate as much **Protocol Risk** as possible, generally we run tests in a subset of browsers, and use a lowest-common-denominator approach to choosing protocol and language features.

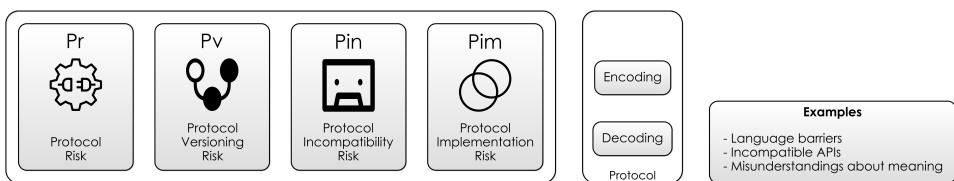


Figure 13.9: Communication Protocols Risks

Messages

Although Shannon’s Communication Theory is about transmitting **Messages**, messages are really encoded **Ideas** and **Concepts**, from an **Internal Model**.

Internal Model Assumption Risk

When we construct messages in a conversation, we have to make judgements about what the other person already knows. When talking to children, it’s often hard work because they *assume* that you have knowledge of everything they do. This is called Theory Of Mind²⁴: the

²⁴https://en.wikipedia.org/wiki/Theory_of_mind



- Risks caused by the difficulty of composing and interpreting messages in the communication process.

Figure 13.10: Message Risk

appreciation that your knowledge is different to other people's, and adjusting your messages accordingly.

When teaching, this is called The Curse Of Knowledge²⁵: teachers have difficulty understanding students' problems *because they already understand the subject*. For example, if I want to tell you about a new JDBC Driver²⁶, this pre-assumes that you know what JDBC is: the message has a dependency on prior knowledge.

Message Dependency Risk

A second, related problem is actually **Dependency Risk**, which is covered more thoroughly in the next section. Often, messages assume that you have followed everything up to that point already, otherwise again, your **Internal Model** will not be rich enough to understand the new messages.

This happens when messages get missed, or delivered out of order. In the past, TV shows were only aired once a week at a particular time. So writers were constrained plot-wise by not knowing whether their audience would have seen the previous week's episode. Therefore, often the state of the show would "reset" week-to-week, allowing you to watch it in *any* order.

The same **Message Dependency Risk** exists for computer software: if there is replication going on between instances of an application, and one of the instances misses some messages, you end up with a "Split Brain"²⁷ scenario, where later messages can't be processed because they refer to an application state that doesn't exist. For example, a message saying:

Update user 53's surname to 'Jones'

only makes sense if the application has previously had the message

Create user 53 with surname 'Smith'

²⁵https://en.wikipedia.org/wiki/Curse_of_knowledge

²⁶https://en.wikipedia.org/wiki/JDBC_driver

²⁷[https://en.wikipedia.org/wiki/Split-brain_\(computing\)](https://en.wikipedia.org/wiki/Split-brain_(computing))

Misinterpretation Risk

People don't rely on rigorous implementations of abstractions like computers do; we make do with fuzzy definitions of concepts and ideas. We rely on **Abstraction** to move between the name of a thing and the *idea of a thing*.

While machines only process *information*, people's brains run on concepts and ideas. For people, abstraction is critical: nothing exists unless we have a name for it. Our world is just atoms, but we don't think like this. *The name is the thing*.

"The famous pipe. How people reproached me for it! And yet, could you stuff my pipe? No, it's just a representation, is it not? So if I had written on my picture "This is a pipe", I'd have been lying!" - Rene Magritte, of *The Treachery of Images*²⁸

This brings about **Misinterpretation Risk**: names are not *precise*, and concepts mean different things to different people. We can't be sure that people have the same meaning for concepts that we have.

Invisibility Risk

Another cost of **Abstraction** is **Invisibility Risk**. While abstraction is a massively powerful technique, (as we saw above in the section on **Protocols**, it allows things like the Internet to happen) it lets the function of a thing hide behind the layers of abstraction and become invisible.

Invisibility Risk In Software

As soon as you create a function, you are doing abstraction. You are saying: "I now have this operation. The details, I won't mention again, but from now on, it's called f" And suddenly, "f" hides. It is working invisibly. Things go on in f that people don't necessarily need to understand. There may be some documentation, or tacit knowledge around what f is, and what it does, but it's not necessarily right. Referring to f is a much simpler job than understanding f.

We try to mitigate this via (for the most part) documentation, but this is a terrible deal: because we can't understand the original, (un-abstracted) implementation, we now need to write some simpler documentation, which *explains* the abstraction, in terms of further abstractions, and this is where things start to get murky.

Invisibility Risk is mainly **Hidden Risk**. (Mostly, *you don't know what you don't know*.) But you can carelessly *hide things from yourself* with software:

- Adding a thread to an application that doesn't report whether it's worked, failed, or is running out of control and consuming all the cycles of the CPU.
- Redundancy can increase reliability, but only if you know when servers fail, and fix them quickly. Otherwise, you only see problems when the last server fails.
- When building a webservice, can you assume that it's working for the users in the way you want it to?

²⁸https://en.wikipedia.org/wiki/The_Treachery_of_Images

When you build a software service, or even implement a thread, ask yourself: “How will I know next week that this is working properly?” If the answer involves manual work and investigation, then your implementation has just cost you in **Invisibility Risk**.

Invisibility Risk In Conversation

Invisibility Risk is risk due to information not sent. But because humans don’t need a complete understanding of a concept to use it, we can cope with some **Invisibility Risk** in communication, and this saves us time when we’re talking. It would be *painful* to have conversations if, say, the other person needed to understand everything about how cars worked in order to discuss cars.

For people, **Abstraction** is a tool that we can use to refer to other concepts, without necessarily knowing how the concepts work. This divorcing of “what” from “how” is the essence of abstraction and is what makes language useful.

The debt of **Invisibility Risk** comes due when you realise that *not* being given the details *prevents* you from reasoning about it effectively. Let’s think about this in the context of a project status meeting, for example:

- Can you be sure that the status update contains all the details you need to know?
- Is the person giving the update wrong or lying?
- Do you know enough about the details of what’s being discussed in order to make informed decisions about how the project is going?

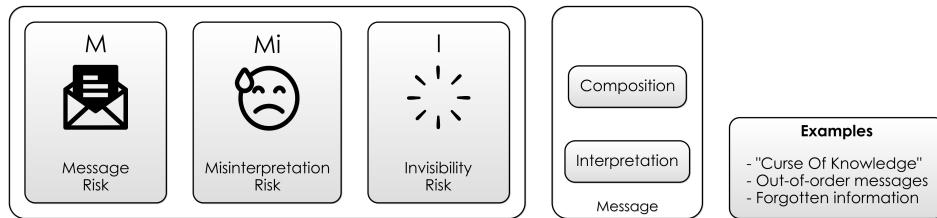
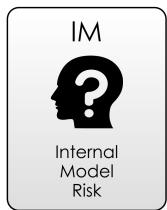


Figure 13.11: Message Risk

Internal Models

So finally, we are coming to the root of the problem: communication is about transferring ideas and concepts from one **Internal Model** to another.

The communication process so far has been fraught with risks, but we have a few more to come.



- Risks to communication caused by the fact that the internal models of the participants are semantically incompatible in some way.

Figure 13.12: Internal Model Risk



- Risk that a party we are communicating with can't be trusted, as it has agency or is unreliable in some other way.

Figure 13.13: Internal Model Risk

Trust & Belief Risk

Although protocols can sometimes handle security features of communication (such as Authentication²⁹ and preventing man-in-the-middle attacks³⁰), trust goes further than this, intersecting with **Agency Risk**: can you be sure that the other party in the communication is acting in your best interests?

Even if the receiver trusts the communicator, they may not trust the message. Let's look at some reasons for that:

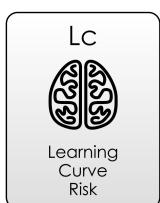
- Weltanschauung (World View)³¹: The ethics, values and beliefs in the receiver's **Internal Model** may be incompatible to those from the sender.
- Relativism³² is the concept that there are no universal truths. Every truth is from a frame of reference. For example, what constitutes *offensive language* is dependent on the listener.
- Psycholinguistics³³ is the study of humans acquire languages. There are different languages and dialects, (and *industry dialects*), and we all understand language in different ways, take different meanings and apply different contexts to the messages.

From the point-of-view of **Marketing Communications** choosing the right message is part of the battle. You are trying to communicate your idea in such a way as to mitigate **Belief Risk** and **Trust Risk**.

Reputational Risk

tbd.

Learning-Curve Risk



- Risks due to the difficulty faced in updating an internal model.

Figure 13.14: Internal Model Risk

²⁹<https://en.wikipedia.org/wiki/Authentication>

³⁰https://en.wikipedia.org/wiki/Man-in-the-middle_attack

³¹https://en.wikipedia.org/wiki/World_view

³²<https://en.wikipedia.org/wiki/Relativism>

³³<https://en.wikipedia.org/wiki/Psycholinguistics>

If the messages we are receiving force us to update our **Internal Model** too much, we can suffer from the problem of “too steep a Learning Curve³⁴” or “Information Overload³⁵”, where the messages force us to adapt our **Internal Model** too quickly for our brains to keep up.

Commonly, the easiest option is just to ignore the information channel completely in these cases.

Reading Code

It's often been said that code is *harder to read than to write*:

“If you ask a software developer what they spend their time doing, they'll tell you that they spend most of their time writing code. However, if you actually observe what software developers spend their time doing, you'll find that they spend most of their time trying to understand code.” - When Understanding Means Rewriting, *Coding Horror*³⁶

By now it should be clear that it's going to be *both* quite hard to read and write: the protocol of code is actually designed for the purpose of machines communicating, not primarily for people to understand. Making code human readable is a secondary concern to making it machine readable.

But now we should be able to see the reasons it's harder to read than write too:

- When reading code, you are having to shift your **Internal Model** to wherever the code is, accepting decisions that you might not agree with and accepting counter-intuitive logical leaps. i.e. **Learning Curve Risk**. (*cf. Principle of Least Surprise*³⁷)
- There is no **Feedback Loop** between your **Internal Model** and the **Reality** of the code, opening you up to **Misinterpretation Risk**. When you write code, your compiler and tests give you this.
- While reading code *takes less time* than writing it, this also means the **Learning Curve** is steeper.

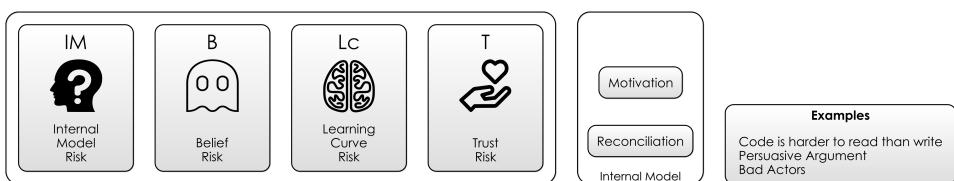


Figure 13.15: Internal Model Risks

³⁴https://en.wikipedia.org/wiki/Learning_curve

³⁵https://en.wikipedia.org/wiki/Information_overload

³⁶<https://blog.codinghorror.com/when-understanding-means-rewriting/>

³⁷https://en.wikipedia.org/wiki/Principle_of_least_astonishment

Communication Risk Wrap Up

So, here's a summary of where we've arrived with our model of communication risk:

There's no point to Communication unless you have someone or something to communicate with! So next it's time to look at **Dependency Risk**.

this seems complex tbd.

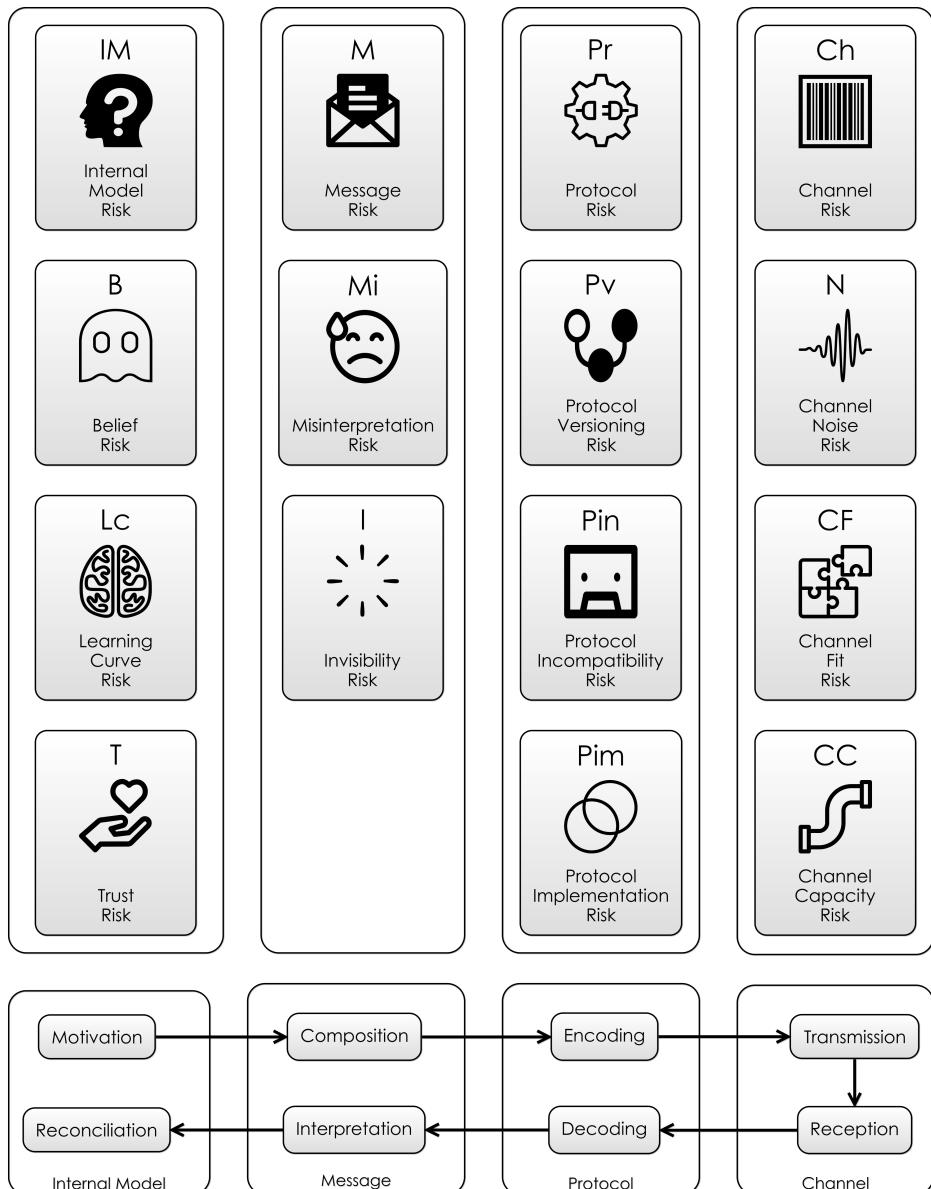


Figure 13.16: Communication 2

Chapter 14

Dependency Risk

Dependency Risk is the risk you take on whenever you have a dependency on something (or someone) else. One simple example could be that the software service you write might depend on a server to run on. If the server goes down, the service goes down too. In turn, the server depends on electricity from a supplier, as well as a network connection from a provider. If either of these dependencies aren't met, the service is out of commission.

Dependencies can be on *events, people, teams, processes, software, services, money*: pretty much *any resource*. Dependencies add risk to any project because the reliability of the project itself is now a function involving the reliability of the dependency.

In order to avoid repetition, and also to break down this large topic, we're going to look at this over 6 sections:

- In this first section will look at dependencies *in general*, and specifically on *events*, and some of the variations on **Dependency Risk**.
- Next, we'll look at **Schedule Risk**, because time and money are key dependencies in any project.
- Then, we'll move on to look specifically at **Software Dependency Risk**, covering using libraries, software services and building on top of the work of others.
- After, we'll take a look at **Process Risk**, which is still **Dependency Risk**, but we'll be considering more organisational factors and how bureaucracy comes into the picture.
- Next, we'll take a closer look at **Boundary Risk** and **Dead-End Risk**. These are the risks you face in choosing the wrong things to depend on.
- Finally, we'll wrap up this analysis with a look at some of the specific problems around working with other people or businesses in **Agency Risk**.

Why Have Dependencies?

Luckily for us, the things we depend on in life are, for the most part, abundant: water to drink, air to breathe, light, heat and most of the time, food for energy.

This isn't even lucky though: life has adapted to build dependencies on things that it can *rely* on.

Although life exists at the bottom of the ocean around hydrothermal vents¹, it is a very different kind of life to us, and has a different set of dependencies given its circumstances.

This tells us a lot about **Dependency Risk** right here:

- On the one hand, depending on something else is very often helpful, and quite often essential. (For example, all animals that *move* seem to depend on oxygen).
- However, as soon as you have dependencies, you need to take into account of their *reliability*. (Living near a river or stream gives you access to fresh water, for example).
- Successful organisms *adapt* to the dependencies available to them (like the thermal vent creatures).
- There is likely to be *competition* for a dependency when it is scarce (think of droughts and famine).

So, dependencies are a trade-off. They give with one hand and take with the other. Our modern lives are full of dependency (just think of the chains of dependency needed for putting a packet of biscuits on a supermarket shelf, for example), but we accept this extra complexity because it makes life *easier*.

Simple Made Easy

In Rich Hickey's talk, Simple Made Easy² he discusses the difference between *simple* software systems and *easy* (to use) ones, heavily stressing the virtues of simple over easy. It's an incredible talk and well worth watching.

But. Living systems are not simple. Not anymore. They evolved in the direction of increasing complexity because life was *easier* that way. In the "simpler" direction, life is first *harder* and then *impossible*, and then an evolutionary dead-end.

Depending on things makes *your job easier*. It's just division of labour³ and dependency hierarchies, as we saw in **Hierarchies and Modularization**.

Our economic system and our software systems exhibit the same tendency-towards-complexity. For example, the television in my house now is *vastly more complicated* than the one in my home when I was a child. But, it contains much more functionality and consumes much less power and space.

Event Dependencies

Let's start with dependencies on *events*.

We rely on events occurring all the time in our lives, and so this is a good place to start in our analysis of **Dependency Risk** generally. And, as we will see, all the risks that apply to events pretty much apply to all the other kinds of dependencies we'll look at.

¹https://en.wikipedia.org/wiki/Hydrothermal_vent

²<https://www.infoq.com/presentations/Simple-Made-Easy>

³https://en.wikipedia.org/wiki/Division_of_labour

Arguably, the event dependencies are the simplest to express, too: usually, a *time* and a *place*. For example: - “I can’t start shopping until the supermarket opens at 9am”, or - “I must catch my bus to work at 7:30am”.

In the first example, you can’t *start* something until a particular event happens. In the latter example, you must *be ready* for an event at a particular time.

Events Mitigate Risk...

Having an event occur in a fixed time and place is **mitigating risk**:

- By taking the bus, we are mitigating our own **Schedule Risk**: we’re (hopefully) reducing the amount of time we’re going to spend on the activity of getting to work.
- Events are a mitigation for **Coordination Risk**: A bus needn’t necessarily *have* a fixed timetable: it could wait for each passenger until they turned up, and then go. (A bit like ride-sharing works). This would be a total disaster from a **Coordination Risk** perspective, as one person could cause everyone else to be really really late. Having a fixed time for doing something mitigates **Coordination Risk** by turning it into **Schedule Risk**. Agreeing a date for a product launch, for example, allows lots of teams to coordinate their activities.
- It’s not entirely necessary to even take the bus: you could walk, or go by another form of transport. But, effectively, this just swaps one dependency for another: if you walk, this might well take longer and use more energy, so you’re just picking up **Schedule Risk** in another way. If you drive, you have a dependency on your car instead. So, there is often an *opportunity cost* with dependencies. Using the bus might be a cheap way to travel. You’re therefore imposing less **Dependency Risk** on a different scarce resource - your money.

But, Events Lead To Attendant Risk

By *deciding to use the bus* we’ve **Taken Action**.



Figure 14.1: Action Diagram showing risks mitigated by having an *event*

However, as we saw in **A Simple Scenario**, this means we pick up **Attendant Risks**.

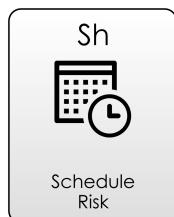
So, we’re going to look at **Dependency Risk** for our toy events (bus, supermarket) from 7 different perspectives, many of which we’ve already touched on in the other sections.

- **Schedule Risk**
- **Reliability Risk**
- **Scarcity Risk**
- **Communication Risk**
- **Complexity Risk**
- **Feature Fit Risk**
- **Dead-End Risk and Boundary Risk**

(Although you might be able to think of a few more.)

Let's look at each of these in turn.

Schedule Risk



- The aspect of dependency risk related to time.

Figure 14.2: Schedule Risk

By agreeing a *time* and *place* for something to happen, you're introducing **Deadline Risk**. Miss the deadline, and you miss the bus, or the start of the meeting or get fined for not filling your tax return on time.

As discussed above, *schedules* (such as bus timetables) exist so that *two or more parties can coordinate*, and **Deadline Risk** is on *all* of the parties. While there's a risk I am late, there's also a risk the bus is late. I might miss the start of a concert, or the band might keep everyone waiting.

Each party can mitigate **Deadline Risk** with *slack*. That is, ensuring that the exact time of the event isn't critical to your plans:

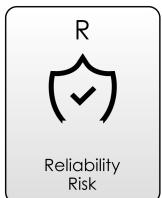
- Don't build into your plans a *need* to start shopping at 9am.
- Arrive at the bus-stop *early*.

The amount of slack you build into the schedule is likely dependent on the level of risk you face: I tend to arrive a few minutes early for a bus, because the risk is *low* (there'll be another bus along soon), however I try to arrive over an hour early for a flight, because I can't simply get on the next flight straight away, and I've already paid for it, so the risk is *high*.

Deadline Risk becomes very hard to manage when you have to coordinate actions with lots of tightly-constrained events. So what else can give? We can reduce the number of *parties*

involved in the event, which reduces risk, or, we can make sure all the parties are in the same place to begin with.

Reliability Risk



- Risks of not getting benefit from a dependency due to its reliability.

Figure 14.3: Reliability Risk

Deadline Risk is really a kind of reliability issue: if you can understand which parties are unreliable, you have a much better handle on your **Deadline Risk**.

Luckily, there is quite a lot of existing science around reliability. For example:

- If a component A depends on component B, unless there is some extra redundancy around B, then A can't be more reliable than B.
- Is A or B a Single Point Of Failure⁴ in a system?
- Are there bugs in B that are going to prevent it working correctly in all circumstances?

This kind of stuff is encapsulated in the science of Reliability Engineering⁵. For example, Failure mode and effects analysis (FMEA)⁶:

"...was one of the first highly structured, systematic techniques for failure analysis. It was developed by reliability engineers in the late 1950s to study problems that might arise from malfunctions of military systems." - FEMA, Wikipedia⁷

This was applied on NASA missions, and then more recently in the 1970's to car design following the Ford Pinto exploding car⁸ affair.

Scarcity Risk

Let's get back to the bus (which, hopefully, is still working). What if, when it arrives, it's already full of passengers? Let's term this, **Scarcity Risk** - the chance that a dependency is

⁴https://en.wikipedia.org/wiki/Single_point_of_failure

⁵https://en.wikipedia.org/wiki/Reliability_engineering

⁶https://en.wikipedia.org/wiki/Failure_mode_and_effects_analysis

⁷https://en.wikipedia.org/wiki/Failure_mode_and_effects_analysis

⁸https://en.wikipedia.org/wiki/Ford_Pinto#Design_flaws_and_ensuing_lawsuits



- Risk of not being able to access a dependency in a timely fashion due to its scarcity.

Figure 14.4: Scarcity Risk

over-subscribed and you can't use it the way you want. This is clearly an issue for nearly every kind of dependency: buses, supermarkets, concerts, teams, services and people.

You could also call this *availability risk* or *capacity risk* of the resource. Here are a selection of mitigations:

- **Buffers:** Smoothing out peaks and troughs in utilisation.
- **Reservation Systems:** giving clients information *ahead* of the dependency usage about whether the resource will be available to them.
- **Graceful degradation:** Ensuring *some* service in the event of over-subscription. It would be no use allowing people to cram onto the bus until it can't move.
- **Demand Management:** Having different prices during busy periods helps to reduce demand. Having "first class" seats means that higher-paying clients can get service even when the train is full. Uber⁹ adjust prices in real-time by so-called Surge Pricing¹⁰. This is basically turning **Scarcity Risk** into a **Market Risk** problem.
- **Queues:** Again, these provide a "fair" way of dealing with scarcity by exposing some mechanism for prioritising use of the resource. Buses operate a first-come-first-served system, whereas emergency departments in hospitals triage according to need.
- **Pools:** Reserving parts of a resource for particular customers.
- **Horizontal Scaling:** allowing a scarce resource to flexibly scale according to how much demand there is. (For example, putting on extra buses when the trains are on strike, or opening extra check-outs at the supermarket.)

Much like **Reliability Risk**, there is science for it:

- Queue Theory¹¹ is all about building mathematical models of buffers, queues, pools and so forth.
- Logistics¹² is the practical organisation of the flows of materials and goods around things like Supply Chains¹³.
- And Project Management¹⁴ is in large part about ensuring the right resources are available at the right times. We'll be taking a closer look at that in the Part 3 sections on

⁹<https://www.uber.com>

¹⁰<https://www.uber.com/en-GB/drive/partner-app/how-surge-works/>

¹¹https://en.wikipedia.org/wiki/Queueing_theory

¹²<https://en.wikipedia.org/wiki/Logistics>

¹³https://en.wikipedia.org/wiki/Supply_chain

¹⁴https://en.wikipedia.org/wiki/Project_management

Communication Risk



- Risks due to the difficulty of communicating with other entities, be they people, software, processes etc.

Figure 14.5: Communication Risk

We've already looked at communication risk in a lot of depth, and we're going to go deeper still in **Software Dependency Risk**, but let's look at some general issues around communicating dependencies. In the **Communication Risk** section we looked at **Marketing Communications** and talked about the levels of awareness that you could have with dependencies. i.e.

- The concept that there is such a thing as **D** which solves my problem isn't something I'd even considered.
- I'd like to use something like **D**, but how do I find it?
- There are multiple implementations of **D**, which is the best one for the task?
- I know **D**, but I can't figure out how to solve my problem in it.

Let's apply this to our Bus scenario:

- Am I aware that there is public transport in my area?
- How do I find out about the different options?
- How do I choose between buses, taxis, cars etc.
- How do I understand the timetable, and apply it to my problem?

Silo Mentality

Finding out about bus schedules is easy. But in a large company, **Communication Risk** and especially **Invisibility Risk** are huge problems. This tends to get called "Silo Mentality"¹⁵, that is, ignoring what else is going on in other divisions of the company or "not invented here"¹⁶ syndrome:

¹⁵https://en.wikipedia.org/wiki/Information_silo#Silo_mentality

¹⁶https://en.wikipedia.org/wiki/Not_invented_here

"In management the term silo mentality often refers to information silos in organizations. Silo mentality is caused by divergent goals of different organizational units." - Silo Mentality, *Wikipedia*¹⁷

Ironically, *more communication* might not be the answer - if channels are provided to discover functionality in other teams you can still run into **Trust Risk** (why should I believe in the quality of this dependency?) Or **Channel Risk** in terms of too low a signal-to-noise ratio, or desperate self-promotion.

Silo Mentality is exacerbated by the problems you would face in *budgeting* if suddenly all the divisions in an organisation started providing dependencies for each other. This starts to require a change to organisational structure towards being a set of individual businesses marketing services to one another, rather than a division-based one. We'll look more closely at these kind of organisational issues in the **Coordination Risk** section.

Complexity Risk

Dependencies are usually a mitigation for **Complexity Risk**, and we'll investigate that in much more detail in **Software Dependency Risk**. The reason for this is that a dependency gives you an **abstraction**: you no longer need to know *how* to do something, (that's the job of the dependency), you just need to interact with the dependency properly to get the job done. Buses are *perfect* for people who can't drive, after all.

But this means that all of the issues of abstractions that we covered in **Communication Risk** apply:

- There is **Invisibility Risk** because you probably don't have a full view of what the dependency is doing. Nowadays, bus stops have a digital "arrivals" board which gives you details of when the bus will arrive, and shops publish their opening hours online. But, abstraction always means the loss of some detail.
- There is **Misinterpretation Risk**, because often the dependency might mistake your instructions. This is endemic in software, where it's nearly impossible to describe exactly what you want up-front.

Fit Risk

Sometimes, the bus will take you to lots of in-between places you *didn't* want to go. This is **Fit Risk** and we saw this already in the section on **Feature Risk**. There, we considered two problems:

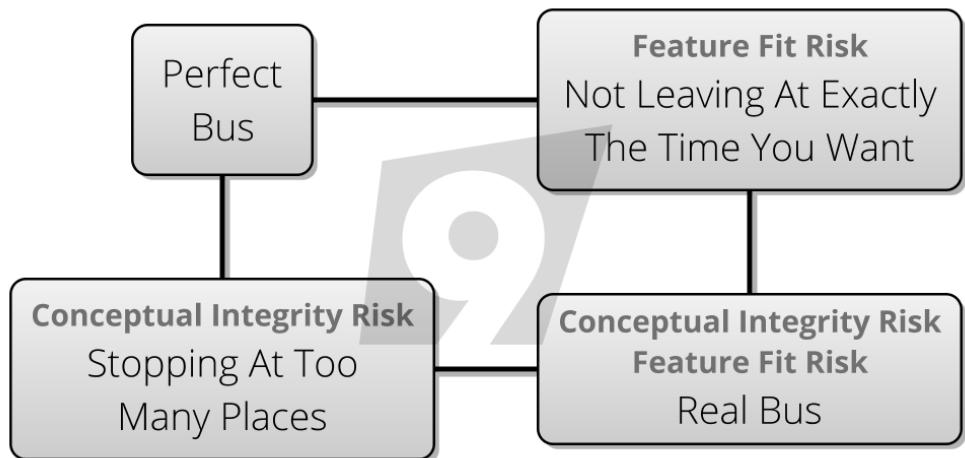
- The feature (or now, dependency) doesn't provide all the functionality you need. This was **Fit Risk**. An example might be the supermarket not stocking everything you wanted.
- The feature / dependency provides far too much, and you have to accept more complexity than you need. This was **Conceptual Integrity Risk**. An example of this might be the supermarket being *too big*, and you spend a lot longer navigating it than you wanted to.

¹⁷https://en.wikipedia.org/wiki/Information_silo#Silo_mentality



- Risks you face by not providing features that fit the needs of your clients.

Figure 14.6: Feature Fit Risk



diagrams rendered by kite9.com

Figure 14.7: Feature Fit: A Two-Dimensional Problem (at least)

Dead-End Risk and Boundary Risk

When you choose something to depend on, you can't be certain that it's going to work out in your favour. Sometimes, the path from your starting point to your goal on the **Risk Landscape** will take you to dead ends: places where the only way towards your destination is to lose something, and do it again another way. This is **Dead End Risk**, which we looked at before.

Boundary Risk is another feature of the **Risk Landscape**: when you make a decision to use one dependency over another, you are picking a path on the risk landscape that *precludes* other choices. After all, there's not really much cost in a **Dead End** if you've not had to follow a path to get to it.

We're also going to look at **Boundary Risk** in more detail later, but I want to introduce it here. Here are some examples:

- If I choose to program some software in Java, I will find it hard to integrate libraries from other languages. The dependencies available to Java software are different to those in Javascript, or C#. Having gone down a Java route, there are *higher risks* associated with choosing incompatible technologies. Yes, I can pick dependencies that use C# (still), but I am now facing greater complexity risk than if I'd just chosen to do everything in C# in the first place.
- If I choose one database over another, I am *limited to the features of that database*. This is not the same as a dead-end: I can probably build what I want to build, but the solution will be "bounded" by the dependency choices I make. One of the reasons we have standards like Java Database Connectivity (JDBC)¹⁸ is to mitigate **Dead End Risk** around databases, so that we can move to a different database later.
- If I choose to buy a bus ticket, I've made a decision not to travel by train, even though later on it might turn out that the train was a better option. Buying the bus ticket is **Boundary Risk**: I may be able to get a refund, but having chosen the dependency I've set down a path on the risk landscape.

Managing Dependency Risk

Arguably, managing **Dependency Risk** is *what Project Managers do*. Their job is to meet the **Goal** by organising the available dependencies into some kind of useful order.

There are *some* tools for managing dependency risk: Gantt Charts¹⁹ for example, arrange work according to the capacity of the resources (i.e. dependencies) available, but also the *dependencies between the tasks*. If task **B** requires the outputs of task **A**, then clearly task **A** comes first and task **B** starts after it finishes. We'll look at this more in **Process Risk**.

We'll look in more detail at project management in the *practices* part, later. But now let's get into the specifics with **Schedule Risk**.

¹⁸https://en.wikipedia.org/wiki/Java_Database_Connectivity

¹⁹https://en.wikipedia.org/wiki/Gantt_chart

Chapter 15

Schedule Risk

Schedule Risk is the term for risks you face because of *lack of time*.

You could also call this “Chronological Risk” or just “Time Risk” if you wanted to.



- The aspect of dependency risk related to time.

Figure 15.1: Schedule Risk

Schedule Risk is very pervasive, and really underlies *everything* we do. People *want* things, but they *want them at a certain time*. We need to eat and drink every day, for example. We might value having a great meal, but not if we have to wait three weeks for it.

And let's go completely philosophical for a second: Were you to attain immortality, you'd probably not feel the need to buy *anything*. You'd clearly have no *needs*, and anything you wanted, you could create yourself within your infinite time-budget. Rocks don't need money, after all.

Let's look at some specific kinds of **Schedule Risk**.



- Risk that a particular set of market conditions

Figure 15.2: Opportunity Risk

Opportunity Risk

Opportunity Risk is really the concern that whatever we do, we have to do it *in time*. If we wait too long, we'll miss the Window Of Opportunity¹ for our product or service.

Any product idea is necessarily of its time: the **Goal In Mind** will be based on observations from a particular **Internal Model**, reflecting a view on reality at a specific *point in time*.

How long will that remain true for? This is your *opportunity*: it exists apart from any deadlines you set yourself, or funding options. It's purely, "how long will this idea be worth doing?"

With any luck, decisions around *funding* your project will be tied into this, but it's not always the case. It's very easy to undershoot or overshoot the market completely and miss the window of opportunity.

The iPad

For example, let's look at the iPad², which was introduced in 2010 and was hugely successful.

This was not the first tablet computer. Apple had already tried to introduce the Newton³ in 1989, and Microsoft had released the Tablet PC⁴ in 1999. But somehow, they both missed the Window Of Opportunity⁵. Possibly, the window existed because Apple had changed the market with their release of the iPhone, which left people open to the idea of a tablet being "just a bigger iPhone".

But maybe now, the iPad's window is closing? We have more *wearable computers* like the Apple Watch⁶, and voice-controlled devices like Alexa⁷ or Siri⁸. Peak iPad was in 2014, according to this graph⁹.

¹https://en.wikipedia.org/wiki/Window_of_opportunity

²https://en.wikipedia.org/wiki/History_of_tablet_computers

³https://en.wikipedia.org/wiki/Apple_Newton

⁴https://en.wikipedia.org/wiki/Microsoft_Tablet_PC

⁵https://en.wikipedia.org/wiki/Window_of_opportunity

⁶https://en.wikipedia.org/wiki/Apple_Watch

⁷https://en.wikipedia.org/wiki/Amazon_Alexa

⁸<https://en.wikipedia.org/wiki/Siri>

⁹<https://www.statista.com/statistics/269915/global-apple-ipad-sales-since-q3-2010/>

So, it seems Apple timed the iPad to hit the peak of the Window of Opportunity.

But, even if you time the Window Of Opportunity correctly, you might still have the rug pulled from under your feet due to a different kind of **Schedule Risk**, such as...

Deadline Risk



- Where the use of a dependency has some kind of deadline, which can be missed.

Figure 15.3: Deadline Risk

Often when running a software project, you're given a team of people and told to get something delivered by a certain date. i.e. you have an artificially-imposed **Deadline** on delivery.

What happens if you miss the deadline? It could be:

- The funding on the project runs out, and it gets cancelled.
- You have to go back to a budgeting committee, and get more money.
- The team gets replaced, because of lack of faith.

.. or something else.

Deadlines can be set by an authority in order to *sharpen focus* and reduce **Coordination Risk**. This is how we arrive at tools like SMART Objectives¹⁰ and KPI's (Key Performance Indicators)¹¹. Time scales change the way we evaluate goals, and the solutions we choose.

In JFK's quote:

"First, I believe that this nation should commit itself to achieving the goal, before this decade is out, of landing a man on the moon and returning him safely to the Earth." - John F. Kennedy, 1961

The 9-year timespan came from an authority figure (the president) and helped a huge team of people coordinate their efforts and arrive at a solution that would work within a given time-frame.

Compare with this quote:

"I love deadlines. I love the whooshing noise they make as they go by." - Douglas Adams¹²

¹⁰https://en.wikipedia.org/wiki/SMART_criteria

¹¹https://en.wikipedia.org/wiki/Performance_indicator

¹²https://en.wikipedia.org/wiki/Douglas_Adams

As a successful author, Douglas Adams *didn't really care* about the deadlines his publisher's gave him. The **Deadline Risk** was minimal for him, because the publisher wouldn't be able to give his project to someone else to complete.

Sometimes, deadlines are set in order to *coordinate work between teams*. The classic example being in a battle, to coordinate attacks. When our deadlines are for this purpose, we're heading towards **Coordination Risk** territory.

Student Syndrome

Student Syndrome¹³ is, according to Wikipedia:

"Student syndrome refers to planned procrastination, when, for example, a student will only start to apply themselves to an assignment at the last possible moment before its deadline." - Wikipedia¹⁴

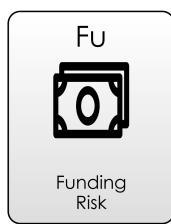
Arguably, there is good psychological, evolutionary and risk-based reasoning behind procrastination: the further in the future the **Deadline Risk** is, the more we discount it. If we're only ever mitigating our *biggest risks*, then deadlines in the future don't matter so much, do they? And, putting efforts into mitigating future risks that *might not arise* is wasted effort.

Or at least, that's the argument. If you're **Discounting the Future To Zero** then you'll be pulling all-nighters in order to deliver any assignment.

So, the problem with **Student Syndrome** is that the *very mitigation* for **Schedule Risk** (allowing more time) is an **Attendant Risk** that *causes Schedule Risk*: you'll work towards the new, generous deadline more slowly, and you'll end up revealing **Hidden Risk** later than you would have with the original, pressing deadline ... and you end up being late because of them.

We'll look at mitigations for this in Part 3's section on **Prioritisation**.

Funding Risk



- A particular scarcity risk, due to lack of funding.

Figure 15.4: Funding Risk

¹³https://en.wikipedia.org/wiki/Student_syndrome

¹⁴https://en.wikipedia.org/wiki/Student_syndrome

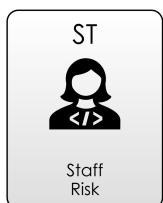
On a lot of software projects, you are “handed down” deadlines from above, and told to deliver by a certain date or face the consequences. But sometimes you’re given a budget instead, which really just adds another layer of abstraction to the **Schedule Risk**: That is, do I have enough funds to cover the team for as long as I need them?

This grants you some leeway as now you have two variables to play with: the *size* of the team, and *how long* you can run it for. The larger the team, the shorter the time you can afford to pay for it.

In startup circles, this “amount of time you can afford it” is called the “Runway”¹⁵: you have to get the product to “take-off” before the runway ends. So you could term this component as “Runway Risk”.

Startups often spend a lot of time courting investors in order to get funding and mitigate this type of **Schedule Risk**. But, this activity usually comes at the expense of **Opportunity Risk** and **Feature Risk**, as usually the same people are trying to raise funds as build the project itself.

Staff Risk



- The aspect of dependency risks related to employing people.

Figure 15.5: Staff Risk

If a startup has a “Runway”, then the chances are that the founders and staff do too, as this article explores¹⁶. It identifies the following risks:

- Company Cash: The **Runway** of the startup itself
- Founder Cash: The **Runway** for a founder, before they run out of money and can't afford their rent.
- Team Cash: The **Runway** for team members, who may not have the same appetite for risk as the founders do.

You need to consider how long your staff are going to be around, especially if you have Key Man Risk¹⁷ on some of them. People like to have new challenges, or move on to live in new places, or simply get bored. The longer your project goes on for, the more **Staff Risk** you will have to endure, and you can't rely on getting the **best staff for failing projects**.

¹⁵<https://en.wiktionary.org/wiki/runway>

¹⁶<https://www.entrepreneur.com/article/223135>

¹⁷https://en.wikipedia.org/wiki/Key_person_insurance#Key_person_definition

In the section on **Coordination-Risk** we'll look in more detail at the non-temporal components of **Staff Risk**.

Red-Queen Risk



- The general risk that the competitive environment we operate within changes over time.

Figure 15.6: Red Queen Risk

A more specific formulation of **Schedule Risk** is **Red Queen Risk**, which is that whatever you build at the start of the project will go slowly more-and-more out of date as the project goes on.

This is named after the Red Queen quote from Alice in Wonderland:

"My dear, here we must run as fast as we can, just to stay in place. And if you wish to go anywhere you must run twice as fast as that." - Lewis Carroll, *Alice in Wonderland*¹⁸

The problem with software projects is that tools and techniques change *really fast*. In 2011, 3DRealms released Duke Nukem Forever after 15 years in development¹⁹, to negative reviews:

"... most of the criticism directed towards the game's long loading times, clunky controls, offensive humor, and overall aging and dated design." - *Duke Nukem Forever*, Wikipedia²⁰

Now, they didn't *deliberately* take 15 years to build this game (lots of things went wrong). But, the longer it took, the more their existing design and code-base were a liability rather than an asset.

Personally, I have suffered the pain on project teams where we've had to cope with legacy code and databases because the cost of changing them was too high. And any team who is stuck using Visual Basic 6.0²¹ is here. It's possible to ignore **Red Queen Risk** for a time, but this is just another form of **Technical Debt** which eventually comes due.

¹⁸<https://www.goodreads.com/quotes/458856-my-dear-here-we-must-run-as-fast-as-we>

¹⁹https://en.wikipedia.org/wiki/Duke_Nukem_Forever

²⁰https://en.wikipedia.org/wiki/Duke_Nukem_Forever

²¹https://en.wikipedia.org/wiki/Visual_Basic

Schedule Risk and Feature Risk

In the section on **Feature Risk** we looked at **Market Risk**, the idea that the value of your product is itself at risk from the moves of the market, share prices being the obvious example of that effect. In Finance, we measure this using *money*, and we can put together probability models based on how much money you might make or lose.

With **Schedule Risk**, the underlying measure is *time*:

- “If I implement feature X, I’m picking up something like 5 days of **Schedule Risk**.”
- “If John goes travelling that’s going to hit us with lots of **Schedule Risk** while we train up Anne.”

... and so on. Clearly, in the same way as you don’t know exactly how much money you might lose or gain on the stock-exchange, you can’t put precise numbers on **Schedule Risk** either.

Schedule Risk, then, is *fundamental* to every dependency. But now it’s time to get into the *specifics*, and look at **Software Dependencies**.

Chapter 16

Software Dependency Risk

In this section, we're going to look specifically at *Software* dependencies, although many of the concerns we'll raise here apply equally to all the other types of dependency we outlined in **Dependency Risk**.

Kolmogorov Complexity: Cheating

In the earlier section on **Complexity Risk** we tackled **Kolmogorov Complexity**, and the idea that your codebase had some kind of minimal level of complexity based on the output it was trying to create. This is a neat idea, but in a way, we cheated. Let's look at how.

We were trying to figure out the shortest (Javascript) program to generate this output:

And we came up with this:

Which had 26 symbols in it

Now, here's the cheat: The `repeat()` function was built into Javascript in 2015 in ECMAScript 6 ¹. If we'd had to program it ourselves, we might have added this:

```
function repeat(s,n) {  
    var a=[];  
    while(a.length<n){  
        a.push(s)  
    }  
}
```

(10 symbols)
(7 symbols)
(9 symbols)
(6 symbols)
(1 symbol)

¹<http://www.ecma-international.org/ecma-262/6.0/>

```

    return a.join('');
}

```

(10 symbols)
(1 symbol)

... which would be an extra **44** symbols (in total **70**), and push us completely over the original string encoding of **53** symbols. So, *encoding language is important*.

Conversely, if ECMAScript 6.0 had introduced a function called `abcdRepeater(n)` we'd have been able to do this:

```

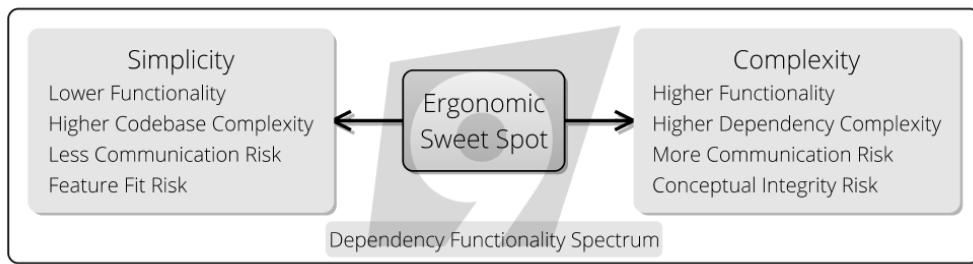
function out() {
    return abcdRepeater(10)
}

```

(7 symbols)
(6 symbols)
(1 symbol)

.. and re-encode to **14** symbols. Now, clearly there are some problems with all this:

1. Clearly, *language matters*: the Kolmogorov complexity is dependent on the language, and the features the language has built in.
2. The exact Kolmogorov complexity is uncomputable anyway (it's the *theoretical* minimum program length). It's just a fairly abstract idea, so we shouldn't get too hung up on this. There is no function to be able to say, "what's the Kolmogorov complexity of string X"
3. What is this new library function we've created? Is `abcdRepeater` going to be part of *every* Javascript? If so, then we've shifted **Codebase Risk** away from ourselves, but we've pushed **Communication Risk** and **Dependency Risk** onto every *other* user of Javascript. (Why these? Because `abcdRepeater` will be clogging up the documentation and other people will rely on it to function correctly.)
4. Are there equivalent functions for every single other string? If so, then compilation is no longer a tractable problem: is `return abcdRepeater(10)` correct code? Well, now we have a massive library of different `XXXRepeater` functions to compile against to see if it is... So, what we *lose* in **Kolmogorov Complexity** we gain in **Runtime Complexity**.
5. Language design, then, is about *ergonomics*. After you have passed the relatively low bar of providing Turing Completeness², the key is to provide *useful* features that enable problems to be solved, without over-burdening the user with features they *don't* need. And in fact, all software is about this.



diagrams rendered by kite9.com

Figure 16.1: Software Dependency Ergonomics: finding the sweet spot between too many features and too few

²https://en.wikipedia.org/wiki/Turing_completeness

Ergonomics Examined

Have a look at some physical tools, like a hammer, or spanner. To look at them, they are probably *simple* objects, obvious, strong and dependable. Their entire behaviour is encapsulated in their form. Now, if you have a drill or sander to hand, look at the design of this too. If it's well-designed, then from the outside it is simple, perhaps with only one or two controls. Inside, it is complex and contains a motor, perhaps a transformer, and is maybe made of a hundred different components.

But outside, the form is simple, and designed for humans to use. This is *ergonomics*³:

"Human factors and ergonomics (commonly referred to as Human Factors), is the application of psychological and physiological principles to the (engineering and) design of products, processes, and systems. The goal of human factors is to reduce human error, increase productivity, and enhance safety and comfort with a specific focus on the interaction between the human and the thing of interest." - Human Factors and Ergonomics, *Wikipedia*⁴

Interfaces

The interface of a tool is the part we touch and interact with. By striving for simplicity, the interface reduces **Communication Risk**.

The interface of a system expands when you ask it to do a wide variety of things. An easy-to-use drill does one thing well: it turns drill-bits at useful levels of torque for drilling holes and sinking screws. But if you wanted it to also operate as a lathe, a sander or a strimmer (all basically mechanical things going round) you would have to sacrifice the ergonomic simplicity for a more complex interface, probably including adapters, extensions, handles and so on.

So, we now have split complexity into two: - The inner complexity of the tool (how it works internally, its own **Kolmogorov Complexity**). - The complexity of the instructions that we need to write to make the tool work (the interface **Kolmogorov Complexity**).

Software Tools

In the same way as with a hand-tool, the bulk of the complexity of a software tool is hidden behind its interface. But, the more complex the *purpose* of the tool, the more complex the interface will be.

Software is not constrained by *physical* ergonomics in the same way as a tool is. But ideally, it should have conceptual ergonomics: ideally, complexity is hidden away from the user behind the Application Programming Interface (API)⁵. This is the familiar concept of **Abstraction** we've already looked at.

That is, the tool should be as simple to use and understand as possible. This is the Principal Of Least Astonishment⁶:

³https://en.wikipedia.org/wiki/Human_factors_and_ergonomics

⁴https://en.wikipedia.org/wiki/Human_factors_and_ergonomics

⁵https://en.wikipedia.org/wiki/Application_programming_interface

⁶https://en.wikipedia.org/wiki/Principle_of_least_astonishment

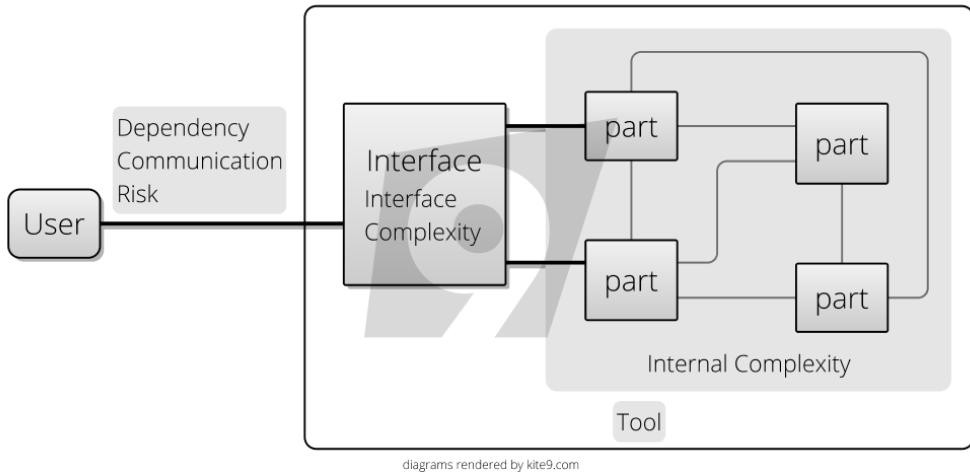


Figure 16.2: Types of Complexity For a Software Dependency

- **The abstractions should map easily to how the user expects the tool to work.** For example, I *expect* the trigger on a drill to start the drill turning.
- **The abstractions should leverage existing idioms and knowledge.** In a new car, I *expect* to know what the symbols on the dashboard mean, because I've driven other cars.
- **The abstractions provide me with only the functions I need.** Because everything else is confusing and gets in the way.

The way to win, then, is to allow a language to be extensible as-needed with features written by third parties. By supplying mechanisms for extension a language can provide insurances against the **Boundary Risk** of adopting it.

Types Of Software Dependencies

There are lots of ways you can depend on software. Here though, we're going to focus on just three main types:

1. **Code Your Own:** write some code ourselves to meet the dependency. 2. **Software Libraries:** importing code from the Internet, and using it in our project. Often, libraries are Open Source (this is what we'll consider here). 3. **Software as a Service:** calling a service on the Internet, (probably via `http`) This is often known as SaaS, or Software as a Service⁷.

All 3 approaches involve a different risk-profile. Let's look at each in turn, from the perspective of which risks get mitigated, and which risks are accentuated.

⁷https://en.wikipedia.org/wiki/Software_as_a_service

1. Code Your Own

Initially, writing our own code was the only game in town: when I started programming, you had a user guide, BASIC and that was pretty much it. Tool support was very thin-on-the-ground. Programs and libraries could be distributed as code snippets *in magazines* which could be transcribed and run, and added to your program. This spirit lives on somewhat in StackOverflow and JSFiddle, where you are expected to “adopt” others’ code into your own project.

One of the hidden risks of embarking on a code-your-own approach is that the features you need are *not* apparent from the outset. What might appear to be a trivial implementation of some piece of functionality can often turn into its own industry as more and more hidden **Feature Risk** is uncovered.

For example, as we discussed in our earlier treatment of **Dead-End Risk**, building log-in screens *seemed like a good idea*. However, this gets out-of-hand fast when you need: - A password reset screen - To email the reset links to the user - An email verification screen - A lost account screen - Reminders to complete the sign up process - ... and so on.

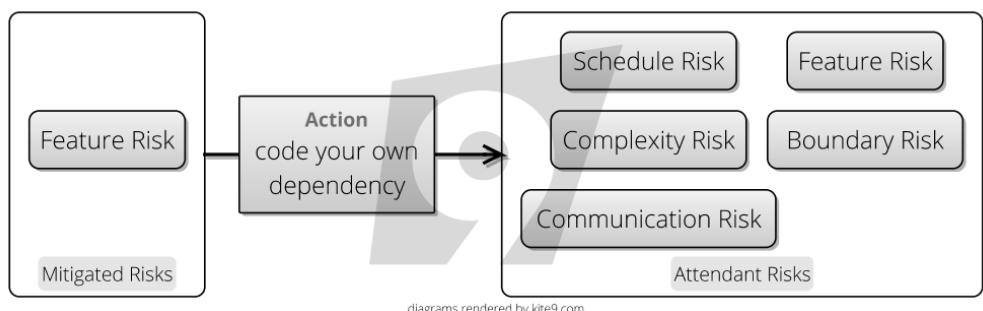


Figure 16.3: Code-Your-Own mitigates immediate feature risk, but at the expense of schedule risk, complexity risk and communication risk. There is also a hidden risk of features you don’t yet know you need.

Unwritten Software

Sometimes, you will pick up a dependency on *unwritten software*. This commonly happens when work is divided amongst team members, or teams.

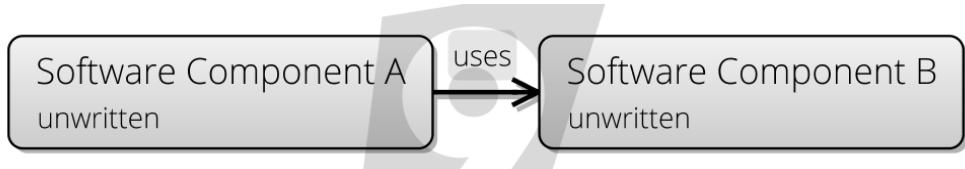


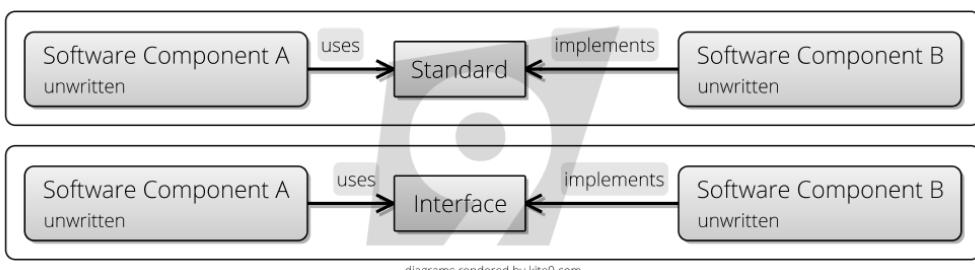
Figure 16.4: Sometimes, a module you’re writing will depend on unwritten code

If a component A of our project *depends* on B for some kind of processing, you might not be able to complete A before writing B. This makes *scheduling* the project harder, and if component A is a risky part of the project, then the chances are you'll want to mitigate risk there first.

But it also hugely increases **Communication Risk** because now you're being asked to communicate with a dependency that doesn't really exist yet, *let alone* have any documentation.

There are a couple of ways to do this:

- **Standards:** If component B is a database, a queue, mail gateway or something else with a standard interface, then you're in luck. Write A to those standards, and find a cheap, simple implementation to test with. This gives you time to sort out exactly what implementation of B you're going for. This is not a great long-term solution, because obviously, you're not using the *real* dependency- you might get surprised when the behaviour of the real component is subtly different. But it can reduce **Schedule Risk** in the short-term.
- **Coding To Interfaces:** If standards aren't an option, but the surface area of B that A uses is quite small and obvious, you can write a small interface for it, and work behind that, using a Mock⁸ for B while you're waiting for finished component. Write the interface to cover only what A *needs*, rather than everything that B *does* in order to minimize the risk of Leaky Abstractions⁹.



diagrams rendered by kite9.com

Figure 16.5: Coding to a standard on an interface breaks the dependency on unwritten software

Conway's Law

If the dependency is being written by another person, another team or in another country, communication risks pile up. When this happens, you will want to minimize *as much as possible* the interface complexity, since the more complex the interface, the worse the **Communication Risk** will be. The tendency then is to make the interfaces between teams or people *as simple as possible*, modularizing along these organisational boundaries.

In essence, this is Conway's Law¹⁰:

“organizations which design systems ... are constrained to produce designs which are copies of the communication structures of these organizations.” — M. Conway,

⁸https://en.wikipedia.org/wiki/Mock_object

⁹https://en.wikipedia.org/wiki/Leaky_abstraction

¹⁰https://en.wikipedia.org/wiki/Conway%27s_law

2. Software Libraries

By choosing a particular software library, we are making a move on the **Risk Landscape** in the hope of moving to place with more favourable risks. Typically, using library code offers a **Schedule Risk and Complexity Risk Silver Bullet**. But, in return we expect to pick up:

- **Communication Risk**: because we now have to learn how to communicate with this new dependency.
- **Boundary Risk** - because now are limited to using the functionality provided by this dependency. We have chosen it over alternatives and changing to something else would be more work and therefore costly.

But, it's quite possible that we could wind up in a worse place than we started out, by using a library that's out-of-date, riddled with bugs or badly supported. i.e. Full of new, hidden **Feature Risk**.

It's *really easy* to make bad decisions about which tools to use because the tools don't (generally) advertise their deficiencies. After all, they don't generally know how you will want to use them.

Software Libraries - Hidden Risks

Currently, choosing software dependencies looks like a "bounded rationality"-type process:

"Bounded rationality is the idea that when individuals make decisions, their rationality is limited by the tractability of the decision problem, the cognitive limitations of their minds, and the time available to make the decision." - Bounded Rationality, *Wikipedia*¹²

Unfortunately, we know that most decisions *don't* really get made this way. We have things like Confirmation Bias¹³ (looking for evidence to support a decision you've already made) and Cognitive Inertia¹⁴ (ignoring evidence that would require you to change your mind) to contend with.

But, leaving that aside, let's try to build a model of what this decision making process *should* involve. Luckily, other authors have already considered the problem of choosing good software libraries, so let's start there.

In the table below, I am summarizing three different sources, which give descriptions of which factors to look for when choosing open-source libraries.

sd1 - Defending your code against dependency problems¹⁵ sd2 - How to choose an open source library¹⁶ sd3 - Open Source - To use or not to use¹⁷

¹¹https://en.wikipedia.org/wiki/Conway%27s_law

¹²https://en.wikipedia.org/wiki/Bounded_rationality

¹³https://en.wikipedia.org/wiki/Confirmation_bias

¹⁴https://en.wikipedia.org/wiki/Cognitive_inertia

¹⁵<https://www.software.ac.uk/resources/guides/defending-your-code-against-dependency-problems>

¹⁶<https://stackoverflow.com/questions/2960371/how-to-choose-an-open-source-library>

¹⁷<https://www.forbes.com/sites/forbestechcouncil/2017/07/20/open-source-to-use-or-not-to-use-and-how-to-choose/#39e67e445a8c>

	 Coordination Risk	 Boundary Risk	 Feature Risk	 Communication Risk	Sources
Is the project "famous"?	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	[sd2] [sd3]
Is there evidence of a large user community on user forums or e-mail list archives?	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	[sd1]
Who is developing and maintaining the project? (Track Record)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	[sd3]
What are the mechanisms for supporting the software (community support, direct email, dedicated support team), and how long will the support be available? The more support, the better			<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	[sd3]
Is the API to your liking?				<input checked="" type="checkbox"/>	[sd2]
Are there examples of using the software successfully in the manner you want to use it?		<input checked="" type="checkbox"/>			[sd1]
Are all the features you need, and think you will need, included?		<input checked="" type="checkbox"/>			[sd1]
How mature is the project?		<input checked="" type="checkbox"/>			[sd2]
In respect to the software licence, do you have the right to use the software in its intended production environment, or the right to distribute it along with your software?			<input checked="" type="checkbox"/>		[sd1]
What is its deprecation or versioning policy? Does it have one? If not then it may be more unstable and features may disappear without warning between versioning, especially if releases are frequent.	<input checked="" type="checkbox"/>		 Regression Risk	<input checked="" type="checkbox"/>	[sd1]
What does the codebase look like?			 Implementation Risk		[sd1]
How frequent are its releases?		<input checked="" type="checkbox"/>			[sd1] [sd2] [sd3]
How well documented is the project?			<input checked="" type="checkbox"/>		[sd2]
Does the software have evidence of a sustainable future (e.g. is there a roadmap)?			<input checked="" type="checkbox"/>		[sd2]
Does the software support open standards? If it does, it will be easier to replace the software should it come to the end of its lifetime	<input checked="" type="checkbox"/>				[sd1]
Does the version you intend to use come from a forked open-source project, or is it from the original source project? If so, which source is more appropriate?	<input checked="" type="checkbox"/>				[sd1]
Are there any alternatives to the software?	<input checked="" type="checkbox"/>				[sd1]
Has your community converged on using a particular software package?	<input checked="" type="checkbox"/>				[sd1]
Totals	1	9	15	8	

Figure 16.6: Software Dependencies

Some take-aways:

- **Feature Risk** is a big concern. How can you be sure that the project will do what you want it to do ahead of schedule? Will it contain bugs or missing features? By looking at factors like *release frequency* and *size of the community* you get a good feel for this which is difficult to fake.
- **Boundary Risk** is also very important. You are going to have to *live* with your choices for the duration of the project, so it's worth spending the effort to either ensure that you're not going to regret the decision, or that you can change direction later.
- Third is **Communication Risk**: how well does the project deal with its users? If a project is “famous”, then it has communicated its usefulness to a wide, appreciative audience. Avoiding **Communication Risk** is also a good reason to pick *tools you are already familiar with*.

Complexity Risk?

One thing that none of the sources consider (at least from the outset) is the **Complexity Risk** of using a solution:

- Does it drag in lots of extra dependencies that seem unnecessary for the job in hand? If so, you could end up in Dependency Hell¹⁸, with multiple, conflicting versions of libraries in the project.
- Do you already have a dependency providing this functionality? So many times, I've worked on projects that import a *new* dependency when some existing (perhaps transitive) dependency has *already brought in the functionality*. For example, there are plenty of libraries for JSON¹⁹ marshalling, but if I'm also using a web framework the chances are it already has a dependency on one already.
- Does it contain lots of functionality that isn't relevant to the task you want it to accomplish? e.g. Using Java when a shell script would do (on a non-Java project)

To give an extreme example of this, I once worked on an application which used Hazelcast²⁰ to cache log-in session tokens for a 3rd party datasource. But, the app is only used once every month, and session IDs can be obtained in milliseconds. So... why cache them? Although Hazelcast is an excellent choice for in-memory caching across multiple JVMs, it is a complex piece of software (after all, it does lots of stuff). By doing this, you have introduced extra dependency risk, cache invalidation risks, networking risks, synchronisation risks and so on, for actually no benefit at all... Unless, it's about **CV Building**.

Sometimes, the amount of complexity *goes up* when you use a dependency for *good reason*. For example, in Java, you can use Java Database Connectivity (JDBC)²¹ to interface with various types of database. Spring Framework²² (a popular Java library) provides a thing called a **JDBCTemplate**. This actually makes your code *more* complex, and can prove very difficult to debug. However, it prevents some security issues, handles resource disposal and makes database access more efficient. None of those are essential to interfacing with the database, but not using them is **Technical Debt** that can bite you later on.

¹⁸https://en.wikipedia.org/wiki/Dependency_hell

¹⁹<https://en.wikipedia.org/wiki/JSON>

²⁰<https://en.wikipedia.org/wiki/Hazelcast>

²¹https://en.wikipedia.org/wiki/Java_Database_Connectivity

²²https://en.wikipedia.org/wiki/Spring_Framework

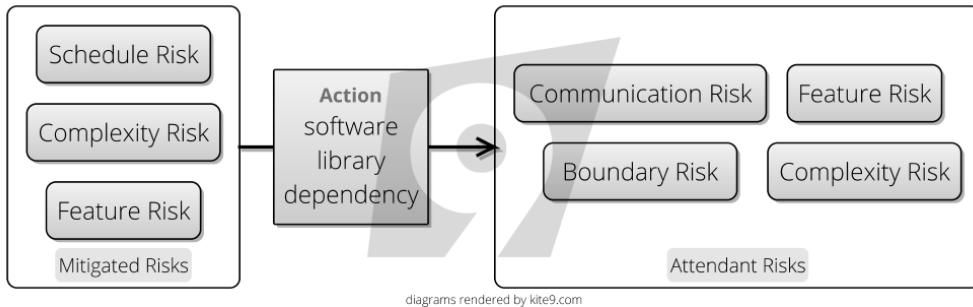


Figure 16.7: Software Libraries Risk Tradeoff

3. Software as a Service

Businesses opt for Software as a Service (SaaS) because:

- It vastly reduces the **Complexity Risk** they face in their organisations. e.g. managing the software or making changes to it.
- Payment is usually based on *usage*, mitigating **Schedule Risk**. e.g. Instead of having to pay for in-house software administrators, they can leave this function to the experts.
- Potentially, you outsource the **Operational Risk** to a third party. e.g. ensuring availability, making sure data is secure and so on.

SaaS is now a very convenient way to provide *commercial* software. Popular examples of SaaS might be SalesForce²³, or GMail²⁴. Both of which follow the commonly-used Freemium²⁵ model, where the basic service is provided free, but upgrading to a paid account gives extra benefits.

By providing the software on their own servers, the commercial organisation has a defence against *piracy*, as well as being able to control the **Complexity Risk** of the their environment (e.g. not having to support *every* version of the software that's ever been released).

Let's again recap the risks raised in some of the available literature:

sd4 - SaaS Checklist - Nine Factors to Consider²⁶ sd5 - How to Evaluate Saas Vendors²⁷

Some take-aways:

- Clearly, **Operational Risk** is now a big concern. By depending on a third-party organisation you are tying yourself to its success or failure in a much bigger way than just by using a piece of open-source software. What happens to data security, both in the data centre and over the internet?
- With **Feature Risk** you now have to contend with the fact that the software will be upgraded *outside your control*, and you may have limited control over which features get added or changed.

²³<https://en.wikipedia.org/wiki/Salesforce.com>

²⁴<https://en.wikipedia.org/wiki/Gmail>

²⁵<https://en.wikipedia.org/wiki/Freemium>

²⁶<https://www.zdnet.com/article/saas-checklist-nine-factors-to-consider-when-selecting-a-vendor/>

²⁷<http://sandhill.com/article/how-to-evaluate-saas-vendors-five-key-considerations/>

	 Operational Risk	 Boundary Risk	 Feature Risk	 Communication Risk	 Schedule Risk	Sources
How does the support process hold up in your trial runs?	<input checked="" type="checkbox"/>					[sd4]
What's the backup plan? (It's vital that you understand how your data are protected, and what redundancies are available should your SaaS provider have an outage.)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>				[sd4] [sd5]
What happens to your data if you sever ties with the vendor?	<input checked="" type="checkbox"/>					[sd4]
Are your current and future user environments supported? (e.g. Browser Compatibility)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>				[sd4]
Can you test in parallel? (i.e. run existing and new SaaS system together)			<input checked="" type="checkbox"/>			[sd4]
How does functionality compare to maturity?			<input checked="" type="checkbox"/>			[sd4]
What's the pricing model? (What might cause a price increase?)	<input checked="" type="checkbox"/>			<input checked="" type="checkbox"/>		[sd4]
What migration and training assistance options are available?				<input checked="" type="checkbox"/>		[sd4]
What integration options are available? (Are there APIs you can use to get at your data?)				<input checked="" type="checkbox"/>		[sd5]
Security (What standards and controls are in place?)	<input checked="" type="checkbox"/>					[sd5]
Service Level Agreements (SLAs) (What are the guarantees? What happens when the service levels are not met?)		<input checked="" type="checkbox"/>				[sd5]
Global Reach. (Is the service usable everywhere you need it?)	<input checked="" type="checkbox"/>					[sd5]
Totals	5	4	3	2	1	

Figure 16.8: Software As A Service Dependencies

- **Boundary Risk** is also a different proposition: you are tied to the software provider by *a contract*. If the service changes in the future, or isn't to your liking, you can't simply fork the code (like you could with an open source project).

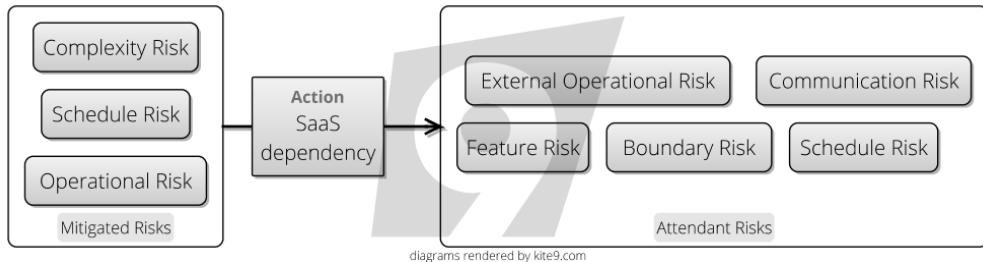


Figure 16.9: Risk Tradeoff From Using _Software as a Service (SaaS)

A Matrix of Options

We've looked at just 3 different ways of providing a software dependency: SaaS, Libraries and code-your-own.

But these are not the only ways to do it, and there's clearly no one *right* way. Although here we have looked just at "Commercial SaaS" and "Free Open Source", in reality, these are just points in a two-dimensional space involving *Pricing* and *Hosting*.

Let's expand this view slightly and look at where different pieces of software sit on these axes:

Pricing	On Premises	3rd Party	In Cloud / Browser	3rd Party	Risk Profile
---------	-------------	-----------	--------------------	-----------	--------------

| Free | OSS Libraries

Tools

Java

Firefox

Linux

Programming Languages

| Freemium

Splunk

Spotify

GitHub

| *

- Low Boundary Risk Drives AdoptionValue In Network Effect

| Advertising Supported | Commercial Software

Lots of phone apps

e.g. Angry Birds

| **Commercial SaaS** Google SearchGmail<

| Monthly / Metered Subscription | **Commercial Software**

Oracle Databases

Windows

Office

| **Commercial SaaS** Office 365SalesForce

|| | *Transferred:

Operational Risk

•

|

- Where there is value in the Network Effect²⁸, it's often a sign that the software will be free, or open source: programming languages and Linux are the obvious examples of this. Bugs are easier to find when there are lots of eyes looking, and learning the skill to use the software has less **Boundary Risk** if you know you'll be able to use it at any point in the future.
- At the other end of the spectrum, clients will happily pay for software if it clearly **reduces complexity**. Take Amazon Web Services (AWS)²⁹. The essential trade here is that you substitute the complexity of hosting and maintaining various pieces of software, in exchange for monthly payments (**Funding Risk** for you). Since the AWS *interfaces* are specific to Amazon, there is significant **Boundary Risk** in choosing this option.
- In the middle there are lots of **substitute options** and therefore high competition. Because of this, prices are pushed towards zero, and therefore often advertising is used to monetarize the product. Angry Birds³⁰ is a classic example: initially, it had demo and paid versions, however Rovio³¹ discovered there was much more money to be made through advertising than from the paid-for app³².

Managing Risks

So far, we've considered only how the different approaches to **Software Dependencies** change the landscape of risks we face to mitigate some **Feature Risk** or other.

But with **Software Dependencies** we can construct dependency networks to give us all kinds of features and mitigate all kinds of risk. That is, *the features we are looking for are to mitigate some kind of risk*.

²⁸https://en.wikipedia.org/wiki/Network_effect

²⁹https://en.wikipedia.org/wiki/Amazon_Web_Services

³⁰https://en.wikipedia.org/wiki/Angry_Birds

³¹https://en.wikipedia.org/wiki/Rovio_Entertainment

³²<https://www.deconstructoroffun.com/blog/2017/6/11/how-angry-birds-2-multiplied-quadrupled-revenue-in-a-year>

For example, I might start using WhatsApp³³ for example, because I want to be able to send my friends photos and text messages. However, it's likely that those same features are going to allow us to mitigate **Communication Risk** and **Coordination Risk** when we're next trying to meet up.

Let's look at some:

Risk	Examples of Software Mitigating That Risk
Coordination Risk	Calendar tools, Bug Tracking, Distributed Databases
Map-And-Territory-Risk	The Internet, generally. Excel, Google, "Big Data", Reporting tools
Schedule-Risk	Planning Software, Project Management Software
Communication-Risk	Email, Chat tools, CRM tools like SalesForce, Forums, Twitter, Protocols
Process-Risk	Reporting tools, online forms, process tracking tools
Agency-Risk	Auditing tools, transaction logs, Timesheet software, HR Software
Operational-Risk	Support tools like ZenDesk, Grafana, InfluxDB, Geneos
Feature-Risk	Every piece of software you use!

Back To Ergonomics

What's clear from this analysis is that software dependencies don't *conquer* any risk - the moves they make on the **Risk Landscape** are *subtle*. Whether or not you end up in a more favourable position risk-wise is going to depend heavily on the quality of the execution and the skill of the implementor.

In particular, *choosing* dependencies can be extremely difficult. As we discussed above, the usefulness of any tool depends on its fit for purpose, it's *ergonomics within a given context*. It's all too easy to pick a good tool for the wrong job:

"I suppose it is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail." - Abraham Maslow, *Toward a Psychology of Being*³⁴

With software dependencies, we often have to live with the decisions we make for a long time. In my experience, given the **Boundary Risks** associated with getting this wrong, not enough time is spent really thinking about this in advance.

Let's take a closer look at this problem in the next section, **Boundary Risk**.

³³<https://en.wikipedia.org/wiki/WhatsApp>

³⁴https://en.wiktionary.org/wiki/if_all_you_have_is_a_hammer,_everything_looks_like_a_nail

Chapter 17

Process Risk

Process Risk, as we will see, is the risk you take on whenever you embark on completing a process.

“Process: A process is a set of activities that interact to achieve a result.” - Process, Wikipedia¹

In the software development world (and the business world generally) processes commonly involve *forms*: If you’re filling out a form (whether on paper or on a computer) then you’re involved in a process of some sort, whether an “Account Registration” process, “Loan Application” process or “Consumer Satisfaction Survey” process. But sometimes, they involve events occurring: a build process² might start after you commit some code, for example.

The Purpose Of Process

Process exists to mitigate other kinds of risk, and for this reason, we’ll be looking at them again in **Part 3: Practices**, where we’ll look at how you can design processes to mitigate risks on your own project.

Until we get there, let’s look at some examples of how process can mitigate other risks:

- **Coordination Risk:** You can often use process to help people coordinate. For example, a Production Line³ is a process where work being done by one person is pushed to the next person when it’s done. A meeting booking process ensures that people will all attend a meeting together at the same place and time, and that a room is available for it.
- **Dependency Risk:** You can use processes to make dependencies explicit and mitigate dependency risk. For example, a process for hiring a car will make sure there is a car available at the time you need it. Alternatively, if we’re processing a loan application, we might need evidence of income or bank statements. We can push this **Dependency**

¹<https://en.wikipedia.org/wiki/Process>

²https://en.wikipedia.org/wiki/Software_build

³https://en.wikipedia.org/wiki/Production_line

Risk onto the person asking for the loan, by making it part of the process and not accepting the application until this has been provided.

- **Complexity Risk:** Working *within a process* can reduce the amount of **Complexity** you have to deal with. We accept that processes are going to slow us down, but we appreciate the reduction in risk this brings. (They allow us to trade **Complexity** for **Schedule risk**). For example, setting up a server might be complex, but filling in a form to do the job might simplify things. Clearly, the complexity hasn't gone away, but it's hidden within the process. Process therefore can provide **Abstraction**. mcdonalds. tbd
- **Operational Risk:** **Operational Risk** encompasses the risk of people *not doing their job properly*. But, by having a process, (and asking, did this person follow the process?) you can draw a distinction between a process failure and a personnel failure. For example, making a loan to a money launderer *could* be a failure of the loan agent. But, if they followed the *process*, it's a failure of the **Process** itself.

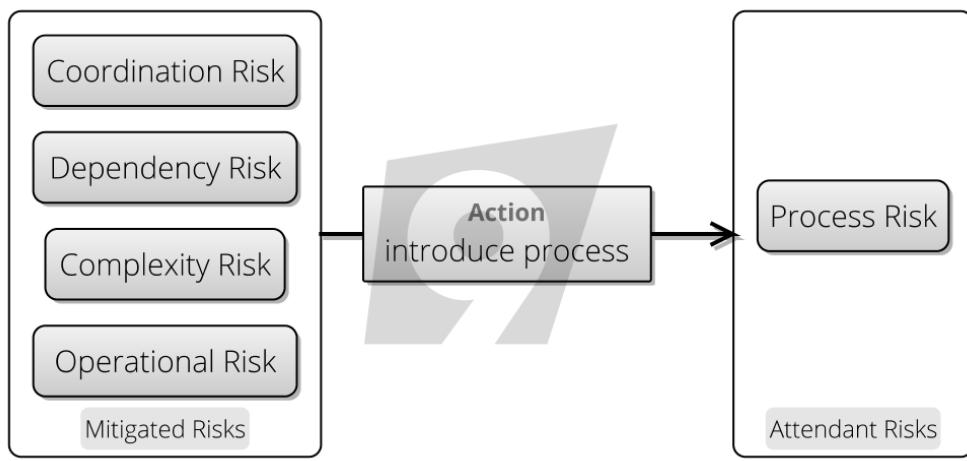


Figure 17.1: Introducing process can mitigate many risks for a team, but there are attendant process risks created.

These are all examples of **Risk Mitigation** for the owners of the process. The *consumers* of the process end up picking up most of the **Process Risks**. Let's see how this comes about.

Evolution Of Business Process

Before we get to examining what constitutes **Process Risks**, let's consider how processes *form*. Specifically, we're going to look at Business Process⁴:

"**Business Process** or **Business Method** is a collection of related, structured activities or tasks that in a specific sequence produces a service or product (serves a particular business goal) for a particular customer or customers." - Business Process, Wikipedia⁵

⁴https://en.wikipedia.org/wiki/Business_process

⁵https://en.wikipedia.org/wiki/Business_process

Business Processes often arise in response to an unmet need within an organisation. And, as we said above, they are usually there to mitigate other risks. Let's look at an example lifecycle of how that can happen.

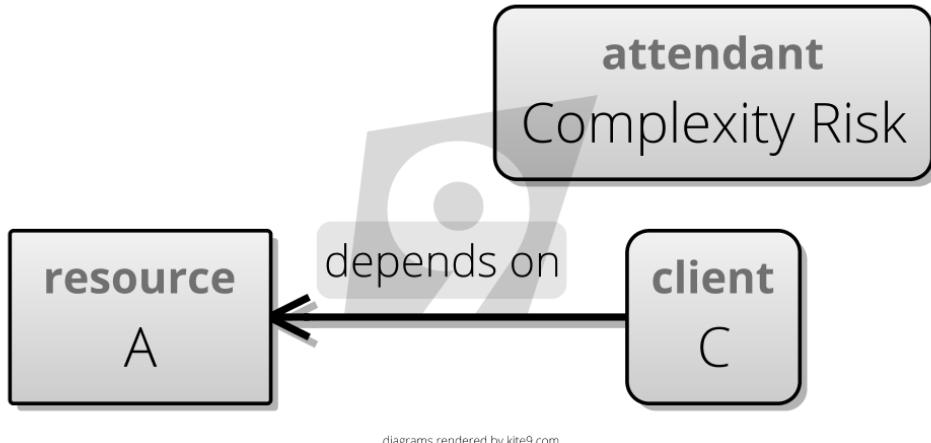
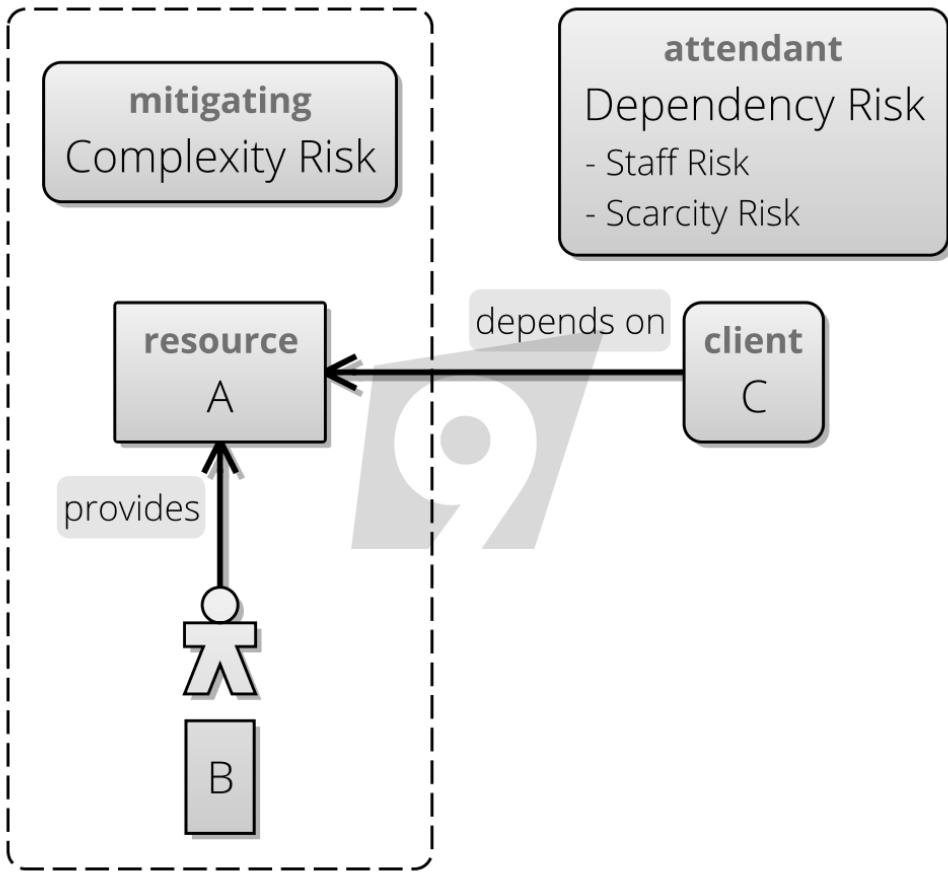


Figure 17.2: Step 0: Clients C need A to do their jobs

- o. Let's say, there exists a group of people inside a company C, which need a certain something A in order to get their jobs done. It might be a producing a resource, or dealing with some source of complexity, or whatever.
1. Person B in a company starts producing A *as a service to others*. This is really useful! It makes the the lives of clients in C much easier as they have an easier path to A than before. B gets busy keeping C happy. No one cares. But then, B goes on holiday. A doesn't get done, and people now care: the **Dependency Risk** is suddenly apparent.
2. Either, B co-opts other people to help, gets given a team (T), or someone else forms a team T containing B to get the job done "properly".
 - T is responsible for doing A, but it needs to supply the company with A reliably and responsibly, otherwise there will be trouble, so they try and please all of their clients as far as possible.
 - This is a good deal for their clients within C, but because there is a lot of variation in what the clients ask for, T end up absorbing a lot of **Complexity Risk** and are overworked.
 - This is attendant **Schedule Risk**: They either need to streamline what they are doing, or get a larger budget, because all this extra **Complexity** impacts their ability to reliably deliver A.
3. T organises bureaucratically, so that there is a controlled process (P) by which A can be accessed. Like a cell, they have arranged a protective barrier around themselves, the strength of which depends on the power conferred to them by control of A.



diagrams rendered by kite9.com

Figure 17.3: Step 1: Person B doing A for company C

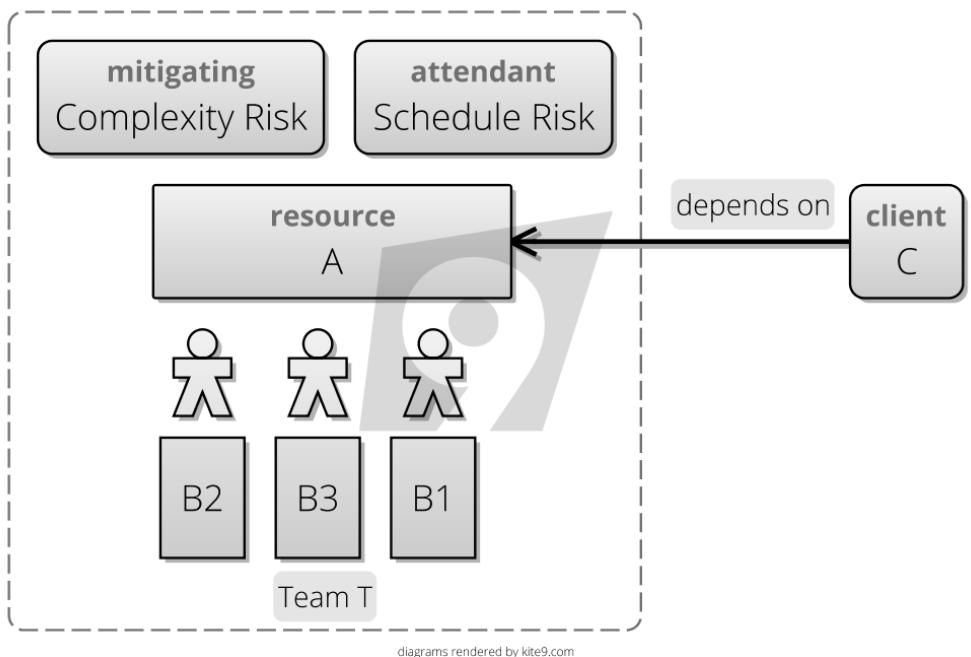


Figure 17.4: Step 2: Team T is created to do A for Company C

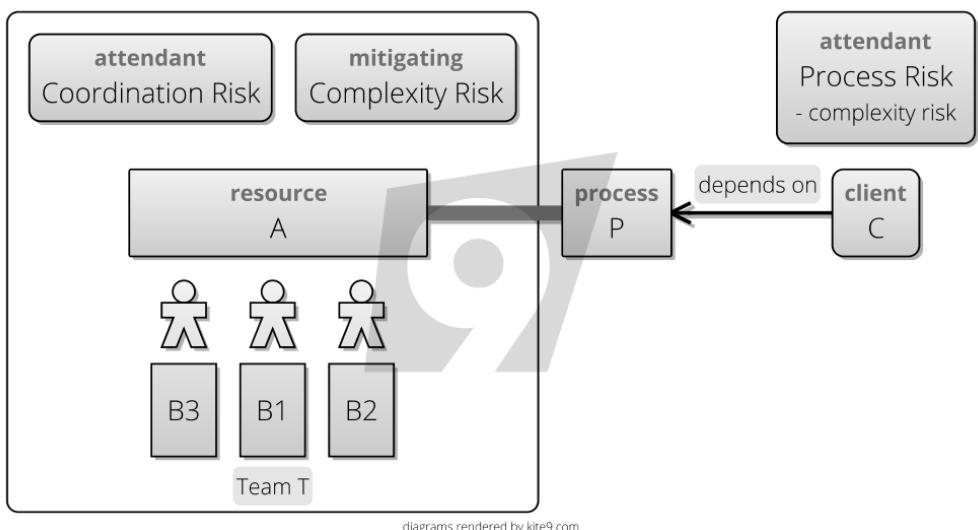


Figure 17.5: Team T protects itself from complexity with a process, P

- P probably involves filling in a form (or following some other **Protocol**).
- T can now deal with requests on a first-come-first-served basis and deal with them all in the same way: **Complexity Risks** are now the problem of the form-filler in C.
- T has mitigated **Schedule Risk** issues by drawing a line around the amount of **Complexity Risk** they are willing to take on.
- C now has **Process Risk**: will their requirements for A be met by T? They have to submit to the process to find out...

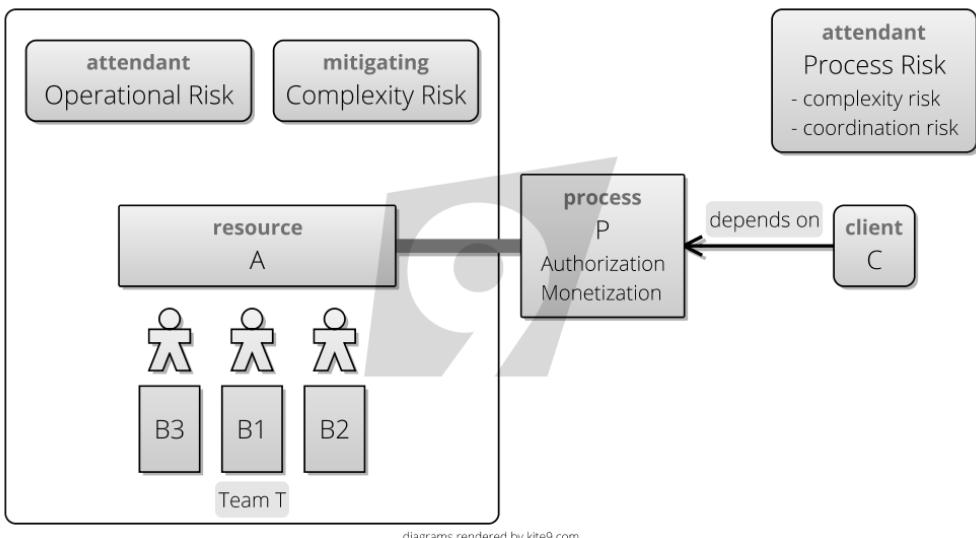


Figure 17.6: Team T protects itself from Coordination issues with signoffs or other barriers

4. But it's hard to make sure the right clients get access to A at the right times, and it's necessary to synchronize access across company C. (A **Coordination Risk** issue.)
- T reacts and sets up sign-off, authorization or monetary barriers around A, moving the **Coordination Risk** issue out of their team.
- But, for C, this *again* increases the **Process Risk** involved in using A.
5. But, there are abuses of A: people either misuse it, or use it too much. (These are **Operational Risks**).
 - T reacts by *increasing* the amount of *process* to use A, mitigating **Operational Risk** within their team, but...
 - This corresponds to greater **Process Risk** for clients in company C.
6. Person D, who has experience working with team T acts as a middleman for customers requiring some variant of A for a subset of C. They are able to help navigate the bureaucratic process (handle with **Process Risk**). The cycle potentially starts again: will D end

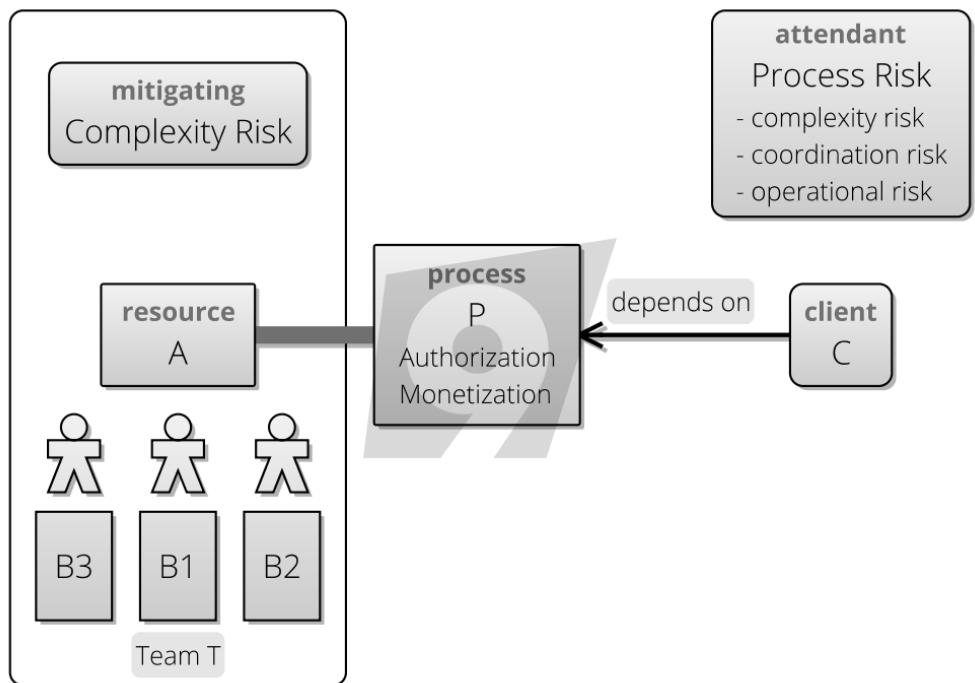


Figure 17.7: Team T increases bureaucratic load, and pushes Process Risk onto C

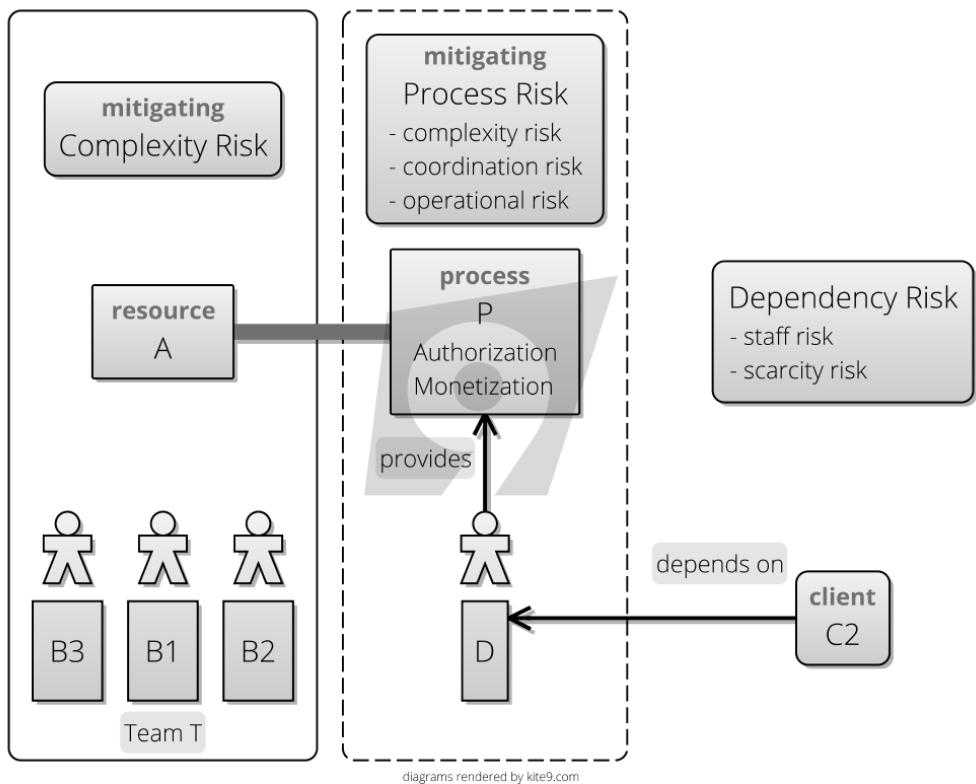


Figure 17.8: Person D acts as a middleman for customers needing some variant of A

up becoming a new team, with a new process?

In this example, you can see how the organisation evolves to mitigate risk around the use (and misuse) of A: First, **Complexity Risk**, then **Coordination Risk**, then **Dependency Risk** and finally, the **Process Risk** was created to mitigate everything else. This is an example of *Process following Strategy*:

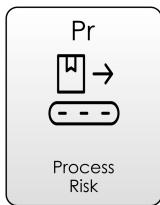
In this conception, you can see how the structure of an organisation (the teams and processes within it, the hierarchy of control) will 'evolve' from the resources of the organisation and the strategy it pursues. Processes evolve to meet the needs of the organisation." - Minzberg, *Strategy Safari*⁶

An Example - Release Process

For many years I have worked in the Finance Industry, and it's given me time to observe how, across an entire industry, process can evolve, both in response to regulatory pressure but also because of organisational maturity, and mitigating risks:

1. Initially, I could release software by logging onto the production accounts with a password that everyone knew, and deploy software or change data in the database.
2. The first issue with this is **bad actors**: How could you know that the numbers weren't being altered in the databases? Production auditing came in so that at least you could tell *who was changing what*, in order to point the blame later.
3. But, there was still plenty of scope for deliberate or accidental damage. I personally managed to wipe production data on one occasion by mistaking it for a development environment. Eventually, passwords were taken out of the hands of developers and you needed approval to "break glass" to get onto production.
4. Change Requests were introduced. This is another approval process which asks you to describe what you want to change in production, and why you want to change it. In most places, this was quite an onerous process, so the unintended consequence was that release cadence was reduced.
5. The change request software is generally awful, making the job of raising change requests tedious and time-consuming. Therefore, developers would *automate* the processes for release, sometimes including the process to write the change request. This allowed them to improve release cadence, at the expense of owning more code.
6. Auditors didn't like the fact that this automation existed, because effectively, that meant that developers could get access to production with the press of a button, effectively taking you back to step 1. So auditing of Change Requests had to happen.

... and so on.



- Risks due to the fact that when dealing with a dependency, we have to follow a particular protocol of communication, which may not work out the way we want.

Figure 17.9: Process Risk

Process Risks

Process Risk, then, is a type of **Dependency Risk**, where you are relying on a process. In a way, it's no different from any other kind of **Dependency Risk**. But **Process Risk** manifests itself in fairly predictable ways:

- **Reliability Risk:** If *people* are part of the process, there's the chance that they forget to follow the process itself, and miss steps or forget your request completely. The reliability is related to the amount of **Complexity Risk** the process is covering.
- **Invisibility Risk:** Usually, processes (like other dependencies) trade **Complexity Risk** for visibility: it's often not possible to see how far along a process is to completion. Sometimes, you can do this to an extent. For example, when I send a package for delivery, I can see roughly how far it's got on the tracking website. But, this is still less-than-complete information, and is a representation of reality.
- **Fit Risk:** You have to be careful to match the process to the outcome you want. Sometimes, it's easy to waste time on the wrong process.
- **Dead-End Risk:** Even if you have the right process, initiating a process has no guarantee that your efforts won't be wasted and you'll be back where you started from. The chances of this happening increase as you get further from the standard use-case for the process, and the sunk cost increases with the length of time the process takes to report back.
- **Agency Risk:** Due to Parkinson's Law, see below.
- **Operational Risk:** Where processes fail, this is often called **Operational Risk**, which we'll address further in its own section.
- **Credit Risk:** Where you pay for something to be done, but then end up without the outcome you want. Let's look at that in more detail.

Processes And Invisibility Risk

Processes tend to work well for the common cases, because *practice makes perfect*. but they are really tested when unusual situations occur. Expanding processes to deal with edge-cases incurs **Complexity Risk**, so often it's better to try and have clear boundaries of what is "in" and "out" of the process' domain.

⁶<http://www.mintzberg.org/books/strategy-safari>

Sometimes, processes are *not* used commonly. How can we rely on them anyway? Usually, the answer is to build in extra **feedback loops** anyway:

- Testing that backups work, even when no backup is needed.
- Running through a disaster recovery scenario at the weekend.
- Increasing the release cadence, so that we practice the release process more.

The feedback loops allow us to perform **Retrospectives and Reviews** to improve our processes.

Bureaucracy Risk



- Risk that a particular dependency devolves into a self-serving bureaucracy due to some kind of agency issue.

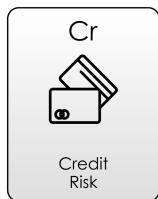
Figure 17.10: Bureaucracy Risk

Where we've talked about process evolution above, the actors involved have been acting in good faith: they are working to mitigate risk in the organisation. The **Process Risk** that concretes along the way is an *unintended consequence*: There is no guarantee that the process that arises will be humane and intuitive. Many organisational processes end up being baroque or Kafkaesque, forcing unintuitive behaviour on their users. This is partly because process design is *hard*, and it's difficult to anticipate all the various ways a process will be used ahead-of-time. So, some of **Bureaucracy Risk** is due to **Complexity**.

But Parkinson's Law⁷ takes this one step further: the human actors shaping the organisation will abuse their positions of power in order to further their own careers (this is **Agency Risk**, which we will come to in a future section):

"Parkinson's law is the adage that "work expands so as to fill the time available for its completion". It is sometimes applied to the growth of bureaucracy in an organization... He explains this growth by two forces: (1) 'An official wants to multiply subordinates, not rivals' and (2) 'Officials make work for each other.'" - Parkinson's Law, Wikipedia⁸

This implies that there is a tendency for organisations to end up with *needless levels of Bureaucratic Risk*.



- A specific version of reliability risk where the risk is due to money owed by one or other party.

Figure 17.11: Credit Risk

Credit Risk

Where the process you depend on is being run by a third-party organisation, (or that party depends on you) you are looking at **Credit Risk** (also known as **Counterparty Risk**):

"A credit risk is the risk of default on a debt that may arise from a borrower failing to make required payments... For example... A business or consumer does not pay a trade invoice when due [or] A business does not pay an employee's earned wages when due" - Credit Risk, *Wikipedia*⁹

Money is *changing hands* between you and the supplier of the process, and often, the money doesn't transfer *at the same time* as the process is performed. Let's look at an example: Instead of hosting my website on a server in my office, I could choose to host my software project with an online provider. I am trading **Complexity Risk** for **Credit Risk**, because now, I have to care that the supplier is solvent.

There's a couple of ways this could go wrong: They may *take my payment*, but then turn off my account. Or, they could go bankrupt, and leave me with the costs of moving to another provider (this is also **Dead-End Risk**).

Mechanisms like insurance¹⁰, contracts¹¹ and guarantees¹² help mitigate this risk at the cost of complexity and expense.

Sign-Offs

Often, Processes will include sign-off steps. The **Sign-Off** is an interesting mechanism:

- By signing off on something for the business, people are usually in some part staking their reputation on something being right. - Therefore, you would expect that sign-off involves a lot of **Agency Risk**: people don't want to expose themselves in career-limiting ways.
- Therefore, the bigger the risk they are being asked to swallow, the more cumbersome and protracted the sign off process.

⁷https://en.wikipedia.org/wiki/Parkinson%27s_law

⁸https://en.wikipedia.org/wiki/Parkinson%27s_law

⁹https://en.wikipedia.org/wiki/Credit_risk

¹⁰https://en.wikipedia.org/wiki/Insurance_policy

¹¹<https://en.wikipedia.org/wiki/Contract>

¹²<https://en.wikipedia.org/wiki/Guarantee>

Often, **Sign Offs** boil down to a balance of risk for the signer: on the one hand, *personal, career risk* from signing off, on the other, the risk of upsetting the rest of the staff waiting for the sign-off, and the **Dead End Risk** of all the effort gone into getting the sign off if they don't.

This is a nasty situation, but there are a couple of ways to de-risk this: - break **Sign Offs** down into bite-size chunks of risk that are acceptable to those doing the sign-off.

- Agree far-in-advance the sign-off criteria. As discussed in **Risk Theory**, people have a habit of heavily discounting future risk, and it's much easier to get agreement on the *criteria* than it is to get the sign-off.

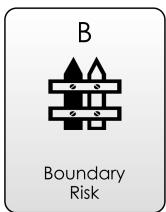
Software Processes

tbd

tbd. processes as a response to legal environment.

Chapter 18

Boundary Risk



- Risks due to the choices we make around dependencies, and the limitations they place on our ability to change.

Figure 18.1: Boundary Risk

In the previous few sections on **Dependency Risk** we've touched on **Boundary Risk** several times, but now it's time to tackle it head-on and discuss this important type of risk.

In terms of the **Risk Landscape**, **Boundary Risk** is exactly as it says: a *boundary*, *wall* or other kind of obstacle in your way to making a move you want to make. This changes the nature of the **Risk Landscape**, and introduces a maze-like component to it. It also means that we have to make *decisions* about which way to go, knowing that our future paths are constrained by the decisions we make.

And, as we discussed in **Complexity Risk**, there is always the chance we end up at a **Dead End**, and we've done work that we need to throw away. In this case, we'll have to head back and make a different decision.

Emergence Through Choice

Boundary Risk is an emergent risk, which exists at the intersection of **Complexity Risk**, **Dependency Risk** and **Communication Risk**. Because of that, it's going to take a bit of time to pick it apart and understand it, so we're going to build up to this in stages.



- The risk that a particular approach to a change will fail.
Caused by the fact that at some level, our internal models are not a complete reflection of reality.

Figure 18.2: Dead-End Risk

Let's start with an obvious example: Musical Instruments. Let's say you want to learn to play some music. There are a *multitude* of options available to you, and you might choose an *uncommon* instrument like a Balalaika¹ or a Theremin², or you might choose a *common* one like a piano or guitar. In any case, once you start learning this instrument, you have picked up the three risks above:

- **Dependency Risk** You have a *physical Dependency* on it in order to play music, so get to the music shop and buy one. - **Communication Risk**: You have to *communicate* with the instrument in order to get it to make the sounds you want. And you have **Learning Curve Risk** in order to be able to do that. - **Complexity Risk**: As a *music playing system*, you now have an extra component (the instrument), with all the attendant complexity of looking after that instrument, tuning it, and so on.

Those risks are true for *any* instrument you choose. However, if you choose the *uncommon* instrument like the Balalaika³, you have *worse Boundary Risk*, because the *ecosystem* for the balalaika is smaller. It might be hard to find a tutor, or a band needing a balalaika. You're unlikely to find one in a friend's house (compared to the piano, say).

Even choosing the Piano has **Boundary Risk**. By spending your time learning to play the piano, you're mitigating **Communication Risk** issues, but *mostly*, your skills won't be transferrable to playing the guitar. Your decision to choose one instrument over another cements the **Boundary Risk**: you're following a path on the **Risk Landscape** and changing to a different path is *expensive*.

Also, it stands to reason that making *any* choice is better than making *no* choice, because you can't try and learn *all* the instruments. Doing that, you'd make no meaningful progress on any of them.

Boundary Risk For Software Dependencies

Let's look at a software example now.

As discussed in **Software Dependency Risk**, if we are going to use a software tool as a dependency, we have to accept the complexity of its **Interface**, and learn the **protocol** of that

¹<https://en.wikipedia.org/wiki/Balalaika>

²<https://en.wikipedia.org/wiki/Theremin>

³<https://en.wikipedia.org/wiki/Balalaika>

interface. If you want to work with it, you have to use its protocol, it won't come to you.

Let's take a look at a hypothetical system structure:

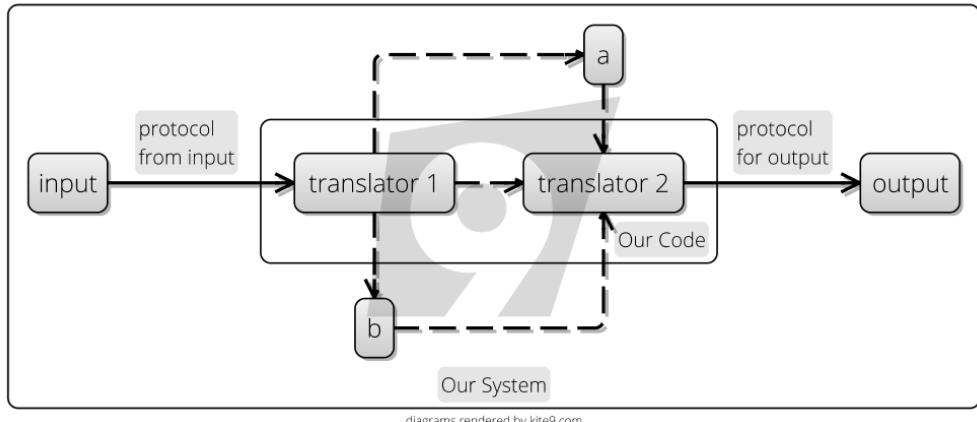


Figure 18.3: Our System receives data from the `input`, translates it and sends it to the `output`. But which dependency should we use for the translation, if any?

In this design, we have are transforming data from the `input` to the `output`. But how should we do it? - We could go via `a`, using the **Protocols** of `a`, and having a dependency on `a`. - We could go via `b`, using the **Protocols** of `b`, and having a dependency on `b`. - We could choose the middle route, and avoid the dependency, but potentially pick up lots more **Complexity Risk** and **Schedule Risk**.

This is a basic **Translation** job from `input` to `output`. Since we are talking about **Translation**, we are clearly talking about **Communication Risk** again: our task in **Integrating** all of these components is *to get them to talk to each other*.

From a **Cyclomatic Complexity** point of view, this is a very simple structure, with low **Complexity**. But the choice of approach presents us with **Boundary Risk**, because we don't know that we'll be able to make them *talk to each other* properly: - Maybe `a` outputs dates in a strange calendar format that we won't understand. - Maybe `b` works on some streaming API basis, that is incompatible with the input protocol. - Maybe `a` runs on Windows, whereas our code runs on Linux.

... and so on.

Boundary Risk Pinned Down

Wherever we integrate dependencies with complex **Protocols**, we potentially have **Boundary Risk**. The more complex the systems being integrated, the higher the risk. When we choose software tools or libraries to help us build our systems, we are trading **Complexity Risk** for **Boundary Risk**. It is:

- The *sunk cost* of the **Learning Curve** we've overcome to integrate the dependency, when it fails to live up to expectations.

- The likelihood of, and costs of changing in the future.
- The rarity of alternatives (or, conversely, the risk of **Lock In**).

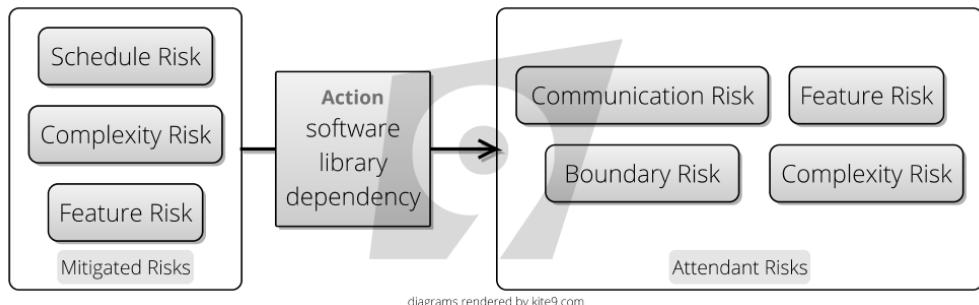


Figure 18.4: The tradeoff for using a library

As we saw in **Software Dependency Risk**, **Boundary Risk** is a big factor in choosing libraries and services. However, it can apply to any kind of dependency: - If you're depending on a **Process or Organisation**, they might change their products or quality, making the effort you put into the relationship worthless. - If you're depending on **Staff**, they might leave, meaning your efforts on training them don't pay back as well as you hoped. - If you're depending on an **Event** occurring at a particular time, you might have a lot of work to reorganise your life if it changes time or place.

Avoiding Boundary Risk Now...

Because of **Boundary Risk**'s relationship to **Learning Curve Risk**, we can avoid accreting it by choose the *simplest* and *fewest* dependencies for any job. Let's look at some examples:

- `mkdirp` is an npm⁴ module defining a single function. This function takes a single string parameter and recursively creating directories. Because the **protocol** is so simple, there is almost no **Boundary Risk**.
- Using a database with a JDBC⁵ driver comes with *some* **Boundary Risk**: but the boundary is specified by a standard. Although the standard doesn't cover every aspect of the behaviour of the database, it does minimize risk, because if you are familiar with one JDBC driver, you'll be familiar with them all, and swapping one for another is relatively easy.
- Using a framework like Spring⁶, Redux⁷ or Angular⁸ comes with higher **Boundary Risk**: you are expected to yield to the framework's way of behaving throughout your application. You cannot separate the concern easily, and swapping out the framework for another is likely to leave you with a whole new set of assumptions and interfaces to deal with.

⁴<https://www.npmjs.com>

⁵https://en.wikipedia.org/wiki/Java_Database_Connectivity

⁶<https://spring.io>

⁷<https://redux.js.org>

⁸<https://angularjs.org>

... And In The Future

Unless your project *ends*, you can never be completely sure that **Boundary Risk** *isn't* going to stop you making a move you want. For example: - `mkdirp` might not work on a new device's **Operating System**, forcing you to swap it out. - You might discover that the database you chose satisfied all the features you needed at the start of the project, but came up short when the requirements changed later on. - The front-end framework you chose might go out-of-fashion, and it might be hard to find developers interested in working on the project because of it.

This third point is perhaps the most interesting aspect of **Boundary Risk**: how can we ensure that the decisions we make now are future-proof? In order to investigate this further, let's look at three things: Plugins, Ecosystems and Evolution (again).

Plugins, Ecosystems and Evolution

On the face of it, WordPress⁹ and [Drupal](<https://en.wikipedia.org/wiki/Drupal>) *should* be very similar: - They are both Content Management Systems¹⁰ - They both use a LAMP (Linux, Apache, MySQL, PHP) Stack¹¹ - They were both started around the same time (2001 for Drupal, 2003 for WordPress) - They are both Open-Source, and have a wide variety of Plugins¹². That is, ways for other programmers to extend the functionality in new directions.

In practice, they are very different. This could be put down to different *design goals*: it seems that WordPress was focused much more on usability, and an easy learning curve, whereas Drupal supported plugins for building things with complex data formats. It could also be down to the *design decisions*: although they both support **Plugins**, they do it in very different ways.

(Side note: I wasn't short of go-to examples for this. I could have picked on Team City¹³ and Jenkins¹⁴ here (Continuous Integration¹⁵ tools), or Maven¹⁶ and Gradle¹⁷ (build tools). All of these support plugins¹⁸, and the *choice* of plugins is dependent on which I've chosen, despite the fact that the platforms are solving pretty much the same problems.)

Ecosystems and Systems

The quality, and choice of plugins for a given platform, along with factors such as community and online documentation is often called its ecosystem¹⁹:

⁹<https://en.wikipedia.org/wiki/WordPress>

¹⁰https://en.wikipedia.org/wiki/Content_management_system

¹¹[https://en.wikipedia.org/wiki/LAMP_\(software_bundle\)](https://en.wikipedia.org/wiki/LAMP_(software_bundle))

¹²[https://en.wikipedia.org/wiki/Plug-in_\(computing\)](https://en.wikipedia.org/wiki/Plug-in_(computing))

¹³<https://en.wikipedia.org/wiki/TeamCity>

¹⁴[https://en.wikipedia.org/wiki/Jenkins_\(software\)](https://en.wikipedia.org/wiki/Jenkins_(software))

¹⁵https://en.wikipedia.org/wiki/Continuous_integration

¹⁶https://en.wikipedia.org/wiki/Apache_Maven

¹⁷<https://en.wikipedia.org/wiki/Gradle>

¹⁸[https://en.wikipedia.org/wiki/Plug-in_\(computing\)](https://en.wikipedia.org/wiki/Plug-in_(computing))

¹⁹https://en.wikipedia.org/wiki/Software_ecosystem

"as a set of businesses functioning as a unit and interacting with a shared market for software and services, together with relationships among them" - Software Ecosystem, *Wikipedia*²⁰

You can think of the ecosystem as being like the footprint of a town or a city, consisting of the buildings, transport network and the people that live there. Within the city, and because of the transport network and the amenities available, it's easy to make rapid, useful moves on the **Risk Landscape**. In a software ecosystem it's the same: the ecosystem has gathered together to provide a way to mitigate various different **Feature Risks** in a common way.

tbd: talk about complexity within the boundary. (increased convenience?)

Ecosystem size is one key determinant of **Boundary Risk**: a *large* ecosystem has a large boundary circumference. **Boundary Risk** is lower because your moves on the **Risk Landscape** are unlikely to collide with it. The boundary *got large* because other developers before you hit the boundary and did the work building the software equivalents of bridges and roads and pushing it back so that the boundary didn't get in their way.

In a small ecosystem, you are much more likely to come into contact with the edges of the boundary. You will have to be the developer that pushes back the frontier and builds the roads for the others. This is hard work.

Evolution

In the real world, there is a tendency for *big cities to get bigger*. The more people that live there, the more services they provide, and therefore, the more immigrants they attract. And, it's the same in the software world. In both cases, this is due to the Network Effect²¹:

"A network effect (also called network externality or demand-side economies of scale) is the positive effect described in economics and business that an additional user of a good or service has on the value of that product to others. When a network effect is present, the value of a product or service increases according to the number of others using it." - Network Effect, *Wikipedia*²²

You can see the same effect in the adoption rates of WordPress and Drupal, shown in the chart below. Note: this is over *all sites on the internet*, so Drupal accounts for hundreds of thousands of sites. In 2018, WordPress is approximately 32% of all websites. For Drupal it's 2%.

Did WordPress gain this march because it was better than Drupal? That's arguable. That it's this way round could be *entirely accidental*, and a result of Network Effect²⁴.

And maybe, they aren't comparable: Given the same problems, the people in each ecosystem have approached them and solved them in different ways. And, this has impacted the 'shape' of the abstractions, and the protocols you use in each. **Complexity emerges**, and the ecosystem gets more complex and opinionated, much like the way in which the network of a city will evolve over time in an unpredictable way.

²⁰https://en.wikipedia.org/wiki/Software_ecosystem

²¹https://en.wikipedia.org/wiki/Network_effect

²²https://en.wikipedia.org/wiki/Network_effect

²⁴https://en.wikipedia.org/wiki/Network_effect

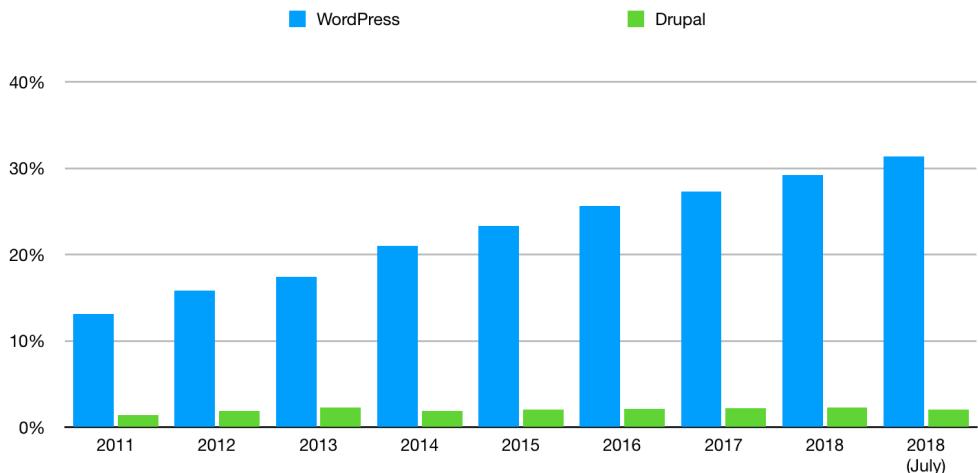


Figure 18.5: Wordpress vs Drupal adoption over 8 years, according to w3techs.com²³

But, by now, if they *are* to be compared side-by-side, WordPress *should be better* due to the sheer number of people in this ecosystem who are... - Creating web sites. - Using those sites. - Submitting bug requests. - Fixing bugs. - Writing documentation. - Building plugins. - Creating features. - Improving the core platform.

But, there are two further factors to consider...

1. The Peter Principle

When a tool or platform is popular, it is under pressure to increase in complexity. This is because people are attracted to something useful, and want to extend it to new purposes. This is known as *The Peter Principle*:

"The Peter principle is a concept in management developed by Laurence J. Peter, which observes that people in a hierarchy tend to rise to their 'level of incompetence'." - The Peter Principle, Wikipedia²⁵

Although designed for *people*, it can just as easily be applied to any other dependency you can think of. Let's look at Java²⁶ as an example of this.

Java is a very popular platform. Let's look at how the number of public classes (a good proxy for the boundary) has increased with each release:

Why does this happen?

- More and more people are using Java for more and more things. It's popularity begets more popularity.
- Human needs are *fractal* in **complexity**. You can always find ways to make a dependency *better* (For some meaning of better). - There is **Feature Drift Risk**: our requirements evolve

²⁵https://en.wikipedia.org/wiki/Peter_principle

²⁶[https://en.wikipedia.org/wiki/Java_\(software_platform\)](https://en.wikipedia.org/wiki/Java_(software_platform))

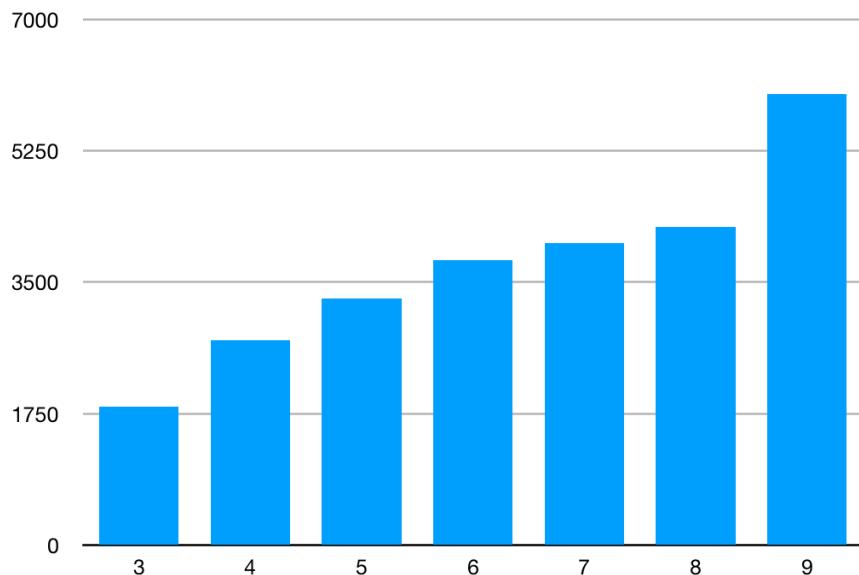


Figure 18.6: Java Public Classes By Version (3-9)

with time. Android Apps²⁷ weren't even a thing when Java 3 came out, for example, yet they are all written in Java now, and Java has had to keep up.

2. Backward Compatibility

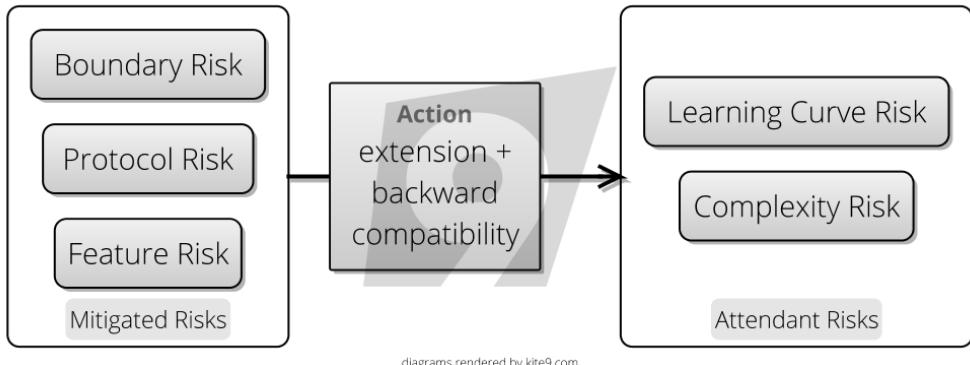
As we saw in **Software Dependency Risk**, The art of good design is to afford the greatest increase in functionality with the smallest increase in complexity possible, and this usually means **Refactoring**. But, this is at odds with **Backward Compatibility**.

Each new version has a greater functional scope than the one before (pushing back **Boundary Risk**), making the platform more attractive to build solutions in. But this increases the **Complexity Risk** as there is more functionality to deal with.

Focus vs Overreach

You can see in the diagram the Peter Principle at play: as more responsibility is given to a dependency, the more complex it gets, and the greater the learning curve to work with it. Large ecosystems like Java react to **Learning Curve Risk** by having copious amounts of literature to read or buy to help, but it is still off-putting.

²⁷https://en.wikipedia.org/wiki/Android_software_development



diagrams rendered by kite9.com

Figure 18.7: The Peter Principle: Backward Compatibility + Extension leads to complexity and learning curve risk

Because **Complexity is Mass**, large ecosystems can't respond quickly to **Feature Drift**. This means that when the world changes, *new systems* will come along to plug the gaps.

This implies a trade-off:

- Sometimes it's better to accept the **Boundary Risk** innate in a smaller system than try to work within the bigger, more complex system.

example:

In the late 80's and 90's there was a massive push towards *building functionality in the database*. **Relational Database Management Systems (RDBMSs)** were all-in-one solutions, expensive platforms that you purchased and built *everything* inside. However, this dream didn't last:

why? (need some research here).

This tbd

tbd. diagram here.

Beating Boundary Risk With Standards

Sometimes, technology comes along that allows us to cross boundaries, like a *bridge* or a *road*. This has the effect of making it easy to go from one self-contained ecosystem to another. Going back to WordPress, a simple example might be the **Analytics Dashboard** which provides Google Analytics²⁸ functionality inside WordPress.

I find, a lot of code I write is of this nature: trying to write the *glue code* to join together two different *ecosystems*.

Standards allow us to achieve the same thing, in one of two ways:

²⁸https://en.wikipedia.org/wiki/Google_Marketing_Platform

- **Mode 1: Abstract over the ecosystems.** Provide a *standard* protocol (a *lingua franca*) which can be converted down into the protocol of any of a number of competing ecosystems.
- **Mode 2: Force adoption.** All of the ecosystems start using the standard for fear of being left out in the cold. Sometimes, a standards body is involved, but other times a “de facto” standard emerges that everyone adopts.

Let's look at some examples:

- ASCII²⁹: fixed the different-character-sets boundary risk by being a standard that others could adopt. Before everyone agreed on ASCII, copying data from one computer system to another was a massive pain, and would involve some kind of translation. Unicode³⁰ continues this work. (**Mode 1**)
- C³¹: The C programming language provided a way to get the same programs compiled against different CPU instruction sets, therefore providing some *portability* to code. The problem was, each different operating system would still have its own libraries, and so to support multiple operating systems, you'd have to write code against multiple different libraries. (**Mode 2**)
- Java³² took what C did and went one step further, providing interoperability at the library level. Java code could run anywhere where Java was installed. (**Mode 2**)
- Internet Protocol³³: As we saw in **Communication Risk**, the Internet Protocol (IP) is the *lingua franca* of the modern internet. However, at one period of time, there were many competing standards. and IP was the ecosystem that “won”, and was subsequently standardized by the IETF³⁴. (**Mode 1**)

Complex Boundaries

As shown in the above diagram, mitigating Boundary Risk involves taking on complexity. The more **Protocol Complexity** there is to bridge the two ecosystems, the more **Complex** the bridge will necessarily be.

Protocol Risk ³⁵ From A	[Protocol Risk][br2] From B	Resulting Bridge Complexity	Example
Low	Low	Simple	Changing from one date format to another.
High	Low	Moderate	Status Dashboard, tbd
High	High	Complex	Object-Relational Mapping (ORM) Tools, (see below)
High + Evolving	Low	Moderate, Versioned	Simple Phone App, e.g. note-taker or calculator

²⁹<https://en.wikipedia.org/wiki/ASCII>

³⁰<https://en.wikipedia.org/wiki/Unicode>

³¹[https://en.wikipedia.org/wiki/C_\(programming_language\)](https://en.wikipedia.org/wiki/C_(programming_language))

³²[https://en.wikipedia.org/wiki/Java_\(programming_language\)](https://en.wikipedia.org/wiki/Java_(programming_language))

³³https://en.wikipedia.org/wiki/Internet_Protocol

³⁴https://en.wikipedia.org/wiki/Internet_Engineering_Task_Force

Protocol Risk ³⁵ From A	[Protocol Risk][br2] From B	Resulting Bridge Complexity	Example
Evolving	High	Complex	Modern browser (see below)
Evolving	Evolving	Very Complex	Google Search, Scala (see below)

From examining the Protocol Risk³⁶ at each end of the bridge you are creating, you can get a rough idea of how complex the endeavour will be: - If it's low-risk at both ends, you're probably going to be able to knock it out easily. Like translating a date, or converting one file format to another. - Where one of the protocols is *evolving*, you're definitely going to need to keep releasing new versions. The functionality of a `Calculator` app on my phone remains the same, but new versions have to be released as the phone APIs change, screens change resolution and so on. - tbd

Where boundaries

tbd Trying to create a complex, fractal surface. User requirements are fractal in nature.

Object-Relational Mapping

For example, **Object Relational Mapping (ORM)** has long been a problem in software. This is [Boundary-Crossing] software trying to bridge the gap between **Relational Databases** and **Object-Oriented Languages** like [Java]. Building a *general* library that does this and is useful tbd said:

'ORM is the vietnam of ...' -

This is a particularly difficult problem because the two ecosystems are so *rich* and *complex* in the functionality they expose. But what are the alternatives?

- Either back to building functionality within the database again, using stored procedures
- Building **Object Oriented Databases**. It's interesting that neither of these really worked out.
- Custom-building the bridge between the systems, one database call at-a-time in your own software.

This is tbd hobson's choice, there is strong debate about whether **ORM** is a worse trade off mitigated **Boundary Risk** for attendant **Complexity Risk** or not, and clearly will depend on your circumstances.

Scala

Mapping between complex boundaries is especially difficult if the **Boundaries** are evolving and changing as you go. This means in ecosystems that are changing rapidly, you are unlikely to be able to create lasting bridges between them. Given that **Java** is an old, large and complex

³⁵

³⁶

ecosystem, you would imagine that it would have a slow-enough rate of change that abstracting technologies can be built on top of it safely.

Indeed, we see that happening with **Clojure** and **Kotlin**, two successful languages built on top of the **Java Virtual Machine (JVM)** and offering compatibility with it.

Scala is arguably the first mainstream language that tried to do the same thing: it is trying to build a **Functional Programming** paradigm on top of the **Java Virtual Machine (JVM)**, which traditionally has an **Object Oriented** paradigm.

The problem faced by **Scala** is that Java didn't stay still: as soon as they demonstrated some really useful features in Scala (i.e. stream-based processing), Java moved to include this new functionality too. If they hadn't, the developer community would have slowly drifted away and used Scala instead.

So, in a sense, Scala is a *success story*: they were able to force change to Java. But, once Java had changed, Scala was in the difficult position of having two sets of competing features in the platform: the existing Scala streams, and the new Java streams.

Clojure can interop with Java because on one side, the boundary is simple: lisp is a simple language which lends itself to reimplementation within other platforms. Therefore, the complexity of the bridge is *simple*: all that needs to be provided is a way to call methods from Java to clojure.

Scala and Java have a complex relationship because Scala creates its own complex boundary: it is syntactically and functionally a broad language with lots of features. And so is Java. Mapping from one to the other is therefore

for interop here. Why is one so different from the other?

Browsers

Web browsers are another surprisingly complex boundary. They have to understand the following **protocols**:

- **HTTP** for loading resources (as we already reviewed in **Complexity Risk**)
- **HTML Pages**, for describing the content of web pages **Complexity Risk**
- Various image formats
- **Javascript** for web-page *interactivity*
- **CSS** for web-page styling, animation and so on.
- ... and several others.

Handling any one of these protocols alone is a massive endeavour, so browsers are built on top of **Software Libraries** which handle each concern, for example, **Networking Libraries**, **Parsers** and so on.

One way of looking at the browser is that it is a *function*, where those elements listed above are the *inputs* to the function, and the output is *what is displayed on the screen*, as shown in the image below.

tbd. browser as a function

There are three specific problems that make this a really complex boundary:

1. All of the standards above are *evolving and improving*. And, although **HTML5** (say) is a reasonably well-specified standard, in reality, web pages tend not to adhere exactly to the letter of it. People make mistakes in the HTML they write, and it's up to the browser to try and figure out what they *meant* to write, rather than what they did write. This makes the *input* to the function extremely complex.
2. Similarly, the *output* of the function is not well defined either, and relies a lot on people's *subjective aesthetic judgement*. For example, if you insert a `<table>` into an HTML page, the specification doesn't say anything about exactly how big the table should be, the size of its borders, the spacing of the content and so on. At least, initially, *none* of this was covered by the **HTML Specification**. The **CSS** specification is over time clearing this up, but it's not *exactly nailed down*, which means...
3. That because there are various different browsers (**Chrome**, **Safari**, **Internet Explorer**, **Microsoft Edge**, **Firefox** etc.) and each browser has multiple different versions, released over a period of many years, you cannot, as a web-page developer know, *a priori* what your web-page will look like to a user.

As developers trying to build software to be delivered over the internet, this is therefore a source of common **Boundary Risk**. If you were trying to build software to work in *all browsers* and *all versions*, this problem would be nearly insurmountable. So, in order to tackle this risk, we do the following:

- We pick a small (but commonly used) subset of browsers, and use features from the specifications that we know commonly work in that subset.
- We test across the subset. Again, testing is *harder than it should be*, because of problem 2 above, that the expected output is not exactly defined. This generally means you have to get humans to apply their *subjective aesthetic judgement*, rather than getting machines to do it.
- There is considerable pressure on browser developers to ensure consistency of behaviour across the implementations. If all the browsers work the same, then we don't face the **Boundary Risk** of having to choose just one to make our software work in. However, it's not always been like this...

Vendor Lock-In

In the late 1990s, faced with the emergence of the nascent **World Wide Web**, and the **Netscape Navigator** browser, **Microsoft** adopted a strategy known as **Embrace and Extend**. The idea of this was to subvert the **HTML** standard to their own ends by *embracing* the standard and creating their own browser (**Internet Explorer**) and then *extending* it with as much functionality as possible, which would then *not work* in **Netscape Navigator**. They then embarked on a campaign to try and get everyone to "upgrade" to **Internet Explorer**. In this way, they hoped to "own" the Internet, or at least, the software of the browser, which they saw as analogous to being the "operating system" of the Internet, and therefore a threat to their own operating system, **Windows**.

There are two questions we need to ask about this, from the point-of-view of understanding **Boundary Risk**:

1. Why was this a successful strategy?
2. Why did they stop doing this?

Let's look at the first question then. Yes, it was a successful strategy. In the 1990s, browser functionality was rudimentary. Developers were *desperate* for more features, and for more control over what appeared on their webpages. And, **Internet Explorer (IE)** was a free download (or, bundled with Windows). By shunning other browsers and coding just for IE, developers pushed **Boundary Risk** to the consumers of the web pages and in return mitigated **Dependency Fit Risk**: they were able to get more of the functionality they wanted in the browser.

It's worth pointing out, *this was not a new strategy*:

- Processor Chip manufacturers had done something similar in the tbds: by providing features (instructions) on their processors that other vendors didn't have, they made their processors more attractive to system integrators. However, since the instructions were different on different chips, this created **Boundary Risk** for the integrators. Intel and Microsoft were able to use this fact to build a big ecosystem around Windows running on Intel chips (so called, Wintel).
- We have two main *mobile* ecosystems: **Apple's iOS** and **Google's Android**, which are both *very* different and complex ecosystems with large, complex boundaries. They are both innovating as fast as possible to keep users happy with their features. Tools like **Xamarin** exist which allow you to build
- Currently, **Amazon Web Services (AWS)** are competing with **Microsoft Azure** and [Google tbd] over building tools for **Platform as a Service (PaaS)** (running software in the cloud). They are both racing to build new functionality, but at the same time it's hard to move from one vendor to another as there is no standardization on the tools.
- As we saw above, Database vendors tried to do the same thing with features in the database. Oracle particularly makes money over differentiating itself from competitors by providing features that other vendors don't have. Tom tbd provides a compelling argument for using these features thus:

tbd.

The next question, is why did Microsoft *stop* pursuing this strategy? It seems that the answer is because they were made to. tbd.

Everyday Boundary Risks

Boundary Risk occurs all the time. Let's look at some ways:

- **Configuration:** When software has to be deployed onto a server, there has to be configuration (usually on the command line, or via configuration property files) in order to bridge the boundary between the *environment it's running in* and the *software being run*. Often, this is setting up file locations, security keys and passwords, and telling it where to find other files and services.
- **Integration Testing:** Building a unit test is easy. You are generally testing some code you have written, aided with a testing framework. Your code and the framework are both written in the same language, which means low boundary risk. But, to *integration*

test you need to step outside this boundary and so it becomes much harder. This is true whether you are integrating with other systems (providing or supplying them with data) or parts of your own system (say testing the client-side and server parts together).

- **User Interface Testing:** If you are supplying a user-interface, then the interface with the user is already a complex, under-specified risky **protocol**. Although tools exist to automate UI testing (such as **Selenium**, these rarely satisfactorily mitigate this **protocol risk**: can you be sure that the screen hasn't got strange glitches, that the mouse moves correctly, that the proportions on the screen are correct on all browsers?)
- **Jobs:** When you pick a new technology to learn and add to your CV, it's worth keeping in mind how useful this will be to you in the future. It's career-limiting to be stuck in a dying ecosystem and need to retrain.
- **Teams:** if you're given license to build a new product within an existing team, are you creating **Boundary Risk** by using tools that the team aren't familiar with?
- **Organisatations:** Getting teams or departments to work with each other often involves breaking down **Boundary Risk**. Often the departments use different tool-sets or processes, and have different goals making the translation harder. tbd

Boundary Risk and Change

You can't always be sure that a dependency now will always have the same guarantees in the future: - **Ownership changes** Microsoft buys Github. What will happen to the ecosystem around github now? - **Licensing changes**. (e.g. Oracle³⁷ buys **Tangosol** who make Coherence³⁸ for example). Having done this, they increase the licensing costs of Tangosol to huge levels, milking the **Cash Cow** of the installed user-base, but ensuring no-one else is likely to use it. - **Better alternatives become available**: As a real example of this, I began a project in 2016 using **Apache Solr**. However, in 2018, I would probably use ElasticSearch³⁹. In the past, I've built websites using Drupal and then later converted them to use WordPress.

Patterns In Boundary Risk

In **Feature Risk**, we saw that the features people need change over time. Let's get more specific about this:

- Human need is **Fractal**. This means that over time, software products have evolved to more closely map to human needs. Software that would have delighted us ten years ago lacks the sophistication we expect today.
- Software and hardware are both improving with time, due to evolution and the ability to support greater and greater levels of complexity.
- Abstractions build too. As we saw in **Process Risk**, we *encapsulate* earlier abstractions in order to build later ones.

³⁷<http://oracle.com>

³⁸https://en.wikipedia.org/wiki/Oracle_Coherence

³⁹<https://en.wikipedia.org/wiki/Elasticsearch>

If all this is true, the only thing we can expect in the future is that the lifespan of any ecosystem will follow an arc through creation, adoption, growth, use and finally either be abstracted over or abandoned.

tbd diagram.

Although our discipline is a young one, we should probably expect to see “Software Archaeology” in the same way as we see it for biological organisms. Already we can see the dead-ends in the software evolutionary tree: COBOL and BASIC languages, CASE systems. Languages like FORTH live on in PostScript, SQL is still embedded in everything

Boundary risk is *inside* and *outside*

Chapter 19

Agency Risk

Coordinating a team is difficult enough when everyone on the team has a single **Goal**. But, people have their own goals, too. Sometimes, the goals harmlessly co-exist with the team's goal, but other times they don't.



- Risks due to the fact that things you depend on have agency, and they have their own goals to pursue.

Figure 19.1: Agency Risk

This is **Agency Risk**. This term comes from finance and refers to the situation where you (the “principal”) entrust your money to someone (the “agent”) in order to invest it, but they don’t necessarily have your best interests at heart. They may instead elect to invest the money in ways that help them, or outright steal it.

“This dilemma exists in circumstances where agents are motivated to act in their own best interests, which are contrary to those of their principals, and is an example of moral hazard.” - Principal-Agent Problem, *Wikipedia*¹

The less visibility you have of the agent’s activities, the bigger the risk. However, the whole *point* of giving the money to the agent was that you would have to spend less time and effort managing it.

Agency Risk clearly includes the behaviour of Bad Actors². But, this is a very strict definition of **Agency Risk**. In software development, we’re not lending each other money, but we are

¹https://en.wikipedia.org/wiki/Principal-agent_problem

²https://en.wiktionary.org/wiki/bad_actor



Figure 19.2: Mitigating Agency Risk Through Monitoring

being paid by the project sponsor, so they are assuming **Agency Risk** by employing us.

As we saw in the previous section on **Process Risk**, **Agency Risk** doesn't just apply to people: it can apply to *running software or whole teams*.

Let's look at some examples of borderline **Agency Risk** situations, in order to sketch out where the domain of this risk lies.

Personal Lives

We can't (shouldn't) expect people on a project to sacrifice their personal lives for the success of the project, right? Except that "Crunch Time"³ is exactly how some software companies work:

"Game development... requires long working hours and dedication from their employees. Some video game developers (such as Electronic Arts) have been accused of the excessive invocation of "crunch time". "Crunch time" is the point at which the team is thought to be failing to achieve milestones needed to launch a game on schedule. " - Crunch Time, *Wikipedia*⁴

People taking time off, going to funerals, looking after sick relatives and so on are all **Agency Risk**, but they should be *accepted* on the project. They are a necessary **Attendant Risk** of having *staff* rather than *slaves*.

The Hero

"The one who stays later than the others is a hero." - Hero Culture, *Ward's Wiki*⁵

Conversely, Heroes put in more hours and try to rescue projects single-handedly, often cutting corners like team communication and process in order to get there.

Sometimes, projects don't get done without heroes. But other times, the hero has an alternative agenda than just getting the project done:

- A need for control, and for their own vision.

³https://en.wikipedia.org/wiki/Video_game_developer#%22Crunch_time%22

⁴https://en.wikipedia.org/wiki/Video_game_developer#%22Crunch_time%22

⁵<http://wiki.c2.com/?HeroCulture>

- A preference to work alone.
- A desire for recognition and acclaim from colleagues.
- For the job security of being a Key Man⁶.

A team *can* make use of heroism, but it's a double-edged sword. The hero can becomes a **bottleneck** to work getting done, and because want to solve all the problems themselves, they **under-communicate**.

Consultancies

When you work with an external consultancy, there is *always* more **Agency Risk** than with a direct employee. This is because as well as your goals and the employee's goals, there is also the consultancy's goals.

This is a good argument for not using consultancies, but sometimes the technical expertise they bring can outweigh this risk.

Also, try to look for *hungry* consultancies: if you being a happy client is valuable to them, they will work at a discount (either working cheaper, harder or longer or more carefully) as a result.

CV Building

This is when someone decides that the project needs a dose of "Some Technology X", but in actual fact, this is either completely unhelpful to the project (incurring large amounts of **Complexity Risk**), or merely less useful than something else.

It's very easy to spot CV building: look for choices of technology that are incongruently complex compared to the problem they solve, and then challenge by suggesting a simpler alternative.

Career Risk

Devil Makes Work

Heroes can be useful, but *underused* project members are a nightmare. The problem is, people who are not fully occupied begin to worry that actually, the team would be better off without them, and then wonder if their jobs are at risk.

The solution to this is "busy-work": finding tasks that, at first sight, look useful, and then delivering them in an over-elaborate way (Gold Plating⁷) that'll keep them occupied. This will leave you with more **Complexity Risk** than you had in the first place.

Even if they don't worry about their jobs, doing this is a way to stave off *boredom*.

⁶https://en.wikipedia.org/wiki/Key_person_insurance

⁷[https://en.wikipedia.org/wiki/Gold_plating_\(software_engineering\)](https://en.wikipedia.org/wiki/Gold_plating_(software_engineering))

Pet Projects

A project, activity or goal pursued as a personal favourite, rather than because it is generally accepted as necessary or important. - Pet Project, *Wiktionary*⁸

Sometimes, budget-holders have projects they value more than others without reference to the value placed on them by the business. Perhaps the project has a goal that aligns closely with the budget holder's passions, or its related to work they were previously responsible for.

Working on a pet project usually means you get lots of attention (and more than enough budget), but due to **Map and Territory Risk**, it can fall apart very quickly under scrutiny.

Morale Risk

Morale, also known as Esprit de Corps is the capacity of a group's members to retain belief in an institution or goal, particularly in the face of opposition or hardship - Morale, *Wikipedia*⁹

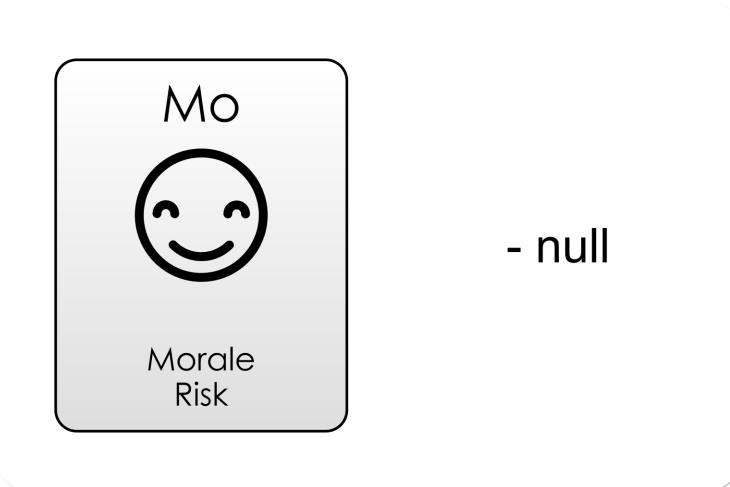


Figure 19.3: Morale Risk

Sometimes, the morale of the team or individuals within it dips, leading to lack of motivation. **Morale Risk** is a kind of **Agency Risk** because it really means that a team member or the whole team isn't committed to the **Goal**, may decide their efforts are best spent elsewhere. **Morale Risk** might be caused by:

- External factors: Perhaps the employees' dog has died, or they're simply tired of the industry, or are not feeling challenged.

⁸https://en.wiktionary.org/wiki/pet_project

⁹<https://en.wikipedia.org/wiki/Morale>

- If the team don't believe a goal is achievable, they won't commit their full effort to it. This might be due to a difference in the evaluation of the risks on the project between the team members and the leader.
- If the goal isn't considered sufficiently worthy, or the team isn't sufficiently valued.
- In military science, a second meaning of morale is how well supplied and equipped a unit is. This would also seem like a useful reference point for IT projects. If teams are under-staffed or under-equipped, this will impact on motivation too.

Hubris & Ego

It seems strange that humans are over-confident. You would have thought that evolution would drive out this trait but apparently it's not so:

"Now, new computer simulations show that a false sense of optimism, whether when deciding to go to war or investing in a new stock, can often improve your chances of winning." - Evolution of Narcissism, *National Geographic*¹⁰

In any case, humans have lots of self-destructive tendencies that *haven't* been evolved away, and we get by.

Development is a craft, and ideally, we'd like developers to take pride in their work. Too little pride means lack of care, but too much pride is *hubris*, and the belief that you are better than you really are. Who does hubris benefit? Certainly not the team, and not the goal, because hubris blinds the team to hidden risks that they really should have seen.

Although over-confidence might be a useful trait when bargaining with other humans, the thesis of everything so far is that **Meeting Reality** will punish your over-confidence again and again.

Perhaps it's a little unfair to draw out one human characteristic for attention. After all, we are **riddled with biases**. There is probably an interesting article to be written about the effects of different biases on the software development and project management processes. (This task is left as an exercise for the reader.)

Software Processes And Teams

Agency Risk doesn't just refer to people - it refers to anything which has agency over its actions.

"Agency is the capacity of an actor to act in a given environment... Agency may either be classified as unconscious, involuntary behavior, or purposeful, goal directed activity (intentional action)." - Agency, *Wikipedia*¹¹

There is significant **Agency Risk** in running software *at all*. Since computer systems follow rules we set for them, we shouldn't be surprised when those rules have exceptions that lead to disaster. For example: - A process continually writing log files until the disks fill up, crashing

¹⁰<https://news.nationalgeographic.com/news/2011/09/110914-optimism-narcissism-overconfidence-hubris-evolution-science-nature/>

¹¹[https://en.wikipedia.org/wiki/Agency_\(philosophy\)](https://en.wikipedia.org/wiki/Agency_(philosophy))

the system. - Bugs causing data to get corrupted, causing financial loss. - Malware infecting a system, and sending your passwords and data to undesirables.

Agency Risk also covers *whole teams* too. It's perfectly possible that a team within an organisation develops **Goals** that don't align with those of the overall organisation. For example: - A team introduces excessive **Bureaucracy** in order to avoid work it doesn't like. - A team gets obsessed with a particular technology, or their own internal process improvement, at the expense of delivering business value. - A marginalised team forces their services on other teams in the name of "consistency". (This can happen a lot with "Architecture", "Branding" and "Testing" teams, sometimes for the better, sometimes for the worse.)

It's About Goals

We've looked here at some illustrative examples of **Agency Risk**. But as we stated at the beginning, **Agency Risk** at any level comes down to differences of **Goals** between the different agents, whether they are *people, teams or software*.

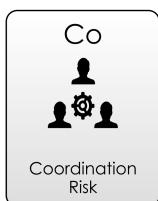
So, having looked at agents *individually*, it's time to look more closely at **Goals**, and the **Attendant Risks** when aligning them amongst multiple agents.

On to **Coordination Risk...**

Chapter 20

Coordination Risk

Coordination Risk is the risk that, a group of people (or processes), maybe with a similar **Goal In Mind** they can fail to coordinate on a way to meet this goal and end up making things worse. **Coordination Risk** is embodied in the phrase “Too Many Cooks Spoil The Broth”: more people, opinions or agents often make results worse.



- Risks that a group of agents cannot work together in a mutually beneficial way, and their behaviour devolves into competition.

Figure 20.1: Coordination Risk

As in **Agency Risk**, we are going to use the term *agent*, which refers to anything with agency¹ in a system to decide it's own fate. That is, an agent has an **Internal Model**, and can **take actions** based on it. Here, we're going to work on the assumption that the agents *are* working towards a common **Goal**, even though in reality it's not always the case, as we saw in the section on **Agency Risk**.

In this section, we'll first build up **A Model Of Coordination Risk** and what exactly coordination means and why we do it. Then, we'll look at some classic **Problems of Coordination**. Then, we're going to consider agency at several different levels (because of **Scale Invariance**). We'll look at: - **Team Decision Making**, - **Living Organisms**, - **Larger Organisations** and the staff within them, - and **Software Processes**.

... and we'll consider how **Coordination Risk** is a problem at each scale.

But for now, let's crack on and examine where **Coordination Risk** comes from.

¹<https://github.com/risk-first/website/wiki/Agency-Risk#software-processes-and-teams>

A Model Of Coordination Risk

Earlier, in **Dependency Risk**, we looked at various resources (time, money, people, events etc) and showed how we could **Depend On Them**, taking on risk. Here, however, we're looking at the situation where there is *competition for those dependencies*, that is, **Scarcity Risk**: other parties want to use them in a different way.

Competition

The basic problem of **Coordination Risk**, then, is *competition*. Sometimes, competition is desireable (such as in sports and in markets), but sometimes competition is a waste and cooperation would be more efficient. Without coordination, we would deliberately or accidentally compete for the same **Dependencies**, which is wasteful.

Why is this wasteful?

One argument could come from Diminishing Returns², which says that the earlier units of a resource (say, chocolate bars) give you more benefit than later ones.

We can see this in the graph below. Let's say A and B compete over a resource, of which there are 5 units available. For every extra A takes, B loses one. The X axis shows A's consumption of the resource, so the biggest benefit to A is in the consumption of the first unit.

As you can see, by *sharing*, it's possible that the *total benefit* is greater than it can be for either individual. But sharing requires coordination. Further, the more competitors involved, the worse a winner-take-all outcome is for total benefit.

Just two things are needed for competition to occur:

- Individual agents, trying to achieve **Goals**.
- Scarce Resources, which the agents want to use as **Dependencies**.

Coordination via Communication

The only way that the agents can move away from competition towards coordination is via **Communication**, and this is where their coordination problems begin.

You might think, therefore, that this is just another type of **Communication Risk** problem, and that's often a part of it, but even with synchronized **Internal Models**, coordination risk can occur. Imagine the example of people all trying to madly leave a burning building. They all have the same information (the building is on fire). If they coordinate, and leave in an orderly fashion, they might all get out. If they don't, and there's a scramble for the door, more people might die.

But commonly, **Coordination Risk** occurs where people have different ideas about how to achieve a **goal**, and they have different ideas because they have different evaluations of the **Attendant Risk**. As we saw in the section on **Communication Risk**, we can only hope to synchronize **Internal Models** if there are high-bandwidth **Channels** available for communication.

²https://en.wikipedia.org/wiki/Diminishing_returns

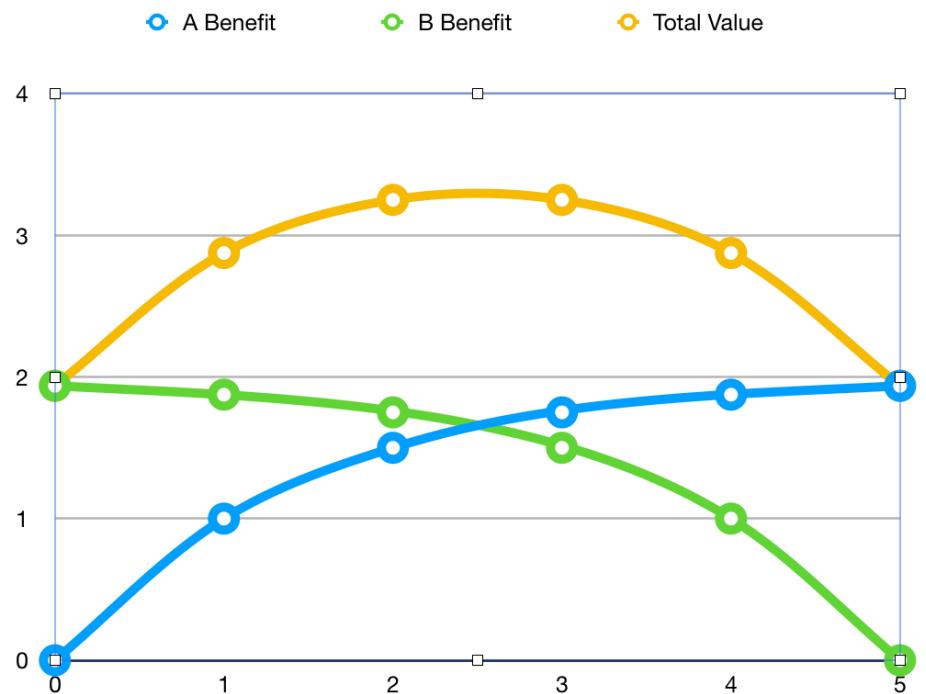


Figure 20.2: Sharing Resources. 5 units are available, and the X axis shows A's consumption of the resource. B gets whatever remains. Total benefit is maximised somewhere in the middle

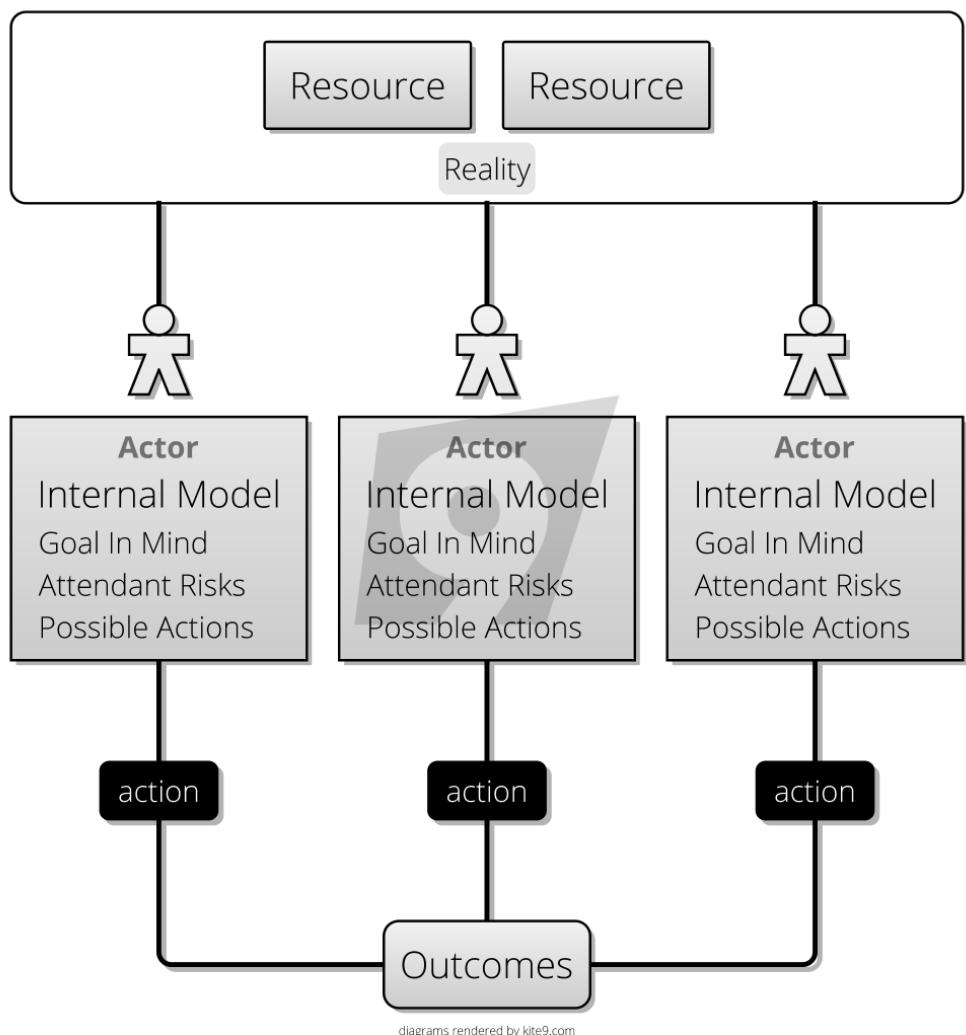


Figure 20.3: A model of competition: scarce resources, and individual agents competing for them.

Problems Of Coordination

Let's unpack this idea, and review some classic problems of coordination, none of which can be addressed without good communication:

1. **Merging Data.** If you are familiar with the source code control system, Git³, you will know that this is a *distributed* version control system. That means that two or more people can propose changes to the same files without knowing about each other. This means that at some later time, Git⁴ then has to merge (or reconcile) these changes together. Git is very good at doing this automatically, but sometimes, different people can independently change the same lines of code and these will have to be merged manually. In this case, a human arbitrator “resolves” the difference, either by combining the two changes or picking a winner.
2. **Consensus.** Making group decisions (as in elections) is often decided by votes. But having a vote is a coordination issue, and requires that everyone has been told the rules:
 - Where will the vote be held?
 - How long do you provide for the vote?
 - What do you do about absentees?
 - What if people change their minds in the light of new information?
 - How do you ensure everyone has enough information to make a good decision?
3. **Factions.** Sometimes, it’s hard to coordinate large groups at the same time, and “factions” can occur. That the world isn’t a single big country is probably partly a testament to this: countries are frequently separated by geographic features that prevent the easy flow of communication (and force). We can also see this in distributed systems, with the “split brain”⁵ problem. This is where a network of processes becomes disconnected (usually due to a network failure between data centers), and you end up with two, smaller networks with different knowledge. We’ll address in more depth later.
4. Resource Allocation⁶: Ensuring that the right people are doing the right work, or the right resources are given to the right people is a coordination issue. On a grand scale, we have Logistics⁷, and Economic Systems⁸. On a small scale, the office’s *room booking system* solves the coordination issue of who gets a meeting room using a first-come-first-served booking algorithm.
5. Deadlock⁹: Deadlock refers to a situation where, in an environment where multiple parallel processes are running, the processing stops and no-one can make progress because the resources each process needs are being reserved by another process. This is a specific issue in Resource Allocation¹⁰, but it’s one we’re familiar with in the computer

³<https://en.wikipedia.org/wiki/Git>

⁴<https://en.wikipedia.org/wiki/Git>

⁵[https://en.wikipedia.org/wiki/Split-brain_\(computing\)](https://en.wikipedia.org/wiki/Split-brain_(computing))

⁶https://en.wikipedia.org/wiki/Resource_allocation

⁷<https://en.wikipedia.org/wiki/Logistics>

⁸https://en.wikipedia.org/wiki/Economic_system

⁹<https://en.wikipedia.org/wiki/Deadlock>

¹⁰https://en.wikipedia.org/wiki/Resource_allocation

science industry. Compare with Gridlock¹¹, where traffic can't move because other traffic is occupying the space it wants to move to already.

6. Race Conditions¹²: A race condition is where we can't be sure of the result of a calculation, because it is dependent on the ordering of events within a system. For example, two separate threads writing the same memory at the same time (one ignoring and overwriting the work of the other) is a race.
7. **Contention:** Where there is **Scarcity Risk for a Dependency**, we might want to make sure that everyone gets fair use of it, by taking turns, booking, queueing and so on. As we saw in **Scarcity Risk**, sometimes, this is handled for us by the **Dependency** itself. However if it isn't, it's the *users* of the dependency who'll need to coordinate to use the resource fairly, again, by communicating with each other.

Team Decision Making

Within a team, **Coordination Risk** is at its core about resolving **Internal Model** conflicts in order that everyone can agree on a **Goal In Mind** and cooperate on getting it done. Therefore, **Coordination Risk** is worse on projects with more members, and worse in organizations with more staff.

If you are engaged in a solo project, do you suffer from **Coordination Risk** at all? Maybe: sometimes, you can feel "conflicted" about the best way to solve a problem. And weirdly, usually *not thinking about it* helps. Sleeping too. (Rich Hickey calls this "Hammock Driven Development¹³"). This is probably because, unbeknownst to you, your subconscious is furiously communicating internally, trying to resolve these conflicts itself, and will let you know when it's come to a resolution.

Vroom and Yetton¹⁴ introduced a model of group decision making which delineated five different styles of decision making within a team. These are summarised in the table below (AI, AII, CI, CII, GII). To this, I have added a sixth (UI), which is the *uncoordinated* option, where everyone competes. In the accompanying diagrams I have adopted the following convention:
- Thin lines with arrow-heads show a flow of *information*, either one-way or two-way.
- Thick lines show a flow of *opinion*.
- Boxes with corners are *decision makers*, whereas curved corners don't have a part in the decision.

Type	People Involved In Decision	Opinions	Channels Of Communication	Coordination Risk	Description
UI	1	1	0	Competition	No Coordination

¹¹<https://en.wikipedia.org/wiki/Gridlock>

¹²https://en.wikipedia.org/wiki/Race_condition

¹³<https://www.youtube.com/watch?v=f84n5oFoZBc>

¹⁴https://en.wikipedia.org/wiki/Vroom–Yetton_decision_model

Type	People Involved In Decision	Opinions	Channels Of Communication	Coordination Risk	Description
AI	1	1	s (One message to each subordinate)	Maximum Coordination Risk ¹⁵	Autocratic, top-down
AII	1	1	2 x s (Messages from/to each subordinate)		Autocratic, with information flow up.
CI	1	1 + s	> 2 x s		Individual Consultations
CII	1	1 + s	> s2		Group Consultation
GII	1 + s	1 + s	> s2	Maximum Communication Risk ¹⁶ , Schedule Risk ¹⁷	Group Consultation with voting

At the top, you have the *least* consultative styles, and at the bottom, the *most*. At the top, decisions are made with just the leader's **Internal Model** but moving down, the **Internal Models** of the rest of the team are increasingly brought into play.

The decisions at the top are faster, but don't do much for mitigating **Coordination Risk**. The ones below take longer, (incurring **Schedule Risk**) but mitigate more **Coordination Risk**. Group decision-making inevitably involves everyone *learning*, and improving their **Internal Models**.

The trick is to be able to tell which approach is suitable at which time. Everyone is expected to make decisions *within their realm of expertise*: you can't have developers continually calling meetings to discuss whether they should be using an Abstract Factory¹⁸ or a Factory Method¹⁹, this would waste time. The critical question is therefore, "what's the biggest risk?" - Is the **Coordination Risk** greater? Are we going to suffer **Dead End Risk** if the decision is made wrongly? What if people don't agree with it? Poor leadership has an impact on **Morale** too. - Is the **Schedule Risk** greater? If you have a 1-hour meeting with eight people to decide a decision, that's *one man day* gone right there: group decision making is *expensive*.

Hopefully, this model shows how *organisation* can reduce **Coordination Risk**. But, to make this work, we need more *communication*, and this has attendant complexity and time costs.

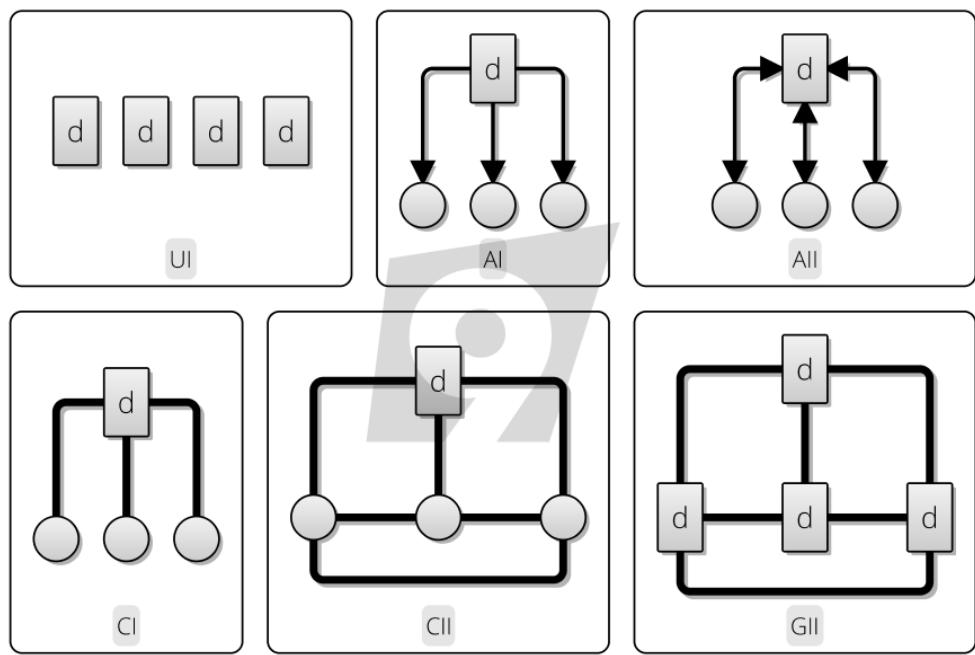
¹⁵

¹⁶

¹⁷

¹⁸https://en.wikipedia.org/wiki/Abstract_factory_pattern

¹⁹https://en.wikipedia.org/wiki/Factory_method_pattern



diagrams rendered by kite9.com

Figure 20.4: Vroom And Yetton Decision Making Styles. “d” indicates authority in making a decision. Thin lines with arrow-heads show information flow, whilst thick lines show *opinions* being passed around.

So, we can draw this diagram of our move on the **Risk Landscape**:

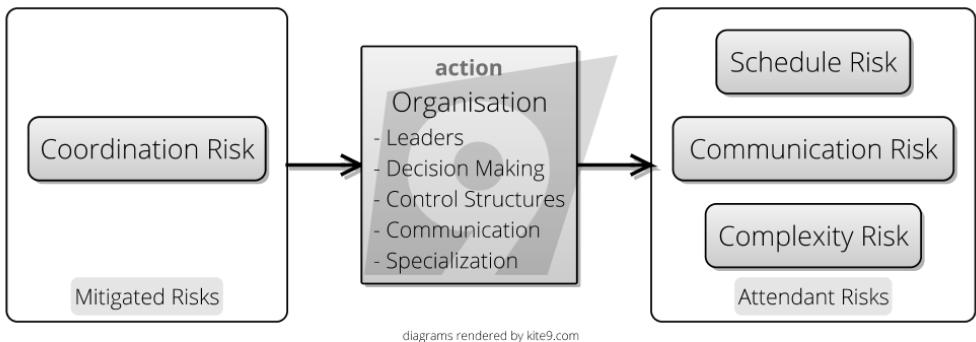


Figure 20.5: Coordination Risk traded for Complexity Risk, Schedule Risk and Communication Risk

Staff As Agents

Staff in a team have a dual nature: they are **Agents** and **Resources** at the same time. The team **depends** on staff for their resource of *labour*, but they're also part of the decision making process of the team, because they have **agency** over their own actions.

Part of **Coordination Risk** is about trying to mitigate differences in **Internal Models**. So it's worth considering how varied people's models can be: - Different skill levels - Different experiences - Expertise in different areas - Preferences - Personalities

The job of harmonizing this on a project would seem to fall to the team leader, but actually people are self-organising to some extent. This process is called Team Development²⁰:

"The forming-storming-norming-performing model of group development was first proposed by Bruce Tuckman in 1965, who said that these phases are all necessary and inevitable in order for the team to grow, face up to challenges, tackle problems, find solutions, plan work, and deliver results." - Tuckman's Stages Of Group Development, Wikipedia²¹

Specifically, this describes a process whereby a new group will form and then be required to work together. In the process, they will have many *disputes*. Ideally, the group will resolve these disputes internally and emerge as a *Team*, with a common **Goal In Mind**.

They can be encouraged with orthogonal practices such as team-building exercises²² (generally, submitting everyone to extreme experiences in order to bond them together). With enough communication bandwidth and detente, a motivated team will self-organise code reviews, information exchange and improve their practices.

As described above, the job of **Coordination** is **Resource Allocation**, and so the skills of staff can potentially be looked at as resources to allocate. This means handling **Coordination Risk** issues like:

²⁰https://en.wikipedia.org/wiki/Tuckman%27s_stages_of_group_development

²¹https://en.wikipedia.org/wiki/Tuckman%27s_stages_of_group_development

²²https://en.wikipedia.org/wiki/Team_building

- People leaving, taking their **Internal Models** and expertise with them **Key Man Risk**.
- People requiring external training, to understand new tools and techniques **Learning-Curve Risk**.
- People being protective about their knowledge in order to protect their jobs **Agency Risk**.
- Where there are mixed ability levels, senior developers not helping juniors as it “slows them down”.
- People not getting on and not helping each other.

“As a rough rule, three programmers organised into a team can do only twice the work of a single programmer of the same ability - because of time spent on co-ordination problems.” - Gerald Weinberg, *The Psychology of Computer Programming*²³

In Living Organisms

Vroom and Yetton’s organisational style isn’t relevant to just teams of people. We can see it in the natural world too. Although *the majority* of cellular life on earth (by weight) is single celled organisms²⁴, the existence of *humans* (to pick a single example) demonstrates that sometimes it’s better to try to mitigate **Coordination Risk** and work as a team, accepting the **Complexity Risk** and **Communication Risk** this entails. As soon as cells start working together, they either need to pass *resources* between them, or *control* and *feedback*.

For example, in the human body, we have various systems²⁵:

- The Respiratory System²⁶ which is responsible for ensuring that Red Blood Cells²⁷ are replenished with Oxygen, as well as disposing of Carbon Dioxide.
- The Digestive System²⁸ which is responsible for extracting nutrition from food and putting them in our Blood Plasma²⁹.
- The Circulatory System³⁰ which is responsible for moving blood cells to all the rest of the body.
- The Nervous System³¹ which is responsible for collecting information from all the parts of the body, dealing with it in the Brain³² and issuing commands.
- The Motor System³³ which contains muscles and bones, and allows us to move about.

... and many others. Each of these systems contains organs, which contain tissues, which contain cells of different types. (Even cells are complex systems containing multiple different, communicating sub-systems.) There is huge **Complexity Risk** here: the entire organism fails

²³https://en.wikipedia.org/wiki/Gerald_Weinberg

²⁴http://www.stephenjaygould.org/library/gould_bacteria.html

²⁵https://en.wikipedia.org/wiki/List_of_systems_of_the_human_body

²⁶https://en.wikipedia.org/wiki/Respiratory_system

²⁷https://en.wikipedia.org/wiki/Red_blood_cell

²⁸https://en.wikipedia.org/wiki/Human_digestive_system

²⁹https://en.wikipedia.org/wiki/Blood_plasma

³⁰https://en.wikipedia.org/wiki/Circulatory_system

³¹https://en.wikipedia.org/wiki/Nervous_system

³²<https://en.wikipedia.org/wiki/Brain>

³³https://en.wikipedia.org/wiki/Motor_system

if one of these systems fail (they are Single Points Of Failure³⁴, although we can get by despite the failure of one lung or one leg say).

Some argue³⁵ that the human nervous system is the most complex known artifact in the universe: there is huge attendant **Communication Risk** to running the human body. But, given the success of humanity as a species, you must conclude that these steps on the evolutionary **Risk Landscape** have benefitted us in our ecological niche.

The key observation from looking at biology is this: most of the cells in the human body *don't get a vote*. Muscles in the motor system have an **AI** or **AII** relationship with the brain - they do what they are told, but there are often nerves to report pain back. The only place where **CII** or **GII** could occur is in our brains, when we try to make a decision and weigh up the pros and cons.

This means that there is a deal: *most* of the cells in our body accede control of their destiny to "the system". Living within the system of the human body is a better option than going it alone. Occasionally, due to mutation, we can end up with Cancer³⁶, which is where one cell genetically "forgets" its purpose in the whole system and goes back to selfish individual self-replication (**UI**). We have White Blood Cells³⁷ in the body to shut down this kind of behaviour and try to kill the rogue cells. In the same way, society has a police force to stop undesirable behaviour amongst its citizens.

Large Organisations

Working in a large organisation often feels like being a cell in a larger organism. Just as cells live and die, but the organism goes on, in the same way, workers come and go from a large company but the organisation goes on. By working in an organisation, we give up self-control and competition and accept **AI** and **AII** power structures above us, but we trust that there is symbiotic value creation on both sides of the employment deal.

Less consultative decision making styles are more appropriate then when we don't have the luxury of high-bandwidth channels for discussion, or when the number of parties rises above a room-full of people. As you can see from the table above, for **CII** and **GII** decision-making styles, the amount of communication increases non-linearly with the number of participants, so we need something simpler. As we saw in the **Complexity Risk** section, hierarchies are an excellent way of economizing on number of different communication channels, and we use these frequently when there are lots of parties to coordinate.

In large organisations, teams are created and leaders chosen for those teams precisely to mitigate **Communication Risk**. We're all familiar with this: control of the team is ceded to the leader, who takes on the role of 'handing down' direction from above, but also 'reporting up' issues that cannot be resolved within the team. In Vroom and Yetton's model, this is moving from a **GII** or **CII** to an **AI** or **AII** style of leadership.

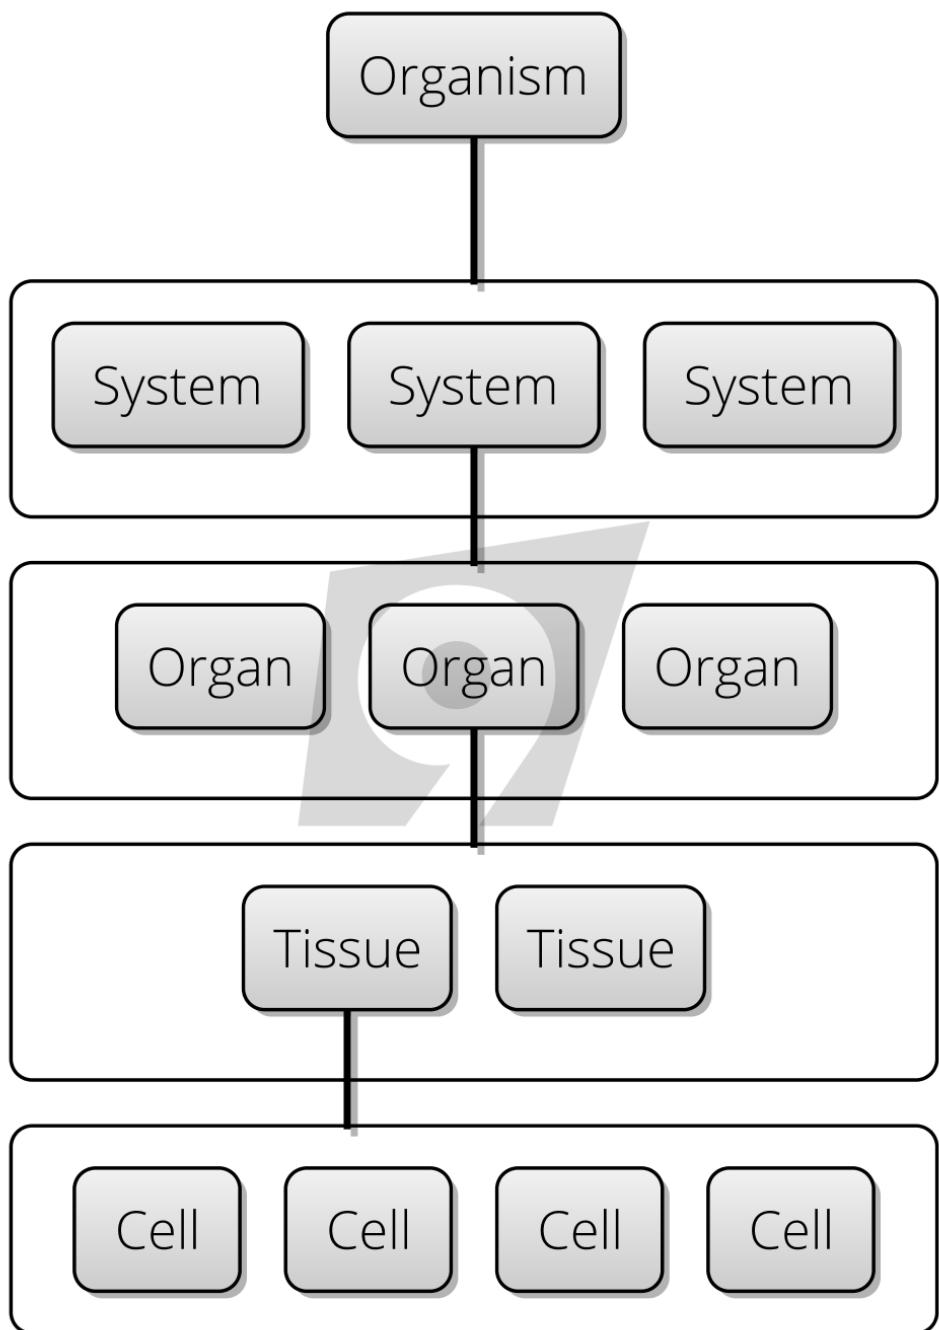
As shown in the diagram above, we end up with a hierarchy of groups, each having its own decision-making style. The team leader at the bottom level is a *decision maker* within his team,

³⁴https://en.wikipedia.org/wiki/Single_point_of_failure

³⁵<https://www.quora.com/What-is-the-most-complex-object-in-the-universe>

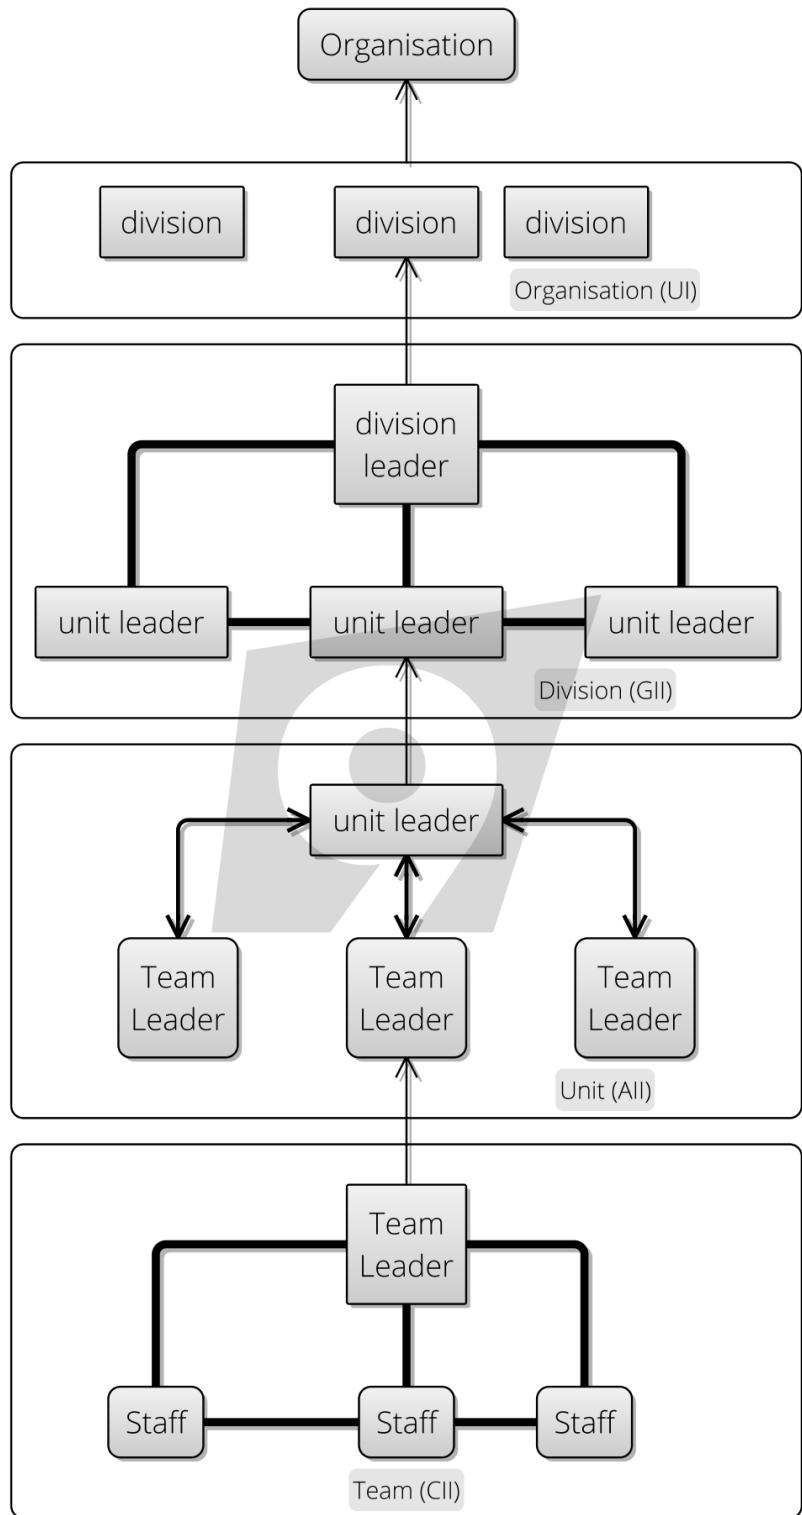
³⁶<https://en.wikipedia.org/wiki/Cancer>

³⁷https://en.wikipedia.org/wiki/White_blood_cell



diagrams rendered by kite9.com

Figure 20.6: Hierarchy of Function in the Human Body



diagrams rendered by kite9.com

Figure 20.7: Hierarchy of Function in an Organisation

but moving up, doesn't have decision making power in the next team up.. and so on.

Sometimes, parts of an organisation are encouraged *not* to coordinate, but to compete. In the diagram above, we have an M-Form³⁸ organisation, composed of *competing divisions*.

Clearly, this is just a *model*, it's not set in stone and decision making styles usually change from day-to-day and decision to decision. The same is not true in our software - *rules are rules*.

In Software Processes

It should be pretty clear that we are applying the **Scale Invariance** rule to **Coordination Risk**: all of the problems we've described as affecting teams, also affect software, although the scale and terrain are different. Software processes have limited *agency* - in most cases they follow fixed rules set down by the programmers, rather than self-organising like people can (so far).

As before, in order to face **Coordination Risk** in software, we need multiple agents all working together. **Coordination Risks** (such as race conditions or deadlock) only really occurs where *more than one thing is happening at a time*. This means we are considering *at least* multi-threaded software and anything above that (multiple CPUs, servers, data-centres and so on).

CAP Theorem

The CAP Theorem³⁹ has a lot to say about **Coordination Risk**. Imagine talking to a distributed database, where your request (*read* or *write*) can be handled by one of many agents.

In the diagram below, we have just two agents 1 and 2, in order to keep things simple. User A writes something to the database, then User B reads it back afterwards.

According to the CAP Theorem⁴⁰, there are three properties we could desire in such a system:

- **Consistency**: Every read receives the most recent value from the last write.
- **Availability**: Every request receives a response.
- **Partition tolerance**: The system can operate despite the isolation (lack of communication with) some of its agents.

The CAP Theorem⁴¹ states that this is a Trilemma⁴². That is, you can only have two out of the three properties.

There are plenty of resources on the internet that discuss this in depth, but let's just illustrate with some diagrams to show how this plays out. In our diagram example, we'll say that *any* agent can receive the read or write. So this might be a GII decision making system, because all the agents are going to need to coordinate to figure out what the right value is to return for a read, and what the last value written was. In these, the last write (setting X to 1) was sent to Agent 1 which then becomes *isolated*, and can't be communicated with, due to network failure. What will User B get back?

³⁸https://en.wikipedia.org/wiki/Multi-divisional_form

³⁹https://en.wikipedia.org/wiki/CAP_theorem

⁴⁰https://en.wikipedia.org/wiki/CAP_theorem

⁴¹https://en.wikipedia.org/wiki/CAP_theorem

⁴²<https://en.wikipedia.org/wiki/Trilemma>

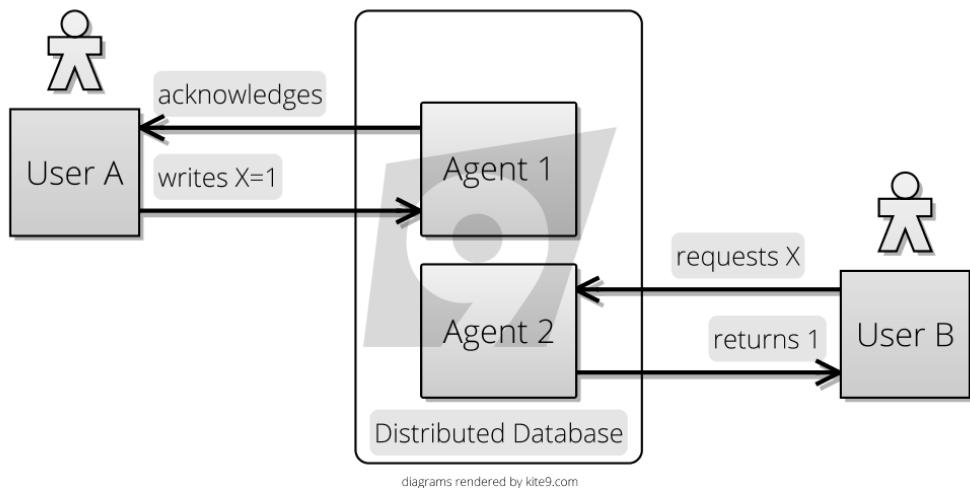
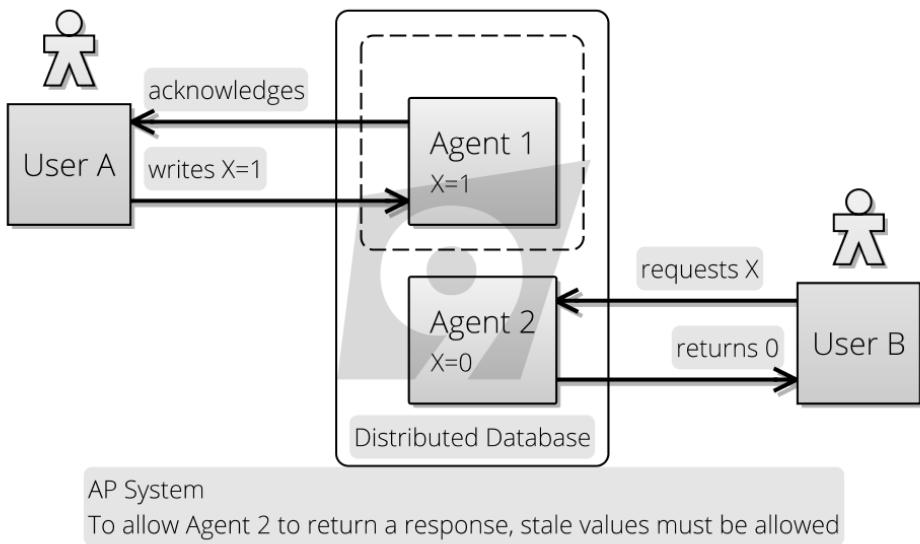


Figure 20.8: User A and User B are both using a distributed database, managed by Agents 1 and 2, whom each have their own Internal Model

With an AP System

With AP, you can see that User B is getting back a *stale value*. AP scenarios lead to Race Conditions⁴³: Agent 1's availability determines what value User B gets back.

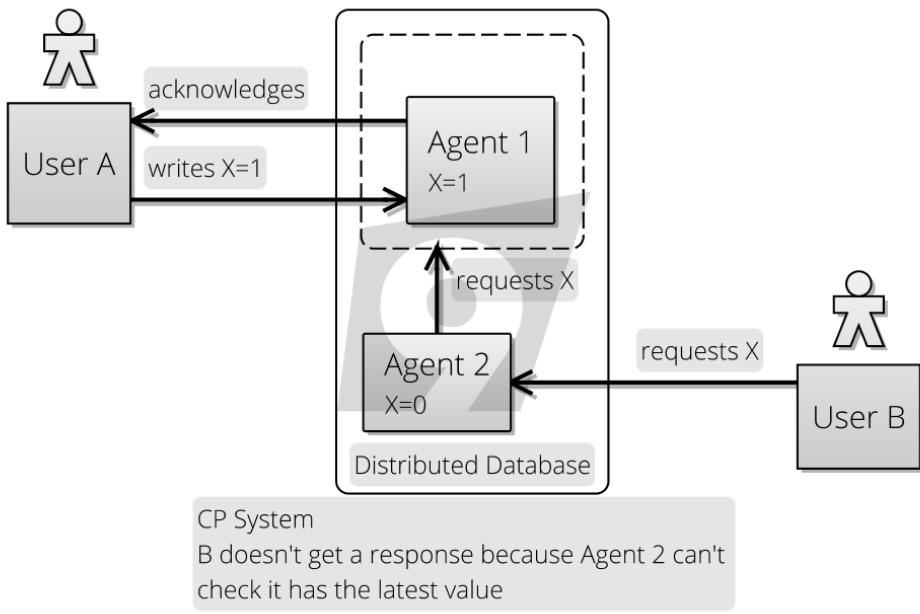
⁴³https://en.wikipedia.org/wiki/Race_condition



diagrams rendered by kite9.com

Figure 20.9: In an AP system, the User B will get back a *stale value* for X

With an CP System



diagrams rendered by kite9.com

Where Agent 2 is left waiting for Agent 1 to re-appear, we are *blocked*. So CP systems lead to Deadlock⁴⁴ scenarios.

With an CA System

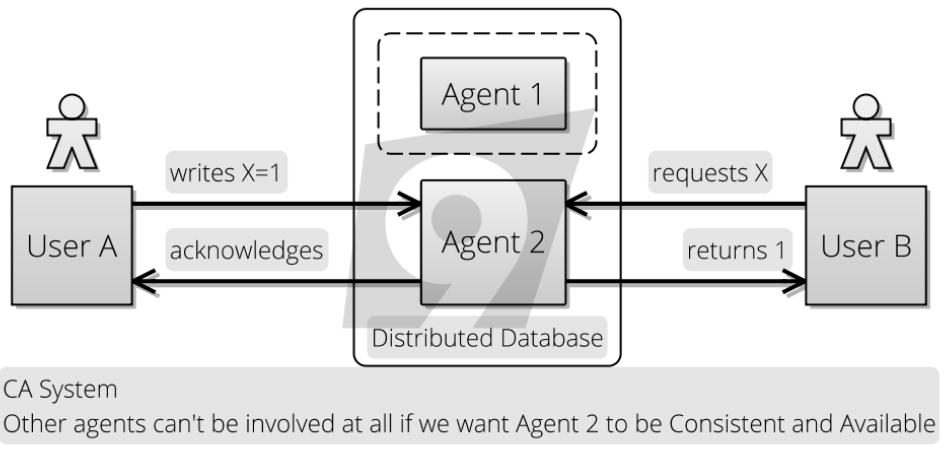


Figure 20.10: In an CA system, we can't have partition tolerance, so in order to be consistent a single Agent has to do all the work

Finally, if we have a CA system, we are essentially saying that *only one agent is doing the work*. (You can't partition a single agent, after all). But this leads to Resource Allocation⁴⁵ and **Contention** around use of the scarce resource of Agent 2's attention. (Both **Coordination Risk** issues we met earlier.)

Some Real-Life Examples

This sets an upper bound on **Coordination Risk**: we *can't* get rid of it completely in a software system, -or- a system on any other scale. Fundamentally, coordination problems are inescapable at some level. The best we can do is mitigate it by agreeing on protocols and doing lots of communication.

Let's look at some real-life examples of how this manifests in software.

ZooKeeper

First, ZooKeeper⁴⁶ is an Open-Source datastore, which is used a lot for coordinating a distributed systems, and storing things like configuration information across them. If the config-

⁴⁴<https://en.wikipedia.org/wiki/Deadlock>

⁴⁵https://en.wikipedia.org/wiki/Resource_allocation

⁴⁶<https://zookeeper.apache.org>

uration of a distributed system gets changed, it's important that *all of the agents in the system know about it*, otherwise... disaster.

This *seems* trivial, but it quickly gets out-of-hand: what happens if only some of the agents receive the new information? What happens if a datacentre gets disconnected while the update is happening? There are lots of edge-cases.

ZooKeeper handles this by communicating inter-agent with its own protocol. It elects a **master agent** (via voting), turning it into an AI-style team. If the master is lost for some reason, a new leader is elected. *Writes* are then coordinated via the **master agent** who makes sure that a *majority of agents* have received and stored the configuration change before telling the user that the transaction is complete. Therefore, ZooKeeper is a CP system.

Git

Second, git⁴⁷ is a (mainly) write-only ledger of source changes. However, as we already discussed above, where different agents make incompatible changes, someone has to decide how to resolve the conflicts so that we have a single source of truth.

The **Coordination Risk** just *doesn't go away*.

Since multiple users can make all the changes they like locally, and merge them later, Git is an AP system: individual users may have *wildly* different ideas about what the source looks like until the merge is complete.

Bitcoin

Finally, Bitcoin (BTC)⁴⁸ is a write-only distributed ledger⁴⁹, where agents *compete* to mine BTC, but also at the same time record transactions on the ledger. BTC is also AP, in a similar way to Git. But new changes can only be appended if you have the latest version of the ledger. If you append to an out-of-date ledger, your work will be lost.

Because it's based on outright competition, if someone beats you to completing a mining task, then your work is wasted. So, there is *huge Coordination Risk*.

For this reason, BTC agents **coordinate** into mining consortia⁵⁰, so they can avoid working on the same tasks at the same time. But this in itself is a problem, because the whole *point* of BTC is that it's competitive, and no one entity has control. So, mining pools tend to stop growing before they reach 50% of the BTC network's processing power. Taking control would be politically disastrous⁵¹ and confidence in the currency (such as there is) would likely be lost.

⁴⁷<https://en.wikipedia.org/wiki/Git>

⁴⁸<https://en.wikipedia.org/wiki/Bitcoin>

⁴⁹https://en.wikipedia.org/wiki/Distributed_ledger

⁵⁰https://en.bitcoin.it/wiki/Comparison_of_mining_pools

⁵¹https://www.reddit.com/r/Bitcoin/comments/5fe9vz/in_the_last_24hrs_three_mining_pools_have_control/

Communication Is For Coordination

So, now we have a fundamental limit on how much **Coordination Risk** we can mitigate. And, just as there are plenty of ways to mitigate **Coordination Risk** within teams of people, organisations or living organisms, so it's the case in software.

Earlier in this section, we questioned whether **Coordination Risk** was just another type of **Communication Risk**. However, it should be clear after looking at the examples of competition, cellular life and Vroom and Yetton's Model that this is exactly *backwards*:

- Most single-celled life has no need for communication: it simply competes for the available resources. If it lacks anything it needs, it dies.
- There are *no* lines of communication on the UI decision-type. It's only when we want to avoid competition, by sharing resources and working towards common goals that we need to communicate.
- Therefore, the whole point of communication *is for coordination*.

In the next section, **Map And Territory Risk**, we're going to look at some new ways in which systems can fail, despite their attempts to coordinate.

Chapter 21

Map And Territory Risk

As we discussed in the section on **Abstraction**, our understanding of the world is entirely informed by the names we give things and the abstractions we create. (In the same way, **Risk-First** is about *identifying patterns* within software development and calling them out.) Our **Internal Models** are a model of the world based on these patterns, and their relationships.

So there is a translation going on here: observations about the arrangement of *atoms* in the world get turned into patterns of *information* (measured in bits and bytes).

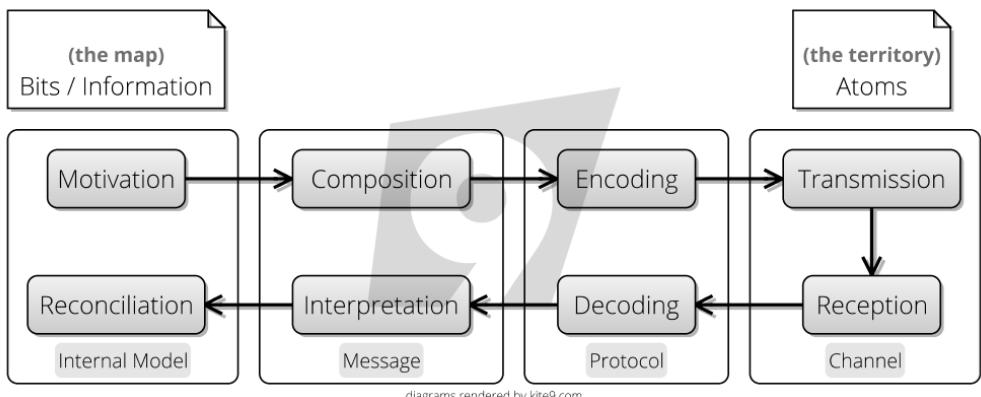
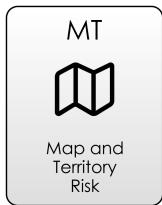


Figure 21.1: Maps and Territories, and Communication happening between them

Map And Territory Risk is the risk we face because we base our behaviour on our **Internal Models** rather than reality itself. It comes from the expression “Confusing the Map for the Territory”, attributed to Alfred Korzybski:

“Polish-American scientist and philosopher Alfred Korzybski remarked that “the map is not the territory” and that “the word is not the thing”, encapsulating his view that an abstraction derived from something, or a reaction to it, is not the thing itself. Korzybski held that many people *do* confuse maps with territories, that is, confuse models of reality with reality itself.” - Map-Territory Relation,



- Risks due to the differences between reality and the internal model of reality, and the assumption that they are equivalent.

Figure 21.2: Map And Territory Risk

In this section, we're going to make a case for analysing **Map and Territory Risk** along the same axes we introduced for **Feature Risk**, that is **Fitness**, **Audience** and **Evolution**. After that, we are going to widen the scope by looking at **Map and Territory Risk** within the context of **machines, people, hierarchies and markets**.

tbd - diagram of how our actions are based on the map, not the territory.

Fitness

In the picture shown here, from the Telegraph newspaper, the driver *trusted* the SatNav to such an extent that he didn't pay attention to the road-signs around him, and ended up getting stuck. This wasn't borne of stupidity, but experience: SatNavs are pretty reliable. *So many times* the SatNav had been right, that the driver stopped *questioning its fallibility*.

So, there are two **Map and Territory Risks** here: - The **Internal Model** of the *SatNav* contained information that was wrong: the track had been marked up as a road, rather than a path.

- The **Internal Model** of the *driver* was wrong: his abstraction of "the SatNav is always right" turned out to be only *mostly* accurate.

Internal Models as Dependencies, Features

What are the risks at play here? We've already looked in detail at the **Dependency Risks** involved in relying on something like a SatNav, in the **Software Dependency Risk** section. But here, we are really looking at the **Internal Models themselves** as a source of **Dependency Risk** too.

We could argue that the SatNav and the Driver's **Internal Model** had bugs in them. That is, they both suffer the **Feature Implementation Risk** we saw in the **Feature Risk** section. If a SatNav has too much of this, you'd end up not trusting it, and getting a new one. With

¹https://en.wikipedia.org/wiki/Map–territory_relation



A van got stuck on a narrow footpath when the driver took a wrong turn while blindly following his sat-nav. Photo: MEN

Figure 21.3: Sat Nav Blunder Sends Asda Van Crashing Narrow Footpath - Telegraph Newspaper

your *personal Internal Model*, you can't buy a new one, but you may learn to *trust certain abstractions less*, as this driver did.

In the **Feature Risk** section, we broke down **Feature Risk** on three axes: **Fitness**, **Evolution** and **Audience**.

Lets do this again and see how each type of **Feature Risk** can manifest in the **Internal Model**:

Dimension

Feature Risk

Map and Territory Examples

Fitness

Conceptual Integrity Risk

Implementation Risk

A filing cabinet containing too much junk.

Learning things that aren't useful.

Knowing how a car works, but actually needing to know how to drive.

Knowing how to program in one language, when another would be more appropriate.

Sat Nav had the wrong route.

Not quite remembering a recipe properly.

Audience

Feature Access Risk

Market Risk

Memes.

Demand for courses.

Metrics.

Echo-chambers.

Shared values which exclude certain people.

Ideas going "out of fashion".

Evolution

Feature Drift Risk

Regression Risk

Knowing outdated tools.

Writing last year's date on the cheque.

The bank sending letters to your old address.

Forgetting things

As with **Features** in a product, Information in an **internal model** has at least these three dimensions:

- **Fitness:** as discussed above with the SatNav example, this is how closely the information matches reality, and how *useful that is to us* (models that contain too much detail are as bad as models with too little).
- **Audience:** is all about how a piece of information is *shared* between many **Internal Models**, and it's this we are going to address further now.
- **Evolution:** is all about how **Internal Models** change when they meet reality, and we'll cover that last.

Audience

We already know a lot about **Internal Models** and audience, as these have been the subject of previous sections: - We know from looking at **Communication Risk** that communication allows us to *share* information between **Internal Models**. - We know from **Coordination Risk** the difficulties inherent in aligning **Internal Models** so that they cooperate. - Job markets show us that there is demand for people with certain *skills*. This demonstrates to us that **Market Risk** is as applicable to **Internal Models** containing certain information as it is to products containing **Features**. This was the focus of the **Ecosystem** discussion in **Boundary Risk**.

... And, we're all familiar with *memes*:

"A meme acts as a unit for carrying cultural ideas, symbols, or practices, that can be transmitted from one mind to another through writing, speech, gestures, rituals, or other imitable phenomena with a mimicked theme." - Meme, *Wikipedia*²

Therefore, we should be able to track the rise-and-fall of *ideas* much as we can track stock prices. And in effect, this is what Google Trends³ does. In the chart below, we can see the relative popularity of two search terms over time. This is probably as good an indicator as any of the audience for an abstraction at any point in time.

Example: Hype Cycles

Most ideas (and most products) have a slow, hard climb to wide-scale adoption. But some ideas seem to disperse much more rapidly and are picked up quickly because they are exciting and promising, having greater "memetic potential" within society. One way this evolution manifests itself in the world is through the Hype Cycle⁴:

"The hype cycle is a branded graphical presentation developed and used by the American research, advisory and information technology firm Gartner, for representing the maturity, adoption and social application of specific technologies. The hype cycle provides a graphical and conceptual presentation of the maturity of emerging technologies through five phases." - Hype Cycle, *Wikipedia*⁵

²<https://en.wikipedia.org/wiki/Meme>

³<https://trends.google.com>

⁴https://en.wikipedia.org/wiki/Hype_cycle

⁵https://en.wikipedia.org/wiki/Hype_cycle

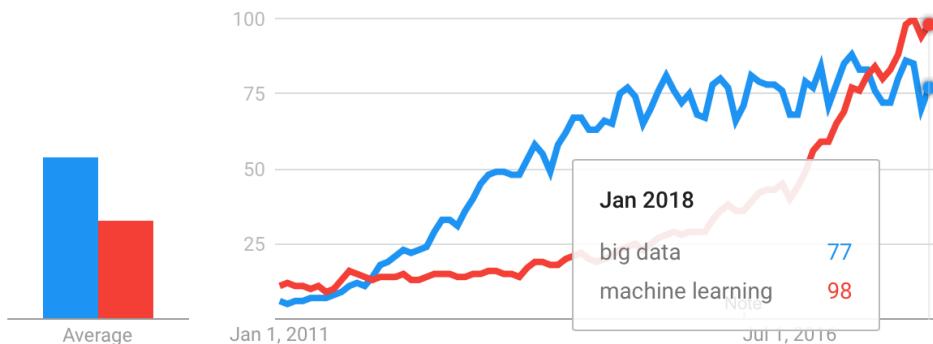


Figure 21.4: Relative popularity of “Machine Learning” and “Big Data” as search terms on Google Trends, 2011-2018

The five phases (and the “Hype” itself) are shown in the chart below, with the thick black line being “Hype”:

Also in this diagram we are showing where the hype originates: - The **saturation** of the idea within the audience (a dotted line). - The **amount known** about the idea by the audience (a **Learning Curve**, if you will, a dashed line).

Both of these are modelled with Cumulative Distribution⁶ curves. From these two things, we can figure out where our maximum **Map and Territory Risk** lies: it’s the point where awareness of an idea is furthest from the understanding of it. This acts as a “brake” on the **hype** around the idea, corresponding to the “Trough of Disillusionment”.

Where the **saturation** and **knowledge** grow together, there is no spike in **Map and Territory Risk** and we don’t see the corresponding “Trough of Disillusionment” at all, as shown in this chart:

Evolution

The section on **Communication Risk** introduced the following model for ideas:

But what happens next? As we saw in **Boundary Risk**, the **Peter Principle** applies, people will use dependencies up to the point when they start breaking down.

Example: Metrics

Let’s dive into a specific example now: someone finds a useful new metric that helps in evaluating performance.

⁶https://en.wikipedia.org/wiki/Cumulative_distribution_function#Use_in_statistical_analysis

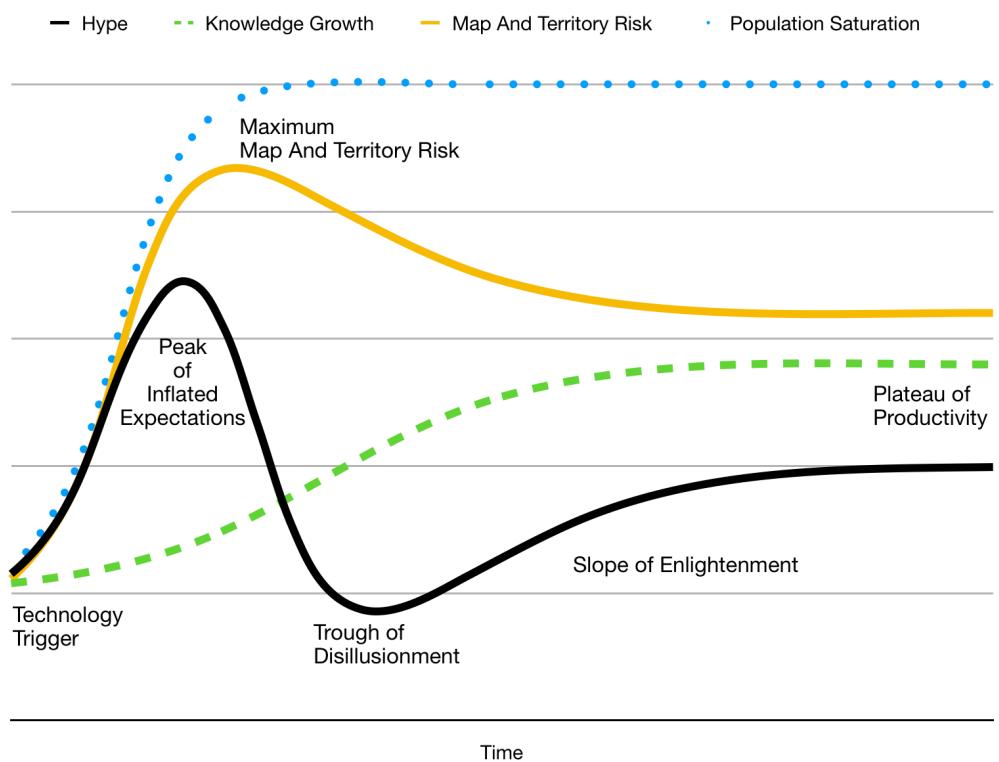


Figure 21.5: Hype Cycle, along with Map & Territory Risk

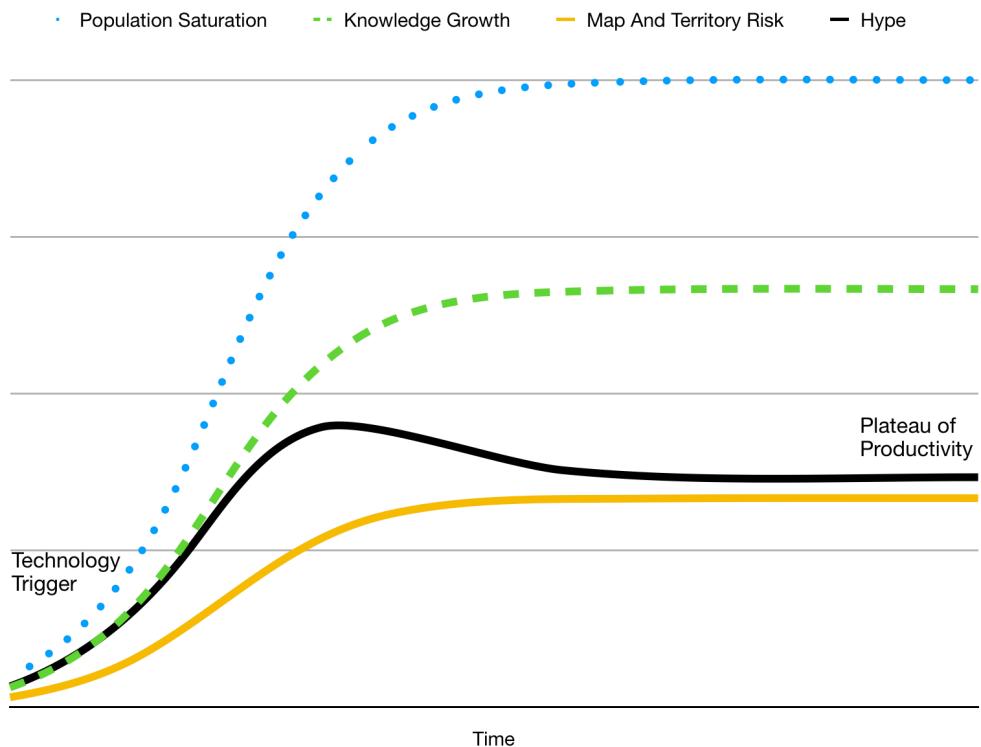


Figure 21.6: Hype Cycle 2: Slower growth of Map and Territory Risk means no “Trough of Disillusionment”

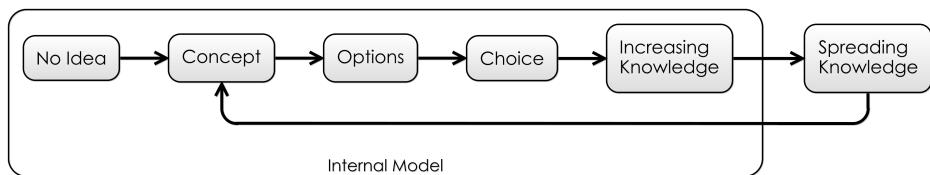


Figure 21.7: Spread of information between Internal Models

It might be:

- **SLOC (Source Lines Of Code)**: i.e. the number of lines of code each developer writes per day/week whatever.
- **Function Points**: the number of function points a person on the team completes, each sprint.
- **Code Coverage**: the number of lines of code exercised by unit tests.
- **Response Time**: the time it takes to respond to an emergency call, say, or to go from a feature request to production.
- **Release cadence**: number of releases a team performs, per month, say.

With some skill, they may be able to *correlate* this metric against some other more abstract measure of success. For example: - “quality is correlated with more releases” - “user-satisfaction is correlated with SLOC” - “revenue is correlated with response time”

Because the *thing on the right* is easier to measure than the *thing on the left*, it becomes used as a proxy (or, Map) for the thing they are really interested in (the Territory). At this point, it’s *easy* to communicate this idea with the rest of the team, and the *market value of the idea is high*: it is a useful representation of reality, which is shown to be accurate at a particular point in time.

But *correlation* doesn’t imply *causation*. The *cause* might be different:

- quality and number of releases might both be down to the simplicity of the product. - user satisfaction and SLOC might both be down to the calibre of the developers. - response time and revenue might both be down to clever team planning.

Metrics are *seductive* because they simplify reality and are easily communicated. But they *inherently contain Map and Territory Risk*: By relying *only* on the metrics, you’re not really *seeing* the reality. The devil is in the detail.

Reality Evolves

In the case of metrics, this is where they start being used for more than just indicators, but as measures of performance or targets: - If a team is *told* to do lots of releases, they will perform lots of releases *at the expense of something else*. - If team members are promoted according to SLOC, they will make sure their code takes up as many lines as possible. - In the UK, ambulances were asked to wait before admitting patients to Emergency wards, in order that hospitals could meet their targets⁷.

Some of this seems obvious: *Of course* SLOC is a terrible measure performance! We’re not that stupid anymore. The problem is, it’s not so much the *choice* of metric, but the fact that *all* metrics merely approximate reality with a few numbers. The map is *always* simpler than the territory, therefore there can be no perfect metrics.

In the same way that **markets evolve to demand more features**, our behaviour evolves to incorporate new ideas. The more popular an idea is, the more people will modify their behaviour as a result of it, and the more the world will change. Will the idea still be useful as the world adapts? Although the **Hype Cycle** model doesn’t cover it, ideas and products all eventually have their day and decline in usefulness.

⁷https://en.wikipedia.org/wiki/NHS_targets

Bad Ideas

There are plenty of ideas which *seem a really good idea at the time* but then end up being terrible. It's only as we *learn about the products* and realize the hidden **Map and Territory Risk** that we stop using them. While SLOC is a minor offender, CFCs⁸ or Leaded Petrol⁹ are more significant examples.

The following Hyph Cycle chart shows an initially promising idea that turns out to be terrible, and there is a “Period of Inoculation” where the population realise their mistake. There is “negative hype” as they work to phase out the offending idea:

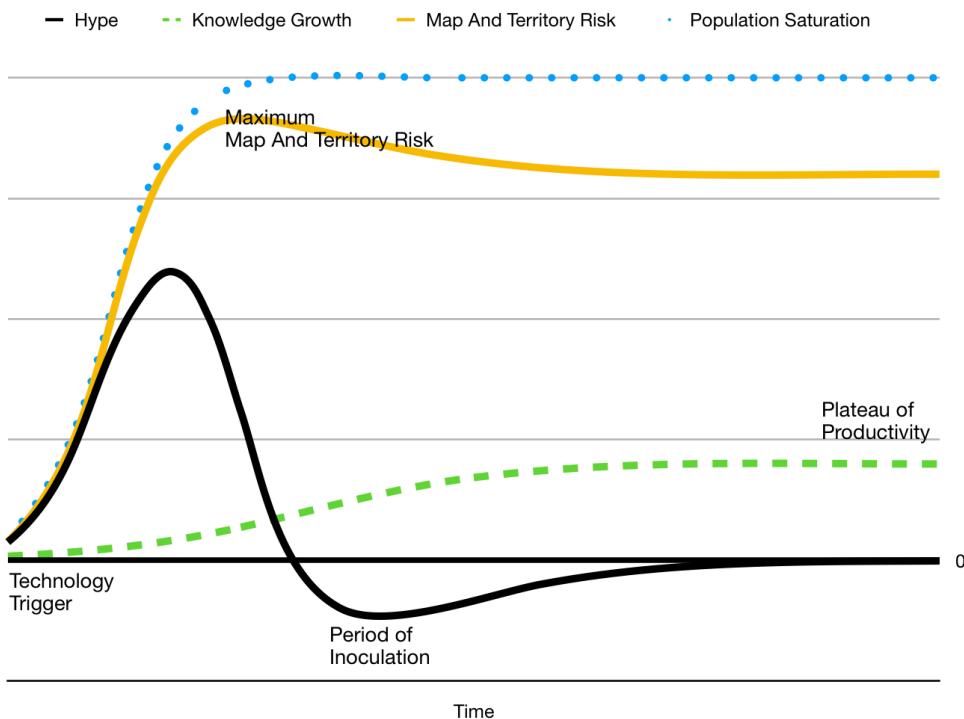


Figure 21.8: Hype Cycle For Something that turns out to be a *bad* idea

Humans and Machines

In the example of the SatNav, we saw how the *quality of Map and Territory Risk* is different for *people* and *machines*. Whereas people *should* be expected show skepticism to new (unlikely) information, our databases accept it unquestioningly. *Forgetting* is an everyday, usually benign part of our human **Internal Model**, but for software systems it is a production crisis involving 3am calls and backups.

⁸<https://en.wikipedia.org/wiki/Chlorofluorocarbon>

⁹<https://en.wikipedia.org/wiki/Tetraethyllead>

For Humans, **Map and Territory Risk** is exacerbated by cognitive biases¹⁰:

“Cognitive biases are systematic patterns of deviation from norm or rationality in judgment, and are often studied in psychology and behavioral economics.” - Cognitive Bias, Wikipedia¹¹

There are *lots* of cognitive biases. But let's just look at a couple that are relevant to **Map and Territory Risk**:

- **Availability Heuristic:** People overestimate the importance of knowledge they have been exposed to.
- **The Ostrich Effect:** Which is where dangerous information is ignored or avoided because of the emotions it will evoke.
- **Bandwagon Effect:** People like to believe things that other people believe. Could this be a factor in the existence of the Hype Cycle?

Hierarchical Organisations

Map And Territory Risk “trickles down” through an organisation. The higher levels have an outsize ability to pervert the incentives at lower levels because once an organisation begins to pursue a “bullshit objective”, the whole company can align to this.

The Huffington Post¹² paints a brilliant picture of how Volkswagen managed to get caught faking their emissions tests. As they point out:

“The leadership culture of VW probably amplified the problem by disconnecting itself from the values and trajectory of society, by entrenching in what another executive in the auto industry once called a “bullshit-castle”... No engineer wakes up in the morning and thinks: OK, today I want to build devices that deceive our customers and destroy our planet. Yet it happened. Why? Because of hubris at the top.” - Otto Scharmer, *Huffington Post*¹³.

This article identifies the following process:

- **De-sensing:** VW Executives ignored *The Territory* society around them (such as the green movement), ensuring their maps were out of date. The top-down culture made it hard for reality to propagate back up the hierarchy.
- **Hubris/Absencing:** They pursued their own metrics of *volume* and *cost*, rather than seeking out others (a la the Availability Heuristic Bias). That is, focusing on their own *Map*, which is *easier* than checking the *Territory*. (See **Hubris** in the **Agency Risk** section).
- **Deception:** Backed into a corner, engineers had no choice but to find “creative” ways to meet the metrics.
- **Destruction:** Eventually, the truth comes out, to the detriment of the company, the environment and the shareholders.

¹⁰https://en.wikipedia.org/wiki/List_of_cognitive_biases

¹¹https://en.wikipedia.org/wiki/List_of_cognitive_biases

¹²https://www.huffingtonpost.com/otto-scharmer/the-fish-rots-from-the-he_b_8208652.html

¹³https://www.huffingtonpost.com/otto-scharmer/the-fish-rots-from-the-he_b_8208652.html

As the article's title summarizes "A fish rots from the head down".

Personal Example

A similar (but less catastrophic) personal story from a bank I worked at, where the objectives end up being mis-aligned *within the company*:

1. My team had been tasked with building automated "smoke tests" of an application. But this was bullshit: We only needed to build these *at all* because the application was so complex. The reason it was so complex was...
2. The application was being designed within a "Framework" constructed by the department. However, the framework was only being used by this one application. Building a "reusable" framework which is only used by a single application is bullshit. But, we had to do this because...
3. The organisational structure was created along a "matrix", with "business function" on one axis and "functional area" on another. Although we were only building the application for a single business function, it was expected to cater with all the requirements from the an entire "functional area". This was bullshit too, because...
4. The matrix structure was largely the legacy of a recent merger with another department. As **Conway's Law** predicts, our software therefore had to reflect this structure. But this was bullshit because...
5. The matrix structure didn't represent reality in any useful way. It was designed to pacify the budget committee at the higher level, and try to demonstrate attributes such as *control* and *governance*. But this was bullshit too, because...
6. The budget that was given to our department was really based on how much fear the budget holders currently had of the market regulators. But this was bullshit too, because...
7. At a higher level, the executives had realised that our division wasn't one of the banks strategic strengths, and was working to close it all down anyway.

When faced with so many mis-aligned objectives, it seemed completely hopeless to concentrate on the task at hand. But then, a colleague was able to nihilistically add one final layer to this onion by saying:

8. "It's all about chasing money, which is bullshit, because life is bullshit."

Picking Fights

It feels like there's no way back from that. All of life might well be a big **Map and Territory** illusion. But let's analyse just a bit: - At each layer, the objectives changed. But, they impacted on the objectives of the layer below. - Therefore, it seems like the more layers you have, the less likely it is that your objectives become inconsistent between the lower and higher levels. - On a new project, it seems like a good idea to model this stuff: does the objective of the work you're about to undertake "align" with the objectives at a higher level?

Trying to spot **Map and Territory Risk** ahead-of-time in this manner seems like a useful way of trying to avoid **Vanity Projects**, and, if you are good at it, allows you to see which **Goals** in the organisation are fragile and likely to change. However, usually, if you are working in a team, you have limited agency to decide which projects you feel are valuable.

This comes down to a personal decision: do you want to spend time working on projects that you know are going in the bin? Some developers have the attitude that, so long as they get paid, it doesn't matter. But others are in it for the satisfaction of the work itself, so this ends up being a personal call. (This theme will be developed further in **Staging and Classifying**.)

Markets

So far, we've considered what happens to individuals, teams and organisations when told to optimise around a particular objective. In **Coordination Risk** we looked at how **Communication** was critical for Coordination to happen. And, as we've already discussed, **Abstraction** is a key part of communication.

The languages we adopt or create are *sets of useful abstractions* that allow us to communicate. But what happens when this goes wrong?

Inadequate Equilibria¹⁴ by Eleizer Yudkowsky, looks at how perverse incentive mechanisms break not just departments, but entire societal systems. He highlights one example involving *academics and grantmakers* in academia:

- It's not very apparent which scientists are better than which other scientists.
- One proxy is what they've published (scientific papers) and where they've published (journals).
- Universities want to attract research grants, and the best way to do this is to have the best scientists.
- Because "best" isn't measureable, they use the proxy.
- Therefore, immense power rests in the hands of the journals, since they can control the money-proxy.
- Therefore, journals are able to charge large amounts of money to universities for subscriptions.

"Now consider the system of scientific journals... Some journals are prestigious. So university hiring committees pay the most attention to publications in that journal. So people with the best, most interesting-looking publications try to send them to that journal. So if a university hiring committee paid an equal amount of attention to publications in lower-prestige journals, they'd end up granting tenure to less prestigious people. Thus, the whole system is a stable equilibrium that nobody can unilaterally defy except at cost to themselves." - Inadequate Equilibria, Eleizer Yudkowsky¹⁵

As the book points out, while everyone *persists* in using an inadequate abstraction, the system is broken. However, **Coordination** would be required for everyone to *stop* doing it this way, which is hard work. (At least within a hierarchy, Maps can get fixed.)

This is a *small example* from a much larger, closely argued book, and it's worth taking the time to read a couple of the chapters on this interesting topic.

As usual, this section forms a grab-bag of examples in a complex topic. But it's time to move on as there is one last stop we have to make on the **Risk Landscape**, and that is to look at

¹⁴<https://equilibriabook.com>

¹⁵<https://equilibriabook.com/molochs-toolbox/>

Operational Risk.

(NB: The Hype Cycle model is available in **Numbers** form here¹⁶.)

(talk about how operational risk is an extension of this). tbd

¹⁶<https://github.com/risk-first/website/blob/master/RiskMatrix.numbers>

Chapter 22

Operational Risk

In this section on Operational Risks, we're going to take our head out of the clouds a bit and start considering the realities of running software systems in the real world. After all, **Coordination Risk** and **Map And Territory Risk** got a bit theoretical. Here, we're going to look at **Operational Risks**, including **Security Risk** and **Reputational Risk**: real-world concerns for anyone running a business.

A Recap

But before we go there, let's try and recap on where we've come so far. So far, we've been looking at risks to *systems in general*:

-
-
-
- Communication Risk: how they mitigate **Coordination Risk** by taking on Communication Risk.

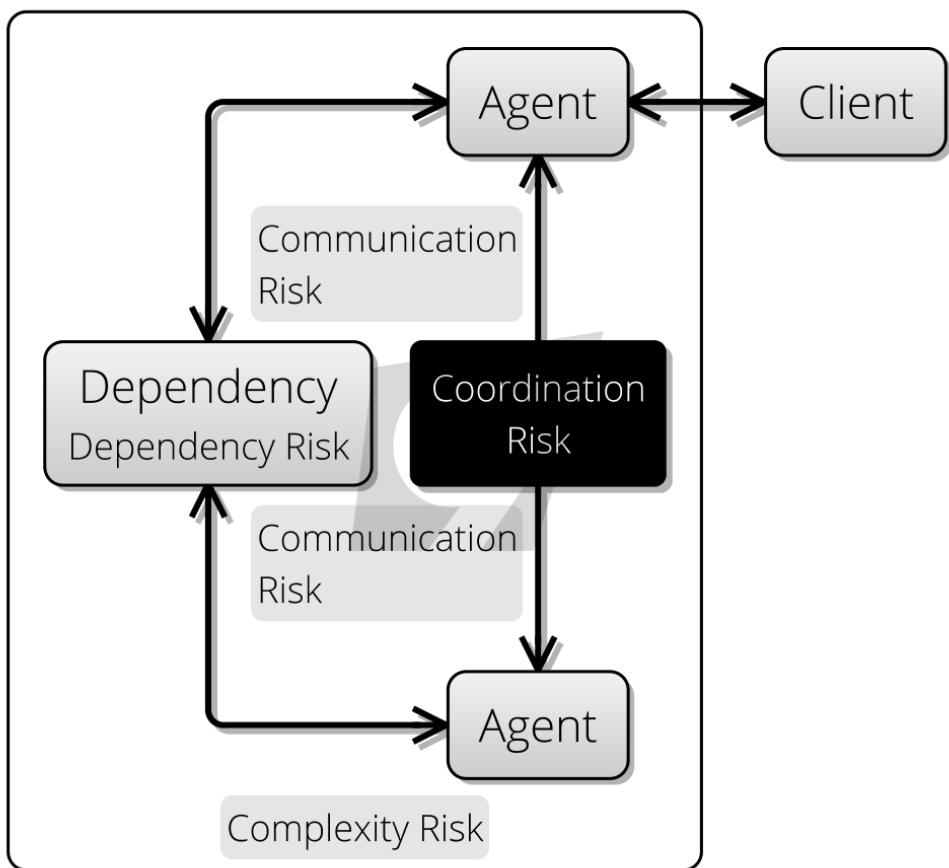
Here is a diagram that shows how these different elements line up:

(tbd, remove outside client)

Operational Risks

It's tempting to take a very narrow view of the dependencies of a system, but **Operational Risks** are often caused by dependencies we don't consider - the *context* within which the system is operating. Here are some examples:

- Staff Dependencies (**Staff Risk**):
 - Freak weather conditions affecting ability of staff to get to work, interrupting the development and support teams.



diagrams rendered by kite9.com

Figure 22.1: Systemic View of Risks

- Reputational damage caused when staff are rude to the customers.
- Infrastructure Dependencies (**Availability Risk**):
 - A data-centre going off-line, causing your customers to lose access.
 - A power cut causing backups to fail.
 - Not having enough desks for everyone to sit at.
- Process Dependencies (**Process Risk**):
 - Regulatory change, which means you have to adapt your business model.
 - Insufficient controls which means you don't notice when some transactions are failing, leaving you out-of-pocket.
 - Data loss because of bugs introduced during an untested release.
- Software Dependencies (**Software Dependency Risk**):
 - Hackers breaking into the system and bringing your service down.
- Agency Dependencies (**Agency Risk**):
 - Suppliers deciding to stop supplying you with something you need.
 - Workers going on strike.
 - Employees trying to steal from the company (bad actors).

.. basically, a long laundry-list of everything that can go wrong due to operating in “The Real World”. Although these issues don’t exist with *ideal* dependencies, pragmatically, when we design our system, we design in features to *strengthen* it against **Operational Risks** such as these. *However*, we don’t get these for free: they come at the cost of extra **Complexity Risk**.

For example, as we saw in **Development Process**, a lot of the processes we put in place (e.g. **Continuous Integration** or various types of **Testing**) are there to mitigate **Operational Risk** problems caused by *releasing buggy software*. **Unit Testing** increases our **Kolmogorov Complexity** (as there is more code) but we accept this extra complexity as the price for mitigating this **Operational Risk**. **User Acceptance Testing** increased our **Schedule Risk**, but again reduced the likelihood of bugs making it into production.

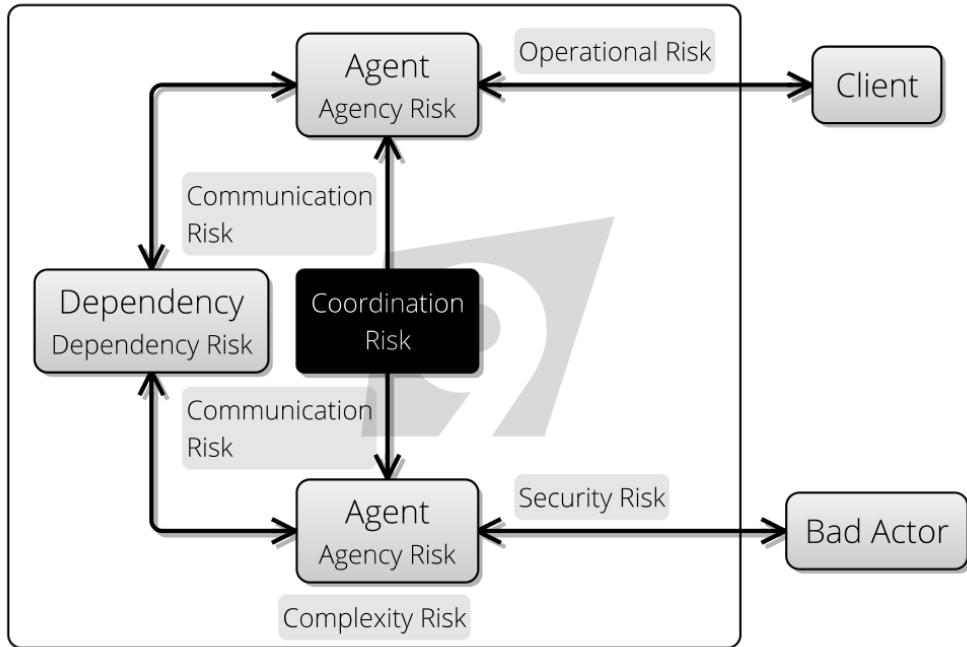
Dependencies are not just things we *use*, then: For a system to run well, it needs to carefully manage unreliable dependencies, and ensure their safety and availability. In the example of the human food system, say, it’s the difference between **Hunter-Gathering** (picking up food where we find it) and **Farming**.

Operational Risk Management

Since we have looked in detail at various types of **Dependency Risk**, in this section we’ll be focusing on the job of **Operational Risk Management**, and the tools it brings to bear on **Dependency Risk**:

“Operational Risk Management is the oversight of Operational Risk, including the risk of loss resulting from inadequate or failed internal processes and systems; human factors; or external events.” - Operational Risk Management, *Wikipedia*¹

¹https://en.wikipedia.org/wiki/Operational_risk_management



diagrams rendered by kite9.com

tbd, need to modify this,

Mitigating Operational Risk

Operational Risk then, is really just an extended view on **Dependency Risk**, with reference to embedding the system of dependencies in the real world. There are various activities then that Operational Risks, then are things that happen *to* our carefully constructed, theoretical system, as shown in this diagram:

diagram: our system -> event's exploit weakness -> effect -> reaction -> recovery -> adaptation

(our system: meeting reality, deployment) (event: detection) (weakness: minimization (see Security risk)) (effect: what happens, how much chaos ensues) (sensing / detection) (reaction: monitoring, etc) (recovery: how we fix it) (adaptation: how the system changes in the future) (prediction: how to forecast failures) pestle, environmental scanning. / anticipation (practicing: e.g. testing failover etc.) (changing outside world)

These are properties not only of software systems, but biological systems, and businesses too. Let's now take each in turn and inspect it further.

Therefore, one of the best defences against **Operational Risk** is dealing with the issues quickly once they happen. This requires:

Good **Feedback Loops** in the form of **Monitoring** and rapid response to issues.

tbd, talk with John about this

Meeting Reality

So in this second model, we are now considering that the world is a dangerous, untrustworthy place where *bad things happen*, either deliberately or accidentally. And, since we don't have a perfect understanding of the world, most of the **Production Risk** we face is **Hidden Risks**.

Putting software into production is **Meeting Reality** in the fullest way possible. The more contact we can give our system with the outside world, the more **Hidden Risks** will materialize. If we observe these and take action to mitigate them, then our system can get stronger. cybernetics, antifragile tbd.

It is tempting to delay **Meeting Reality** as long as possible, to "get your house in order". There is a tension between "you only get one chance to make a first impression" and "gilding the lilly" (perfectionism). In the past I've seen this stated as:

Pressure to ship becomes greater than pressure to improve tbd

A Risk-First reframing of this might be the balance between:

- The perceived Reputational Risk, **Feature Risk** and **Operational Risk** of going to production (pressure to improve)
- The perceived **Schedule Risks** (such as funding, time available, etc) of staying in development (pressure to ship)

The "should we ship?" decision is therefore a complex one. In **Meeting Reality**, we discussed that it's better to do this "sooner, more frequently, in smaller chunks and with feedback". We can meet **Operational Risk** on our own terms by doing so:

Meet Reality...	Techniques
Sooner	Limited Early-Access Programs, Beta Programs, Soft Launches
More Frequently In Smaller Chunks	[Continuous Delivery], Sprints Modular Releases Microservices Feature Toggles Trial Populations
With Feedback	User Communities, Support Groups, Monitoring, Logging, Analytics

External Events

We're familiar with the concept of taking steps on the Risk Landscape, wherein we **take action** to move to a position where the **Attendant Risks** are more acceptable to us. However, now we have to contend with the idea that external events *also* change the risk landscape too. If there is sudden bad weather, we might have a risk of a power-cut, and the losses that might entail to productivity or sales. If there is a change of government, that might impact the contracts we've written, or the security of our servers or staff.

Being *in production* is accepting that the **Risk Landscape** is a volatile, uncaring place. But it's worse than that, since we also have to contend with [Bad Actors], who are deliberately out to exploit weaknesses in the systems we build.

Ordinarily, when we transact with a [Dependency], it should be the case that after the transaction, there is value on both sides of the transaction. This could be, *you do my accounting*, I pay you money. On both sides, financial risks are reduced. If the price is too high, or too low, we see one or other side getting the better deal, and *capturing an unfair share of the value*.

With a [Bad Actor], we're in a situation more like a zero-sum game: value is *taken* from one party and *transferred* to the other. These are exactly the dependency relationships that societies *don't condone*: there is net zero or *negative* value in the transaction.

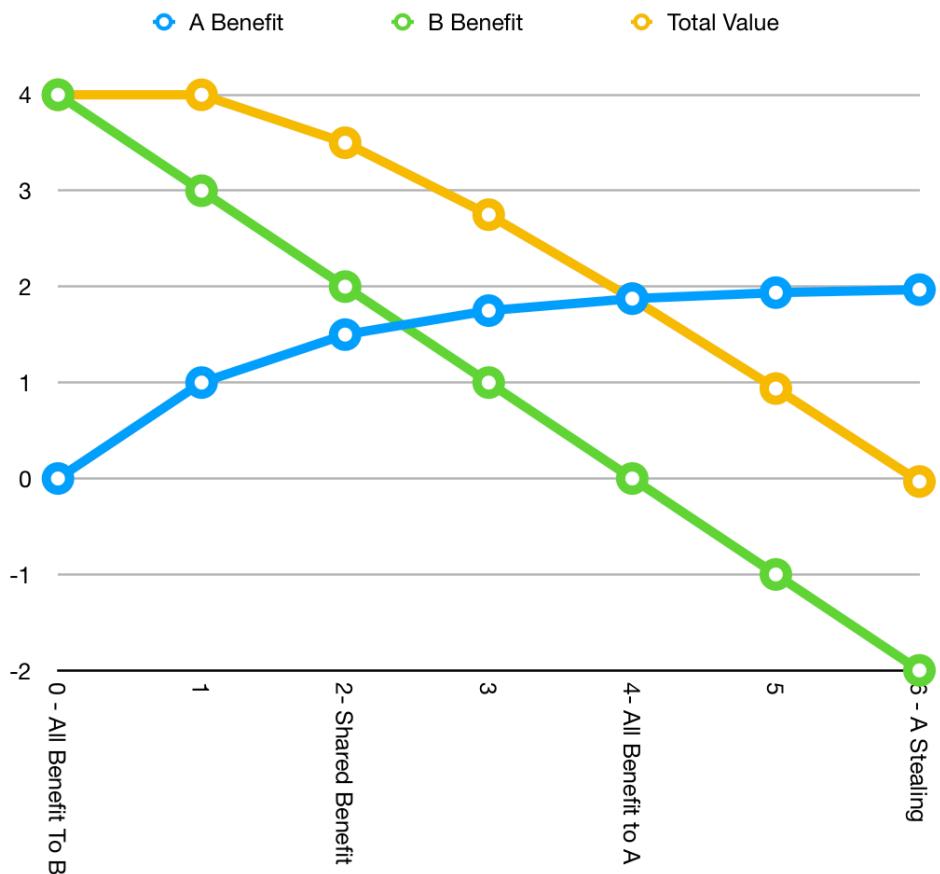


Figure 22.2: In this diagram, B does a deal with A, the value of the product B gets from A is always 4, but the price varies with the X axis. When the price is zero, B captures all the value, but as the price increases, it becomes a worse and worse deal for B, and the net value heads towards zero or negative.

- Regulatory Risk Legal Risk (Pestle?)

Weaknesses

Security of supply

Complex systems (ones which contain multiple, interacting parts, like the ones in the above diagrams) have to contend with their external environments, and try to minimize the ways in which they get interrupted from outside either by *Bad Actors* or external events. In the tbd

Interestingly, security is handled in very similar ways at all sorts of levels:

- **Walls:** defences *around* the complex system, to protect its parts from the external environment.
- **Doors:** ways to get *in* and *out* of the complex system, possibly with *locks*.
- **Guards:** to make sure only the right things go in and out. (i.e. to try and keep out *Bad Actors*).
- **Police:** to defend from *within* the system, against **Agency Risk** and *invaders*.
- **Subterfuge:** Hiding, camouflage, disguises, pretending to be something else. tbd

These work various levels in our own bodies: our *cells* have *cell walls* around them, and *cell membranes* that act as the guards to allow things in and out. Our *bodies* have *skin* to keep the world out, and we have *mouths*, *eyes*, *pores* and so on to allow things in and out. We have an *immune system* to act as the police.

Our societies work in similar ways: in medieval times, a city would have walls, guards and doors to keep out intruders. Nowadays, we have customs control, borders and passports.

We're waking up to the realisation that our software systems need to work the same way: we have **Firewalls** to protect our organisations, we lock down *ports* on servers to ensure there are the minimum number of doors to guard and we *police* the servers ourselves with monitoring tools and anti-virus software.

- Security Risk - Hacking - Denial Of Service - Security, Trust and Complexity - oWASp
- tbd, How much do compilers do for you? Now, they prevent many kinds of security error. Libraries too.

Operational Risk

When processes fail, this is called *Operational Risk*:

tbd - Wikipedia definition

This is a very specific name for **Reliability Risk** with regard to processes. In the UK each year, X number of people are killed in car accidents. If you regard driving a car from A to B as a process, then you could say that car accidents are Operational Risk. Why do we tolerate such costly operational risk in the UK. Could it be reduced? Well, yes. There are lots of ways. One way is that we could just reduce the speed limit.

It is interesting that we *don't* do that: although we know the driving process fails, and fails in a way that is costly to human lives, as a society we value the freedom, the economic efficiency and time savings that come from not mitigating this operational risk. Changing the speed

limit would have its own risks, of course: there would be a complicated transition to manage. However, if ten times as many people were killed in car accidents, and it was shown that reducing the speed limit would help, maybe it would be done. The **Operational Risk** would outweigh the **Schedule Risk**.

The point of this is that we *accept Operational Risk* as we go. However, if opportunities rise to mitigate it, which don't leave us with a net risk increase elsewhere, we'll make those improvements.

tbd. diagram version 3.

Mitigating **Security Risk** is a trade-off. You can spend *a lot* of time and effort on this, only to never face the

secrets: how to mitigate this

Effect / Impact

Sometimes, it's possible to measure the impact of Operational Risks. For example, if a software system fails, and leaves customers unable to access it, this can have a measurable financial impact in lost revenues or damages. Car recall example tbd. - fight club

Impact is usually proportional to some of the below variables:

- Number of customers affected.
- Number of transactions affected.
- Size of the transactions
- Length of time systems were affected.

stuff that can go wrong in production

changing stuff in production is harder than changing it in test, as you have to *migrate*.

all the costs of breaking stuff, and damaging the running of the business.

reputation damage (you only get one chance to make a first impression)

- You don't know all the ways the software will get used in production.
- Different browsers, versions of code, accessibility.
- Can you support all the users? Is there enough kit? Will you know?

Correlation - Upgrades (tell story of Research upgrade that went wrong because we were upgrading at the same time as an outage) - Single points of failure.

reputational damage

Reaction & Recovery

Reliability Risk

- Feedback Loops - Bug reports, feedback - Quality of feedback - Internal Controls - Agency Risk meets Production Risk (bad actors, controls)
 - Contingency Planning
 - Disaster Recovery
 - Performance Degradation / Runaway processes (Performance Risk)
 - Support (trade off – promptness vs ability)

Sometimes, the reaction of the company makes things worse – streisand effect? others?

- Poor monitoring, visibility risk meets operational risk (otherwise, it doesn't matter)
- Correlation (need a good example here)
- Monitoring Tools and Logs

Prevention

- How we learn from our mistakes
- You can't know everything
- Reality changes anyway

Performance Risk

There is a lot more to Operational Risk. Here, we've touched on it, and sketched the edges of it enough for it to be familiar and fit in our framework.

Reputational Risk

High-Profile Cases

Maturity

<https://math.nist.gov/IFIP-UQSC-2011/slides/Oberkampf.pdf> <https://www.bsimm.com/framework/intelligence-models.html> ISO27001

OWASP

Chapter 23

Staging And Classifying

Our tour is complete.

We've collected on this journey around the **Risk Landscape** a (hopefully) good, representative sample of **Risks** and where to find them. But if we are good collectors, then before we're done we should **Stage** our collection on some solid [Mounting Boards], and do some work in classifying what we've seen.

tbd collecting image

If you've been reading closely, you'll notice that a number of themes come up again and again within the different sections. For example, concepts like **Fit**, **Abstraction**, **Evolution**. Although we've been looking at patterns of risk across software projects, it's time to look at the *patterns within the patterns*.

First, A Recap

tbd. list of risks, broken down

Some Observations

1. The Power Of Abstractions

Abstraction appears as a concept continually throughout the book, whether we are looking at **Communication**, **Complexity Metrics**, **Map and Territory Risk** or how it causes **Boundary Risk**. And, so far, we've looked at some complicated examples of abstractions, such as **network protocols**, **dependencies on technology** or **Business Processes**.

There's a good reason for this repetition. Abstraction is at the heart of *everything we do within software*. So, let's now *generalize* what is happening with abstraction, but have in mind a *really simple example*: giving a name to something.

Inventing an Abstraction means:

- **Creating a Feature.** So, at the simplest end, you might be simply *naming a pattern* of behaviour we see in the real world, such as “Binge Watching” or “Remote Working”, or naming your dog, “Alfie”. These abstractions are **Features** in the sense that other people can choose to use them, if they fit their requirements.
- **Creating a Boundary.** By naming something, you *implicitly* create a boundary, because the world is now divided into “things which *are X*” and “things which *are not X*”. Sometimes, this abstraction may literally end up having a physical boundary to enforce this division (such as, “My Property / Not My Property”). **Boundary Risk** is created by abstractions.
- **Creating a Protocol.** At the very simplest level (again), this is just introducing *new words to a language*. Therefore, we create **Protocol Risk**: what if the person we are communicating with _doesn’t know this word?
- **Increasing Complexity.** Because, the more words we have, the more complex the language is.

tbd, diagram.

Choosing an Abstraction means:

- **Overcoming a Learning Curve****:** Because you have to *learn* a name in order to use it (whether a function, a dog, or the name of someone at a party).
- **Accepting Boundary Risks.** Just using *a single word* means accepting the whole *ecosystem* of the language the word is in. Using *french words* means the **Boundary Risk** of the **French Language**.
- **Accepting Map And Territory Risk.** Because the word refers to the *concept* of the thing, and *not the thing itself*.

tbd. diagram

Depending On an Abstraction means:

- **Living with Dependency Risk:** We depend on a word in our language, or a function in our library, or a service on the internet. But all of these things are *unreliable*. The word might not communicate what you want it to, or be understood by the audience, the function might not work, the service might be down.
- **Accepting Evolution.** That language *changes and evolves*, and the words you are using now might not always mean what you want them to mean. Software too changes and evolves. We’ve seen this in **Red Queen Risk**, **Feature Drift Risk**, **Regression Risk** and **Protocol Risk**.
- **Accepting Coordination Risk.** Where a **Dependency** has **Agency**, we need to consider how to work with it.

tbd diagram.

2. Your Feature Risk is Someone Else's Dependency Risk

In the **Feature Risk** section, we looked at the problems of *supplying a dependency to someone else*: you've got to satisfy a demand **Market Risk**, and ensure a close fit with requirements **Conceptual Integrity Risk**. The section on **Production Risk** went further, looking at specific aspects of being the supplier of an IT service as a dependency.

However, over the rest of the **Dependency Risk** sections, we looked at this from the point of view of *being a client to someone else*: you want to find trustworthy, reliable dependencies that don't give up when you least want them to.

So **Feature Risk** is **Dependency Risk**: they're *two sides of the same coin*. You face **Dependency Risk** when you're a client, **Feature Risk** when you're the supplier.

We've looked at three dimensions of **Feature Risk**: - Fit - Evolution - Audience
(recap this)

Dependency Risk has three dimensions too:

- Schedule (things happening in time, running out of stuff. e.g I need enough money to get this done, etc. I need enough patience, enough loyalty, trust, entente.)
- Technology (depending on software, hardware, etc.)
- Organisation (arrangements of **people, processes**, product)
- Agency (people/machines/processes/organisations doing what you asked, working with you)
-

Using a dependency requires learning a *protocol*. You have to learn to use it. Maybe it learns you. This requires changes to your internal model.

Internal Model - Communication – Dependency Goal

3. Coordination Risk As Eden

- similar to _threading/deadlocking issues
- Coordination is how you deal with abstractions. and this means communication.

tbd diagram: abstractions -> agency -> coordination -> communication

One thing that should be apparent is that there are similarities in the risks between all the kinds of

- draw a diagram of this system. mark on all the different risks using numbers. mention specifically that since this is a diagram, it is a "map".

Towards A Periodic Table Of Risks

tbd, diagram, explanatory text.

Patterns

As we said **at the start**, Risk-First is all about developing *A Pattern Language*. We can use the terms like “*Feature Risk*” or “*Learning Curve Risk*” to *explain* phenomena we see on software projects. If we want to **De-Risk** our work, we need to be able to explain what the risks are, and what we expect to do about them.

So, lastly in part 2 let’s put our language to work, and look at some past project failures. Can we apply our lexicon to them?

On to **Stories Of Failure**.

Chapter 24

Stories Of Failure

2. Failure Modes

- Understanding Failure: what exactly does it mean to fail?
 - Personal Failures
 - * CapsLock: complexity, not using tools.
 - * Configuration Tool (Complexity, feature fit, bugs in hibernate (dependency risk, then dead-end risk), difficulty mapping domain model)
 - * Wide Learning (Funding, but also complexity), did we know what we were building? Agency risk
 - * AreAye - needless complexity XMLBox
 - * Agora: Notes / Typing. (Complexity Risk) Archipelago
 - * PDC: website redesign. funding. i.e. schedule risk
 - * Hawk: complexity risk in the software. but actually, they made it work. off-shoring.
 - * Dark: market/feature fit?
 - * Jio: marketing / market fit / Complexity in spades. algorithmic complexity
 - * DSL: complexity (code generation). complexity = layers. team dynamics.
 - * REF: complexity. agency risk. failure of goals. m&t.
 - * REF Testing: complexity risk. communication risk?
 - * HSC: Trader Comments: feature fit.
 - * HSC: Takeover of Symph: Complexity (of change)
 - * TT: Feature Fit
 - Boehm.

https://wwwx.cs.unc.edu/~welch/class/comp145/media/docs/Boehm_Term_NE_Fail.pdf

<https://www.worksoft.com/top-software-failures-of-2017-so-far>

<https://sites.hks.harvard.edu/m-rccb/ethiopia/Publications/Top%20Reasons%20Why%20Systems%20Fail>

JC Example 1: Compensation Workbench, rules to uplift payrolls annually based on certain

definitions. e.g tied into performance grading, min wage. Was done off s/s, out of HR system. Analysis done on S/s. But wanted to do it in the HR database. It took a lot longer than it should. Underestimated the complexity. Excel could be infinitely complex, but if you standardize, you lose that. The customer hadn't got experience of big IT projects - scope creep. No real sense of prioritisation and focusing on what was important. PM was nice, but didn't understand delivery management. (i.e. managing risks and issues). Had short timescale, lead delivery resource underestimated. 2m exercise. Lasted 5m in the end, but not with full functionality.

Part III

Preview

Glossary

Abstraction

Feedback Loop

Goal In Mind

Internal Model

The most common use for **Internal Model** is to refer to the model of reality that you or I carry around in our heads. You can regard the concept of **Internal Model** as being what you *know* and what you *think* about a certain situation.

Obviously, because we've all had different experiences, and our brains are wired up differently, everyone will have a different **Internal Model** of reality.

Alternatively, we can use the term **Internal Model** to consider other viewpoints: - Within an organisation, we might consider the **Internal Model** of a *team of people* to be the shared knowledge, values and working practices of that team. - Within a software system, we might consider the **Internal Model** of a single processor, and what knowledge it has of the world. - A codebase is a team's **Internal Model** written down and encoded as software.

An internal model *represents* reality: reality is made of atoms, whereas the internal model is information.

Meet Reality

Risk

Attendant Risk

Hidden Risk

Mitigated Risk

Take Action